

ENGSCI YEAR 3 FALL 2022 NOTES

BRIAN CHEN

Division of Engineering Science

University of Toronto

<https://chenbrian.ca>

brianchen.chen@mail.utoronto.ca

Contents

I ECE349: Introduction to Energy Systems	1
1 Admin stuff	1
1.1 Lecture 1	1
1.1.1 Mark breakdown	1
2 AC Steady State Analysis	1
2.1 Lecture 2	1
2.1.1 TODO	1
2.2 Lecture 3	3
3 AC Power	4
3.1 Lecture 4	5
3.1.1 Root Mean Squared (RMS) Values	6
3.2 Lecture 5: Multi-Phase AC	7
3.3 Lecture 6: Y and Delta connections	10
3.4 Lecture 7: DC-DC conversion	12
3.5 Lecture 8: Average output voltage	15
3.5.1 Filtering	17
3.6 Lecture 9	18
3.6.1 Inductor VoH-Seconds balance	20
3.6.2 Capacitor charge balance	20
3.7 Lecture 10	21
3.7.1 Design Process	23
3.8 Lecture 11	24
3.9 Lecture 12	26
II ECE352: Computer Organization	26
4 Admin stuff	26
4.1 Lecture 1	26
4.1.1 Mark breakdown	26
5 Preliminary	26
5.1 Lecture 2: Using binary quantities to represent other things	27
5.1.1 Floating Point Numbers	27

6	Memory and NIOS II	28
6.1	Lecture 3: Behavioural Model of Memory	28
6.1.1	Memory	29
6.1.2	Physical Interface	30
6.2	Lecture 4: NIOS II Programming Model	31
6.2.1	Adding Two Numbers	32
6.2.2	Adding two numbers using memory	34
7	Assembly Basics	36
7.1	Lecture 5: Simple Control Flow	36
7.2	Lecture 6, 7: For loops and arrays	38
7.3	Lecture 8: Subroutines	40
7.4	Lecture 9	45
7.5	Lecture 10: Recursive Subroutines	47
7.6	Lecture 11: Structs and recursive structures	48
8	GPIO	53
8.1	Lecture 12: Devices	53
8.1.1	The PIT implementation	56
8.1.2	Memory	57
8.2	Lecture 12: UART	59
8.3	Lecture 13: Interrupts & UART	63
8.4	Lecture 14: Timer	64
8.5	Lecture 15: Code Races	66
8.5.1	Buffer Overflow Stack Attacks	67
8.6	Lecture 16: Emulating instructions in software	67
8.7	Lecture 17: A single cycle processor	73
8.8	Lecture 18: Modifying the single cycle processor	83
8.9	Lecture 19: Multi-cycle processor	83
8.10	Lecture 20: Modifying the multi-cycle processor	90
8.11	Caches	90
8.11.1	Implementations	91
8.11.2	Handling Writes	93
III	ECE355: Signal Analysis and Communication	93
9	Admin and Preliminary	93
9.1	Lecture 1	93
9.1.1	Mark Breakdown	93
10	Transformations	94
10.1	Lecture 2	94
10.2	Lecture 3	95
10.2.1	General Continuous Complex Exponential Signals	95

11 Basic Signals	96
11.1 Lecture 4: Step and Impulse Functions	96
11.2 Lecture 5	98
11.2.1 Types of systems	99
11.2.2 System properties	100
12 LTI systems	101
12.1 Lecture 6: Discrete LTI systems	101
12.2 Lecture 7: Continuous LTI systems	102
12.3 Lecture 12	103
12.4 Lecture 15: Which periodic signals have a Fourier representation?	104
12.5 Fourier series representation of continuous periodic systems	105
IV ECE358: Foundations of Computing	106
13 Admin and Preliminary Maths	106
13.1 Lecture 1	106
13.1.1 Mark Breakdown	106
13.2 Complexities	106
13.3 Lecture 3: Logs & Sums	109
13.3.1 Functional Iteration	109
13.4 Lecture 4: Induction & Contradiction	110
13.4.1 Induction	110
13.4.2 Contradiction	111
13.5 Lecture 5: Recurrences	112
13.5.1 Recurrence Trees	113
13.5.2 Substitution	114
13.5.3 Master Theorem	114
14 Graphs	115
15 Trees	117
16 Lecture 7, 8: Probability and Counting	117
17 Sorting	118
17.1 Heapsort	118
17.2 Placeholder for Quick, counting, selection, and radix sort	119
17.3 BSTs	119
18 Amortized Analysis	120
18.1 Skewed Heaps	122
19 Minimum Spanning Trees	122
20 Shortest Paths	123

21 Splay Trees	125
22 Maximum Flow	129
23 Theory of Computation	132
23.1 Finite Automata	132
23.2 Turing Machines	134
23.3 Complexity Classes	135
23.4 NP Completeness	135
23.4.1 Solving NPC problems	136
V ECE360: Electronics	138
24 Admin and Preliminary	138
24.1 Lecture 1	139
24.1.1 Mark Breakdown	139
24.1.2 Diodes	139
25 Diodes	140
25.1 Lecture 2	140
25.2 Lecture 3	142
26 Lecture 4 & 5: Forward conducting diodes	144
27 Small-Signal Model	146
28 Lecture 6: Small signal model, cont'd	147
28.0.1 Deriving small-signal resistance	148
29 Lecture 7	150
30 Rectifiers	151
30.1 Lecture 8	153
30.2 Lecture 9	156
30.3 Lecture 10	159
30.4 Lecture 11: Introduction to transistors and MOSFETs	162
30.5 Lecture 12: MOSFET Operating Regions	164
30.6 Lecture 13: MOSFET current-voltage characteristics	167
VI MAT389: Complex Analysis	167
31 Complex Numbers	167
31.1 Lecture 1	167
31.2 Lecture 2	169
31.2.1 Functions on complex planes	170
31.2.2 Exponential Functions	171

31.3	Lecture 3: Exponent and Logarithm	173
31.3.1	Exponential	173
31.3.2	Logarithm	174
31.3.3	Powers	175
31.4	Lecture 4	175
31.4.1	Properties of complex derivative	177
31.5	Power series	181
31.6	Lecture 5	182

VII ECE444: Software Engineering 182

32 Preliminary 182

32.1	Lecture 1, 2	183
------	--------------	-----

33 Project Management 183

33.1	Lecture 3	183
33.1.1	Agile	184
33.2	I dropped this course	184

VIII ESC301: Seminar 184

ECE349: *Introduction to Energy Systems*

SECTION 1

Taught by Prof. P. Lehn

Admin stuff

SUBSECTION 1.1

Lecture 1

First lecture was logistical info + a speil about how power systems are one of the great modern wonders. Course will cover sinusoidal AC power systems (1, 3 phase), power systems (dc-dc, dc-ac conversion), and magnetic systems (transformers, actuators, and synchronous machines)

1.1.1 Mark breakdown

- 50 % Final
- 25 % Midterm
- 5 % Quiz
- 15 % Labs
- 5 % Assignments

SECTION 2

AC Steady State Analysis

SUBSECTION 2.1

Lecture 2

2.1.1 TODO

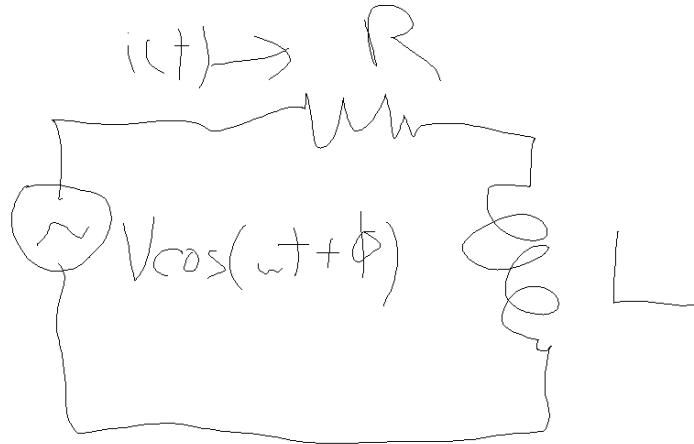
- Review Thomas 669-600

What we have learnt prior for differential equations enables us to arrive at analytical solutions to linear stable AC systems with phasors. A homogeneous and particular solution will be produced. If there's a stable homogeneous solution, $\rightarrow 0$ as $t \rightarrow \infty$. The full solution would be the addition of the two via super position.

We generally use this approach to solve circuits since it's an efficient way to solve circuits and make them into essentially DC circuits.

Recall, for a general phasor \hat{P}

- $\frac{d\hat{P}}{dt} = jw\hat{P}$
- $\int \hat{P} = \frac{1}{jw} \hat{P}$



$$Ri + L \frac{di}{dt} = V \cos(\omega t + \phi) \quad (2.1)$$

But this is a pain to solve. It can be made simpler by applying phasors

$$V \cos(\omega t + \phi) = \operatorname{Re}\{V e^{j(\omega t + \phi)}\} \quad (2.2)$$

Take the real part of \hat{I} :

$$R\hat{I} + L \frac{d\hat{I}}{dt} = V e^{j(\omega t + \phi)} \quad (2.3)$$

And therefore by inspection the solution is of format $\hat{I} e^{j\omega t}$, where \hat{I} is a phasor. Noting that \hat{I} contains only amplitude and phase,

$$\begin{aligned} R\hat{I} e^{j\omega t} + L \frac{d}{dt}(\hat{I} e^{j\omega t}) &= V e^{j\omega t + \phi} \\ R\hat{I} + L\hat{I}j\omega &= V e^{j\phi} \end{aligned} \quad (2.4)$$

And now reconstructing:

$$\begin{aligned} \hat{I} &= \frac{V}{\sqrt{R^2(\omega L)^2}} e^{j(\omega t + \phi - \tan^{-1} \frac{\omega L}{R})} \\ i(t) &= \operatorname{Re} \left\{ \hat{I} \right\} \end{aligned} \quad (2.5)$$

And therefore

$$\hat{I} = \frac{V}{\sqrt{R^2(\omega L)^2}} \cos \left(\omega t + \phi - \tan^{-1} \left(\frac{\omega L}{R} \right) \right) \quad (2.6)$$

The steps to solving a phasor problem are:

Notation: $X e^{j\phi} \leftrightarrow X | \phi$

- Define phasor: $V \cos(\omega t + \phi) \leftrightarrow V | \phi$
- Map L, C into phasor domain; find impedances
 - $v = L \frac{di}{dt} \leftrightarrow \hat{V} = j\omega L \hat{I}$

$$-i = C \frac{dv}{dt} \leftrightarrow \hat{V} = \frac{1}{j\omega C} \hat{I}$$

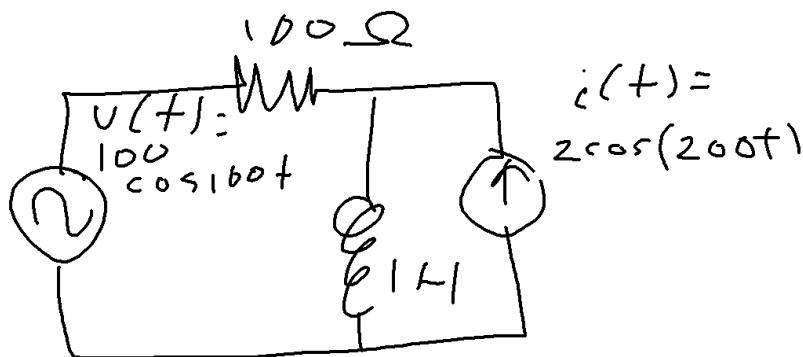
- Do mesh analysis to find \hat{I} ; $\hat{I} = \frac{\hat{V}}{\sum \text{impedances}}$

- Reconstruct $i(t)$ from \hat{I}

SUBSECTION 2.2

Lecture 3

Phasors allow us to solve circuits with multiple sources of differing frequencies.



To find the current $i(t)$ over the inductor we can find its response due to the voltage and current sources and then apply superposition.

- $I_1 = \frac{100|0}{100+j100} = 0.707|-45^\circ \rightarrow i_1(t) = 0.707 \cos(100t - 45^\circ)$
- $I_2 = \frac{100}{100+j200} 2|0 = 0.894|-65^\circ \rightarrow i_2(t) = 0.894 \cos(200t - 63^\circ)$
- $i(t) = i_1(t) + i_2(t) = 0.707 \cos(100t - 45^\circ) + 0.894 \cos(200t - 63^\circ)$

Non-sinusoidal stimulus may be solved by decomposing the signal with Fourier transforms. For example, square waves:

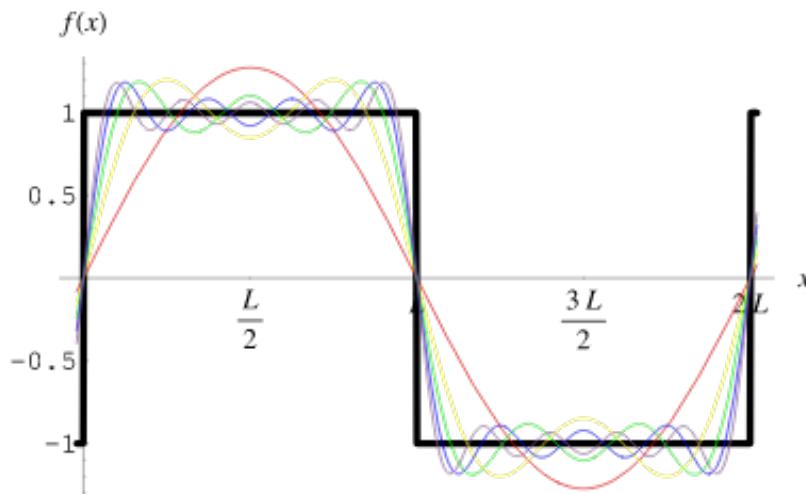


Figure 1. square waves with Fourier series superimposed

The general form of a Fourier transform is given as:

$$v_{equiv}(t) = a_o + \sum_{n=1}^{\infty} a_k \cos(nw_o t) + b_k \cos(nw_o t) \quad (2.7)$$

Where:

$$\begin{aligned} a_o &= \frac{1}{T} \int_0^T v(t) dt \\ a_k &= \frac{2}{T} \int_0^T v(t) \cos(nw_o t) dt \\ b_k &= \frac{2}{T} \int_0^T v(t) \sin(nw_o t) dt \end{aligned} \quad (2.8)$$

Armed with Fourier series and superposition we may now model a non-sinusoidal signal as a superposition of an infinite sum of sources. About half the work can be cut in half by recognizing that \sin lags \cos by 90° , so

$$\begin{aligned} a_o &= \frac{1}{T} \int_0^T v(t) dt \\ a_k &= \frac{2}{T} \int_0^T v(t) \cos(nw_o t) dt \\ b_k &= \frac{2}{T} \int_0^T v(t) \cos(nw_o t - 90^\circ) dt \end{aligned} \quad (2.9)$$

SECTION 3

AC Power

Definition 1 **Instantaneous Power:** $p(t) = v(t) \times i(t) [W, \frac{J}{s}]$

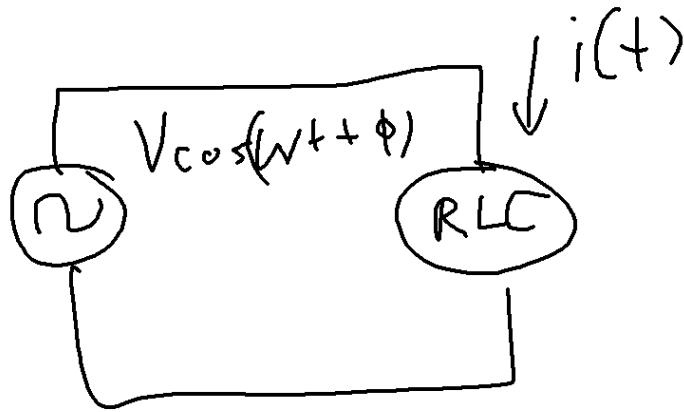
Example For a circuit with a voltage source, $v(t) = V \cos(\omega t)$ and a resistor Ω , $i(t) = I \cos(\omega t)$,

$$p(t) = VI \cos^2(\omega t) = \frac{VI}{2}(1 + \cos(2\omega t))$$

Definition 2 **Average Power over Cycle:** $P(t) = \frac{1}{T} \int_0^T p(t) dt = \frac{VI}{2}$

If we were to plot the instantaneous power we see that due to the sinusoidal response there are times where 0 power is supplied. This will always be true for a single phase power supply; real-world supplies always have multiple phases; this is why computer PSUs always contain a ton of capacitors.

Definition 3 **Reactive Power:** Q



If $\phi_i = 0$ and taking $\phi = \phi_v$,

$$\phi = \phi_v - \phi_i$$

$$\begin{aligned}
 p(t) &= V \cos(\omega t + \phi) * I \cos(\omega t) \\
 &= \frac{VI}{2} \cos(\phi) + \cos(2\omega t + \phi) \\
 &= \underbrace{\frac{VI}{2} \cos(\phi)(1 + \cos(2\omega t))}_{\text{real power e.g. heat}} - \underbrace{\frac{VI}{2} (\sin \phi \sin 2\omega t)}_{\text{stored and released back to source}}
 \end{aligned} \tag{3.1}$$

Taking the average power of the reactive power we get

$$P_{avg} = \frac{VI}{2} \cos \phi \tag{3.2}$$

Another quantity, reactive power, can be defined with regards to the energy sloshing back and forth:

$$Q = \frac{VI}{2} \sin \phi \tag{3.3}$$

SUBSECTION 3.1

Lecture 4

Definition 4

Displacement factor

$$DF \equiv \cos \phi$$

Where ϕ is the angle measured from the \hat{V} to the \hat{I} phasors. A $DF = 1$ means the system is transferring the most power possible.

- Lagging DF: ϕ is -'ve
- Leading DF: ϕ is +'ve

We can also write P, Q from phasors.

$$\begin{aligned}
P &= \frac{\hat{V}\hat{I}}{2} \cos(\phi_v - \phi_i) \\
&= \frac{1}{2} \hat{V}\hat{I} \operatorname{Re}\{e^{j\phi_v - \phi_i}\} \\
&= \frac{1}{2} \operatorname{Re}\{V e^{j\phi_v} \cdot I e^{-j\phi_i}\} \\
&= \frac{1}{2} \operatorname{Re}\{\hat{V}\hat{I}^*\}
\end{aligned} \tag{3.4}$$

$$\begin{aligned}
Q &= \frac{\hat{V}\hat{I}}{2} \sin(\phi_v - \phi_i) \\
&\vdots \\
&= \frac{1}{2} \operatorname{Im}\{\hat{V}\hat{I}^*\}
\end{aligned} \tag{3.5}$$

The expressions for Q and P are basically the same so we define a new quantity, complex power, which incorporates both the imaginary and real components.

Reference: Thomas 16.1, 16.2, 16.3

Definition 5

Complex Power

$$S = \frac{1}{2} \hat{V}\hat{I}^* = P + jQ \tag{3.6}$$

3.1.1 Root Mean Squared (RMS) Values

A RMS value measures the average power of a periodic signal. Consider a circuit with an AC voltage source with $v(t)$ flowing through a resistor.

The power is:

$$p(t) = v(t)i(t) = v(t) \frac{v(t)}{R} = \frac{1}{R} v(t)^2 \tag{3.7}$$

The average power is therefore

$$P = \frac{1}{R} \int_0^T v(t)^2 dt \tag{3.8}$$

Evaluating this becomes easier by defining an useful quantity, RMS voltage

$$v_{rms} = \sqrt{\frac{1}{T} \int_0^T v(t)^2 dt} \tag{3.9}$$

And using v_{rms} we get a nice expression for average power,

$$P = \frac{1}{R} v_{rms}^2 \tag{3.10}$$

More generally speaking we can define RMS values of sinusoidal signals

Definition 6

$$v_{rms} = \sqrt{\frac{1}{2\pi} \int_0^{2\pi} (\hat{V} \cos wt)^2 dt} = \frac{1}{\sqrt{2}} \hat{V} \tag{3.11}$$

Plugging this into the expressions for complex power:

$$S = \underline{V} \cdot \underline{I}^* \quad (3.12)$$

Where \underline{V} , \underline{I}^* are RMS phasors at a given common frequency.

And more generally yet, the RMS values of non-sinusoidal signals can be found with help of a Fourier expansion

Definition 7

Let $v(t) = \hat{v}_0 + \hat{v}_1 \cos(wt + \phi_1) + \dots$

Then,

$$v_{rms} = \sqrt{\hat{v}^2 + \frac{\hat{v}_1^2}{\sqrt{2}} + \frac{\hat{v}_2^2}{\sqrt{2}} \dots} = \sqrt{v_0^2 + \sum_{n=1}^{\infty} \left(\frac{V_m^2}{2} \right)} \quad (3.13)$$

And if V_m is the *rms* value,

$$v_{rms} = \sqrt{v_0^2 + \sum_{n=1}^{\infty} V_m^2} \quad (3.14)$$

One thing to watch out for is that we could have a system with high voltage but near-zero current transferring little to no power. To account for this we look at the power factor 'PF'

Definition 8

Power Factor

$$PF \equiv \frac{\text{average power}}{\text{rms voltage} \cdot \text{rms current}} \quad (3.15)$$

For a sinusoidal V, I :

$$\begin{aligned} PF &= \frac{\frac{1}{2} \hat{V} \hat{I} \cos \phi}{\frac{\hat{V}}{\sqrt{2}} \cdot \frac{\hat{I}}{\sqrt{2}}} \\ &= \cos \phi \end{aligned} \quad (3.16)$$

If signals are not harmonics $PF = DF$. This can be a source of confusion.

For non-sinusoidal systems this becomes more difficult because, unlike pure signals, they may contain harmonics. In these systems either V, I , or both may contain harmonics. Generally in household power V is clean but I contains harmonics.

The effect of harmonics in currents is that $I_{\text{harmonics}}$ causes a higher I_{rms} ; there is more current but no higher power! This reduces PF and causes $PF < DF$.

To summarize,

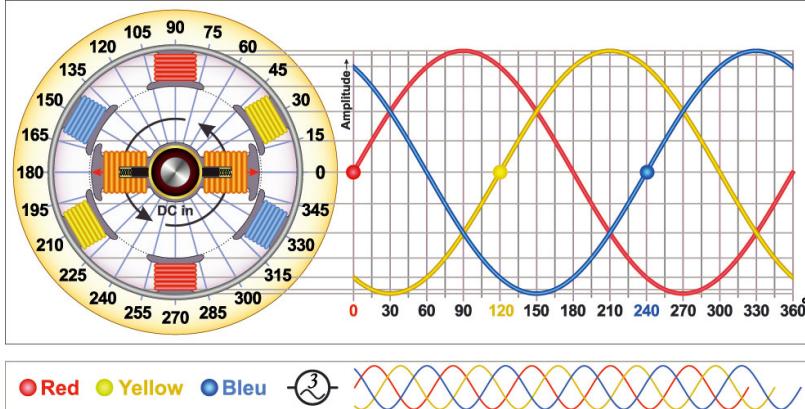
- In ideal systems, $PF = DF$ if all V and I are at one frequency.
- $PF = 1$ means no energy sloshing between load and source.

Harmonics are bad because they reduce the usefulness of the system. There are very tight standards for how many harmonics one is allowed to inject into the system via generators or loads. See: textbook 16.6

SUBSECTION 3.2

Lecture 5: Multi-Phase AC

AC power is generated by spinning a magnet between some coils. Some pixies get excited and by some Maxwell's equations and EMF and ECE259 we get a voltage induced in the coils.



Current from a generator with a single pair of coils is *single phase*. Most generators, like the one in the picture, have three pairs of coils and therefore generate three phases. For a typical three-phase setup with coils arranged at 0° , 120° , 240° , the voltages can be found with a little bit of trigonometry.

$$\begin{aligned} v_a(t) &= \sqrt{2}V \cos \omega t \rightarrow \underline{V_a} = V|0 \\ v_b(t) &= \sqrt{2}V \cos(\omega t - 120) \rightarrow \underline{V_b} = V|-120^\circ \\ v_c(t) &= \sqrt{2}V \cos(\omega t - 240) \rightarrow \underline{V_c} = V|240^\circ \end{aligned} \quad (3.17)$$

And similar expressions may be derived for single-phase and two-phase power.

The reason why three-phase power is typically used has to do with efficiency of power transfer relative to the amount of wires and copper needed. Whereas single-phase power requires two wires to carry power and two-phase power requires four wires, three-phase power can be transmitted over 6, 4, or three wires. This saves a lot of copper as three-phase 3ϕ power can carry 50% more power than 2ϕ power for less copper.

In 3ϕ systems the voltages sum to zero; $v_a(t) + v_b(t) + v_c(t) = 0 \forall t$

Four-wire three-phase power:

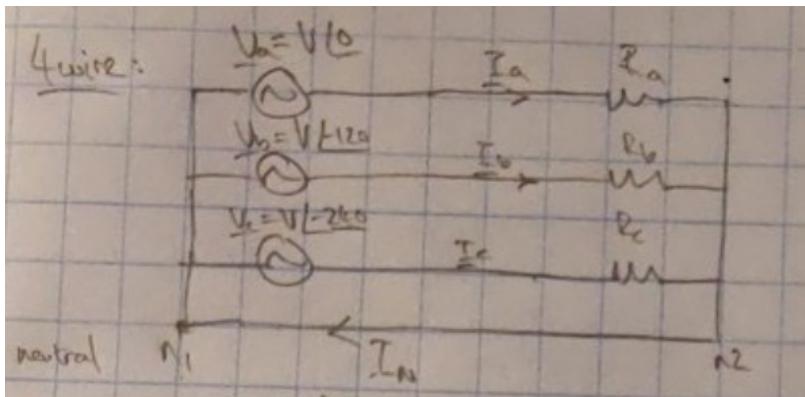


Figure 2. A four wire system for three phase power

If $R_a = R_b = R_c = R$, we have a balanced load and then the currents are related by $i_n = i_a(t) + i_b(t) + i_c(t) = 0$.

Three-wire systems drop the fourth neutral wire since it carries no current. This can be problematic if the load is not balanced; though $I_a + I_b + I_c = 0$ will still hold true since there is

Proof: just substitute the condition earlier that voltages sum to zero and the fact that all resistances are equal into $I = \frac{V}{R}$

no return path, the nodes at either end of the AC source and resistor pair would have differing voltages.

If the system is balanced then to save ourselves from drawing everything out all the time, only a single diagram is drawn for a characteristic phase and then solved once.

Example

$$Z_a = Z_b = Z_c = Z \quad (3.18)$$

$$\underline{V_b} \underline{V_a} e^{\frac{-j_2 n}{3}} \quad (3.19)$$

$$\underline{V_c} \underline{V_a} e^{\frac{-j_4 n}{3}} \quad (3.20)$$

Then,

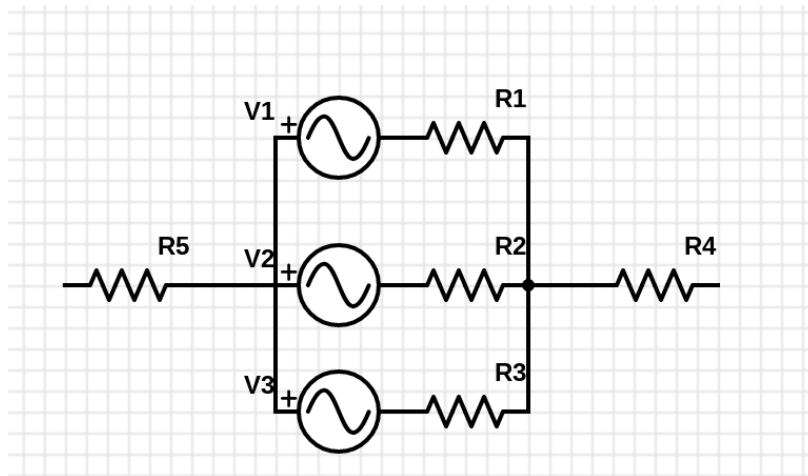
$$\underline{I_a} = \frac{\underline{V_a}}{Z} \quad (3.21)$$

$$\underline{I_b} = \frac{\underline{V_b}}{Z} = \frac{\underline{V_a}}{Z} e^{\frac{-j_2 n}{3}} \quad (3.22)$$

$$\underline{I_b} = \frac{\underline{V_c}}{Z} = \frac{\underline{V_a}}{Z} e^{\frac{-j_4 n}{3}} \quad (3.23)$$

And then the solutions for phase b and c are the same as that for phase a except for the $120^\circ, 240^\circ$ offsets.

Because of this property, instead of drawing three separate diagrams for each phase, we can just draw one diagram and then solve for the other two phases.



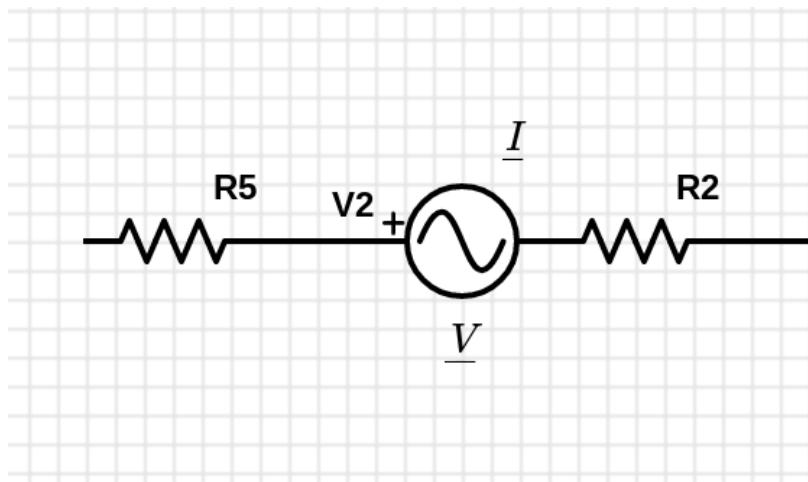


Figure 3. Condensed diagram for three phase power

SUBSECTION 3.3

Lecture 6: Y and Delta connections

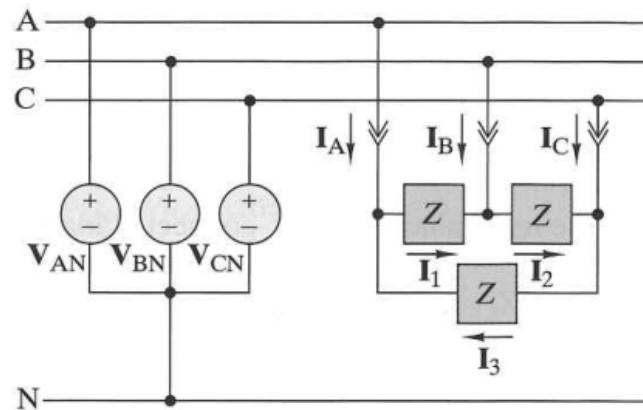


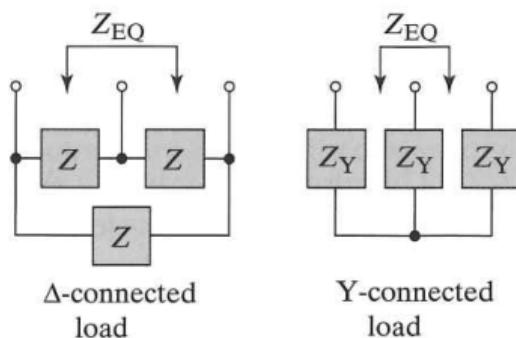
Figure 4. Three-phase system with a Y connected source and a Δ connected load

The current to ground I_N is always 0 for a Δ connected load. I_N is zero for a Y connected load if Y has no neutral connection and the load/source are balanced.

Theorem 1

$Y - \Delta$ conversion

Any Δ load can be converted to an equivalent Y connected load

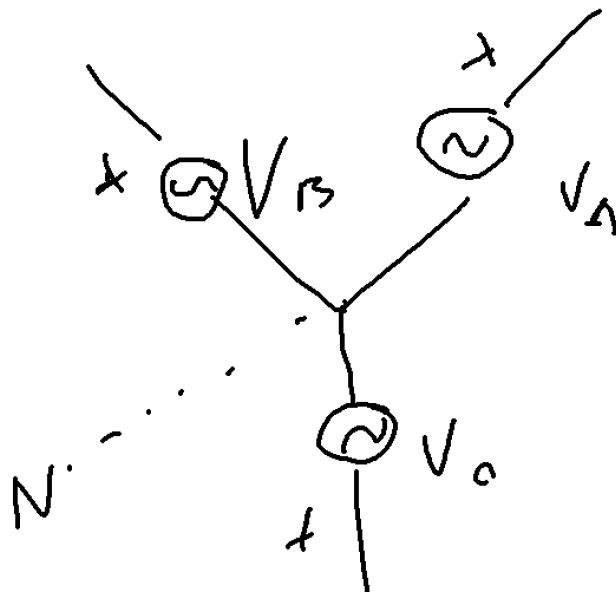


The derivation is in the textbook.

Sources can be connected in Y or Δ configurations.

Definition 9

Line currents are the currents on the lines a, b, c . Phase currents are the currents immediately beside the sources.

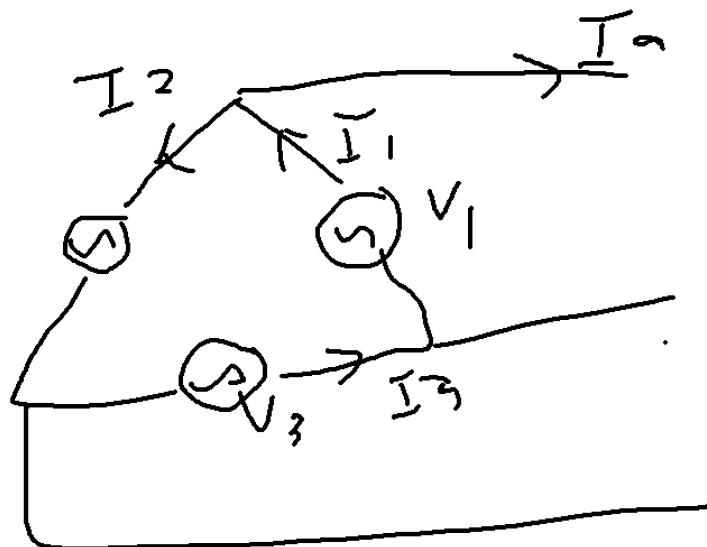


A Y connected source has the neutral line in the center.

$$\tilde{V}_{AB} = \tilde{V}_A - \tilde{V}_B = |\tilde{V}_A|(1 - 1[e^{-\frac{2\pi}{3}}]) = \sqrt{3}|\tilde{V}_A|e^{\frac{j\pi}{6}}$$

Angle differences are
0, 120, 240 degrees

The line and phase current are the exact same in the case of a Y connected source. However, the line and phase voltages are different and are related by (3.24).



A Δ connected source has no neutral point.

$$I_a = I_1 - I_2 = I_1(1 - 1|e^{-\frac{2\pi}{3}}) = \sqrt{3}I_1e^{\frac{j\pi}{6}} \quad (3.25)$$

The line and phase voltages are the exact same in the case of a Δ connected source. However, the line current and phase currents are different and are related by (3.25).

Since neutral is not always available, we write 3ϕ systems based on their line-to-line voltages. For example, a 208V system has 120V_{rms} on each phase.

Example A Y connected three-phase source has a line-to-line voltage, i.e. $V_{A \rightarrow B}$ of 208V but each source has 120V since $\frac{208}{\sqrt{3}} = 120V$. A similar argument can be applied for the Δ sources.

$$I = \frac{10}{\sqrt{3}} = 5.8A \quad (3.26)$$

SUBSECTION 3.4

Lecture 7: DC-DC conversion

Let's say we want to supply 10V DC to a load but our supply is at 20.5V. We have a few options. We can use a resistor to drop the voltage via a voltage divide, which we learned about in ECE159. However this is not only not very efficient, but also non-responsive to changes in resistance load and will provide incorrect supply voltage if the load is not as anticipated. Another option is to use **transistors** to regulate the power instead.

semiconductors to process power \Leftrightarrow power electronics

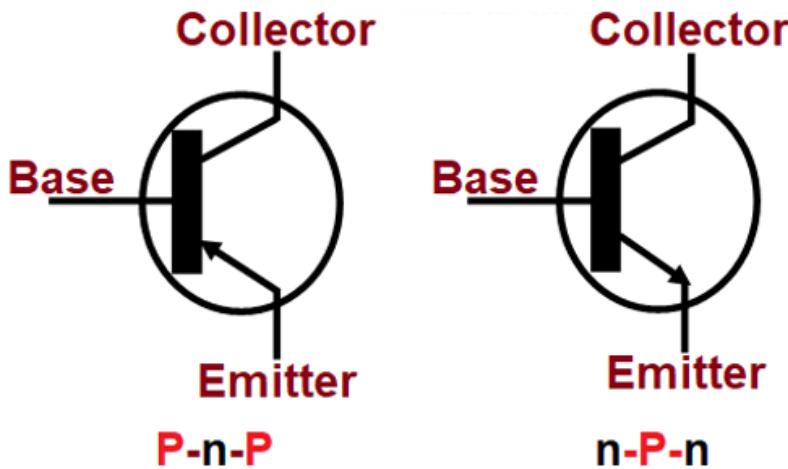


Figure 5. Transistors have a Base, Collector, and Emitter. At this point in the course, we just need to know that we can change the $V - I$ response by fiddling with the base current i_B

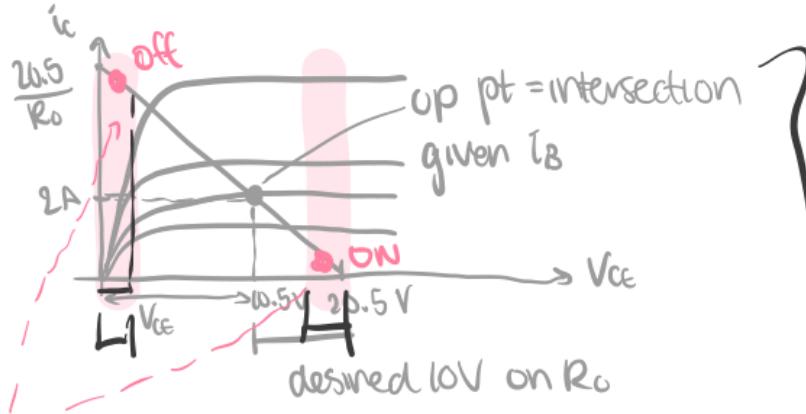


Figure 6. Thanks sherry for this figure

Finding the operating point that we want the transistor to be at can then be done graphically by solving a system of two equations with two unknowns, namely finding the intersection between the transistor response at a given i_B and the desired voltage/current of the load. Then the correct i_B can be selected.

However this is not very efficient since a lot of power is being wasted in the transistor. For a 20.5V source and a 10V load at 2A, $V_{CE} = 10.5V$, so $P_{loss} = 10.5 * 2 = 21W$ being lost in the transistor which results in the rather poor $\frac{20}{20.5*2=41} \rightarrow \eta = 49\%$ efficiency. So the transistor is a little bit better but we still lose half the power – which means that in a device this would halve the battery life and require a huge heat sink. How can we do better?

Switch-mode power is a way to maximize η by switching between two states that correspond to low power loss at a high frequency



The left state is the 'off' state and the right state is the 'on' state. Switching rapidly between 20.5V, 0A and 0.5V, 4A results in what is, on average what we want. This will also produce the same power output while increasing efficiency

What about the timing between the off/on states? This can be found by taking an integral to find the relationship between T_{on} and T_{off} with respect to power.

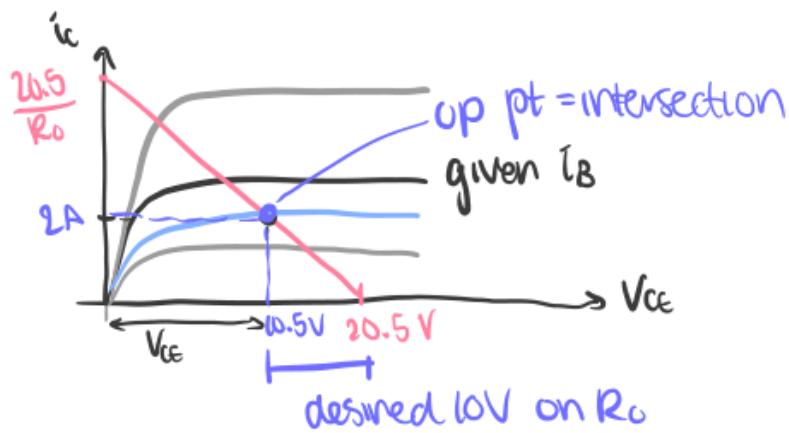
$$\begin{aligned}
 P &= \frac{1}{T} \int_0^T p(t) dt \\
 &= \frac{1}{T} \int_0^{T_{on}} i_c^2 R_o dt \\
 &= \frac{1}{T} \int_0^{T_{on}} 4^2 \cdot 5 dt \\
 &= 80 \frac{T_{on}}{T}
 \end{aligned} \tag{3.27}$$

So in order to transfer the same 20W to the 5Ω load, we would use $T_{on}/T = \frac{1}{4}$, which would correspond to $P_{in} = V_i \cdot T_{on}/T = 20.5W$. It then follows that $\eta = \frac{20}{20.5} = 97.6\%$ which is much better¹

¹ compared to the 49% before

A few points of caution:

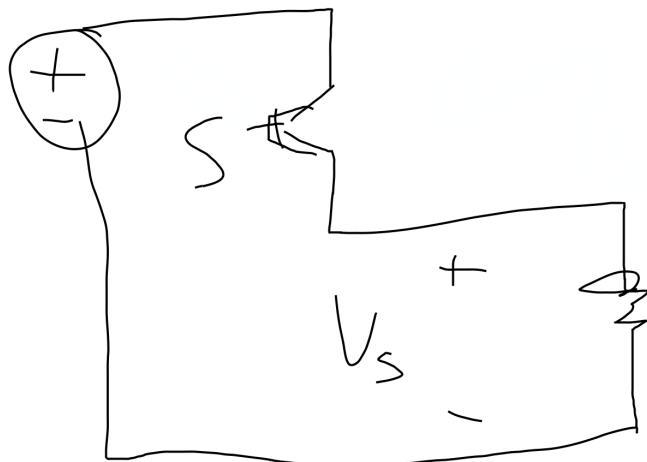
- A pulsing output voltage/power can be dangerous for some loads. For example motors aren't a fan of it. In these cases a filter can be applied to leave behind the average power.
- If we want to supply average voltage/power an on-time of 50% would be required.



SUBSECTION 3.5

Lecture 8: Average output voltage

Consider a circuit with a transistor and a load.



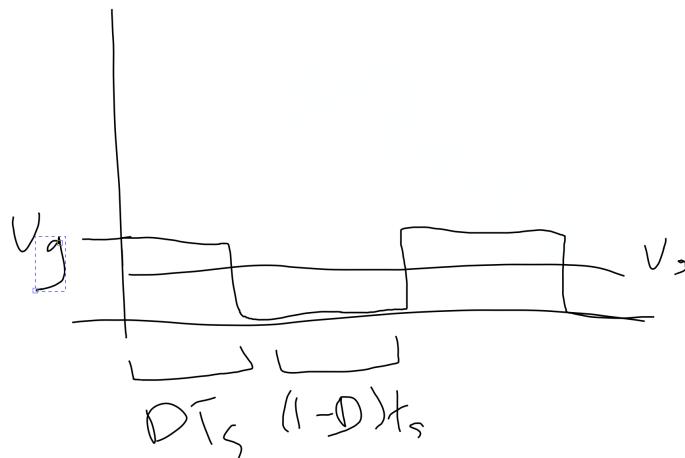
A switching signal S taking on a value from 0, 1 periodically over time is applied to the transistor

Definition 10

Duty cycle: the fraction of time that S is 1

$$D = \frac{T_{on}}{T_s} \quad (3.28)$$

The response of the circuit, where V_g is the source voltage and V_s being the voltage at the load would look like this:



One problem that we want to characterize is the ripple voltage, i.e. unwanted harmonics

$$v_s(t) = V_s + \sum_{k=1}^{\infty} V_s^k \sin\left(k \frac{2\pi}{T_s} t\right) \quad (3.29)$$

The V_s term is desirable, and the harmonic terms V_s^k are undesirable.

$$V_s^k = \frac{2V_g}{\pi} \sin(k\pi D) \quad (3.30)$$

Basically all Fourier analysis for power electronics will basically always result in an expression with a bunch of terms on one end and then a sin or cos. $\frac{1}{k} \frac{2V_g}{\pi}$ is the envelope of the harmonic terms. Note how the envelope decreases with k .

Plotting out the spectrum of signals we get:

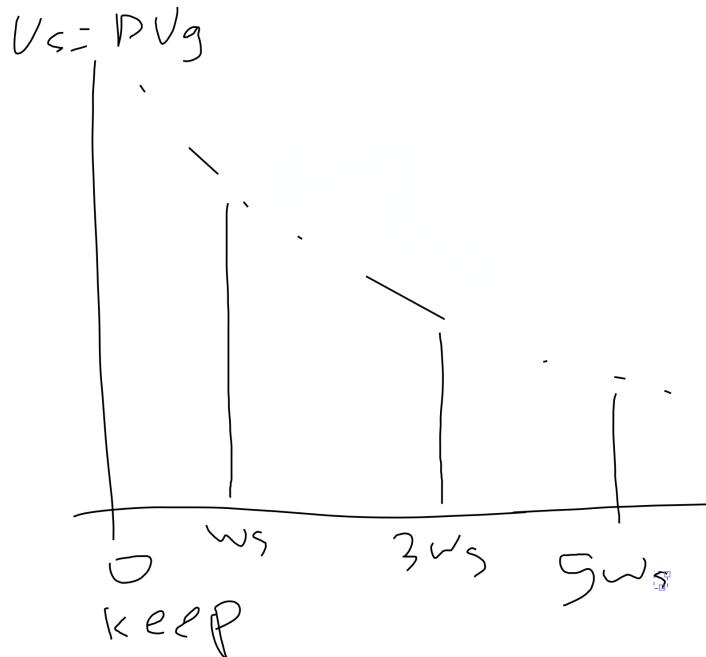
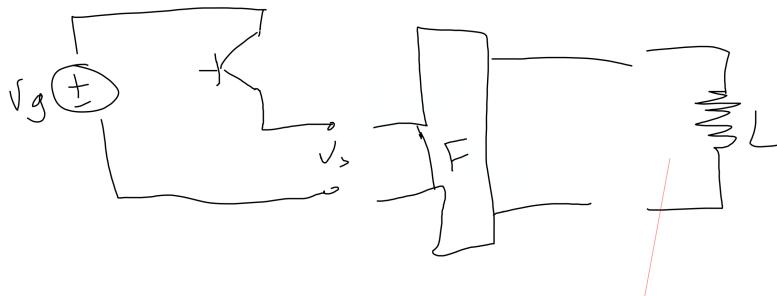


Figure 7. Voltage vs frequency. We want to keep the $\omega = 0$ term and then eliminate the rest with a filter. Note switching frequency $\omega_s = \frac{2\pi}{T_s}$

3.5.1 Filtering

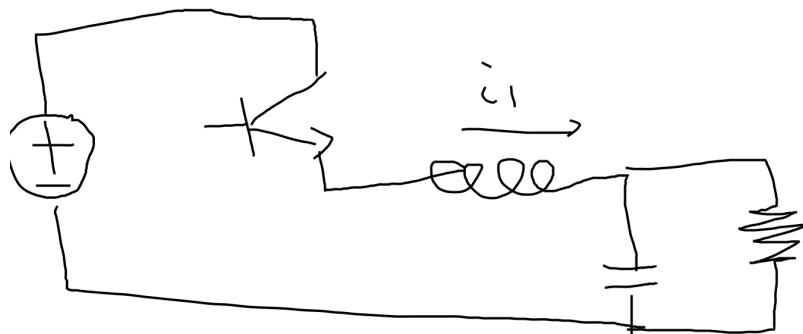
We'll need a "low pass" filter for dc/dc filtering.



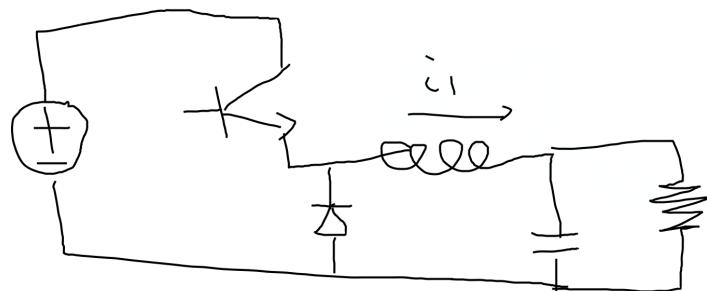
Choices for filters that may work for this system are the RC , LR , and LC filters. The LC is most promising because it is second order and it's theoretically lossless because there is no resistor.

We are going to want to maintain continuity of capacitor voltages and inductor currents. Things work out pretty well while we're on our duty cycle, but once the duty cycle turns off and the switch is opened, $i \rightarrow 0$. This is a problem because this will cause us to lose all the energy stored in the inductor every time the switch is opened, which can be a lot given the frequency that these devices operate at.

$$E = \frac{1}{2} L i^2 \quad (3.31)$$



In order to handle this, we must give the current somewhere to go while off-duty cycle; an energy-recovery diode which ensures that i_L is continuous and therefore E_L is lost.



Comment

Switching should never yield discontinuity in i_L or v_c .

Some design considerations:

1. Filters all store energy during 'on' states and releasing it during 'off' states in the magnetic fields of inductors and electric fields of capacitors.
2. Faster switching \leftrightarrow less stored energy needed \leftrightarrow smaller capacitor and inductor needed
 - Switching too fast would cause the system to spend too much time during switching in undesirable high-loss states

Capacitors tend to be cheaper so designs are biased to store energy in them instead of inductors

Solving these DC-DC converters with Fourier analysis and superposition every time can be a pain, so we'll be looking at other methods next lecture. Plus, Fourier analysis requires an known $v(t)$ or $i(t)$ sources, which is not always the case. To motivate this, let's consider the case of an impulse being applied to the system or if the circuit contains diodes.

Periodic mapping using differential equations is a powerful method for solving transients and steady states from solving piecewise linear D.E.s. This is a whole ton of work so a third method was developed; small-ripple approximation. It only works for steady state² as it is an approximate solution to the differential equations based on periodic mapping

SUBSECTION 3.6

Lecture 9

A periodic system is in steady state if

² Well, we only use it for steady-state

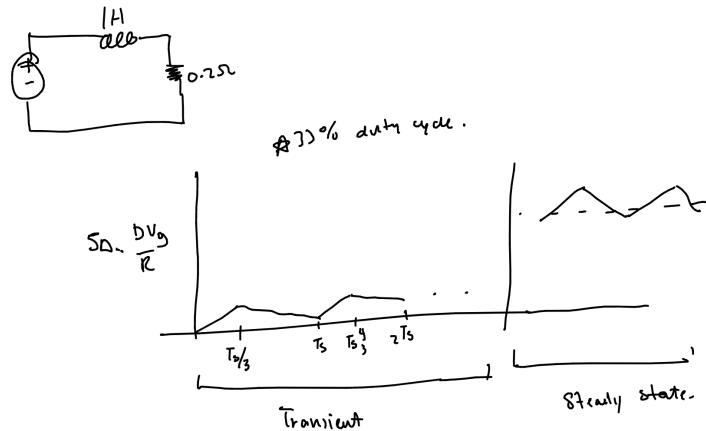
1. all input sources are periodic

2. All states are periodic

3. switching events are periodic

Consider a simple LR circuit with a 33% duty cycle.

Switching events being literally switching parts of the circuit on/off, i.e. with a transistor



The system will have an initial transient response before settling down into a steady state. How can we find this steady state without having to solve for the transients until we get there?

Definition 11

Steady state: the net change of state variables over T_s is zero

An easier way to find the steady state is to use **state space averaging**. Let's look at a RL circuit and its' steady-state response.

In power electronics this is usually termed small ripple assumption

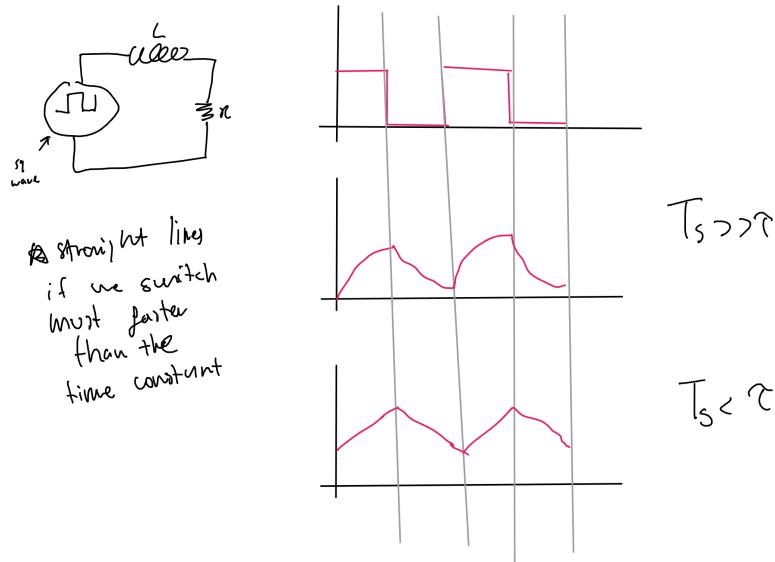


Figure 8. Note how the steady-state response becomes a bunch of straight lines if we switch faster than the time constant τ

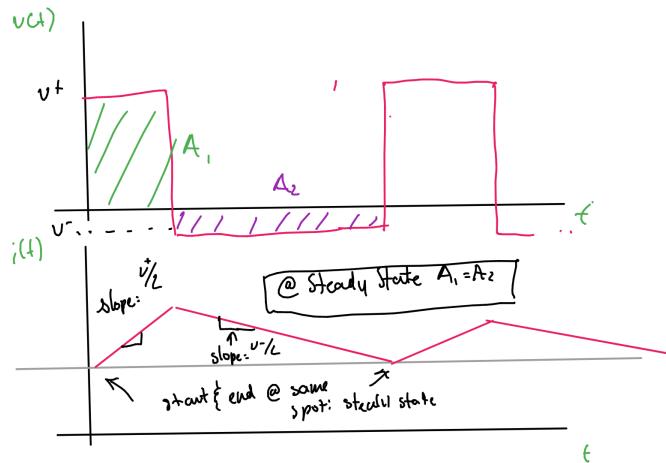
3.6.1 Inductor VoH-Seconds balance

$$\begin{aligned}
 v_l &= L \frac{di_l}{dt} \\
 i_l(t) &= \frac{1}{L} \int v_l(t) dt \\
 i_l(t + T_s) &= \frac{1}{L} \int_{t_0}^{t_0 + T_s} v_l(t) dt + i_l(t_0) \\
 \text{Note: } 0 &= \frac{1}{L} \int_{t_0}^{t_0 + T_s} v_l(t) dt
 \end{aligned} \tag{3.32}$$

So, in steady state $i_l(t_0 + T_s) = i_l(t_0)$

$$0 = \frac{1}{L} \int_{t_0}^{t_0 + T_s} v_l(t) dt \tag{3.33}$$

This implies that, graphically, the area in the positive part of the $v(t)$ plot is equal to that of the negative part.



An analogous relationship can be found for capacitors.

3.6.2 Capacitor charge balance

$$\begin{aligned}
 i_c(t) &= \frac{dv_c(t)}{dt} \\
 v_c(t_0 + T_s) &= \frac{1}{C} \int_{t_0}^{t_0 + T_s} i_c(t) dt + v_c(t_0)
 \end{aligned} \tag{3.34}$$

Capacitor balance gives current * time which is charge

And at steady state

$$0 = \frac{1}{C} \int_{t_0}^{t_0 + T_s} i_c(t) dt \tag{3.35}$$

The relationship between the areas above/after the current curve and the voltage response returning to the same spots is the same as that for inductors, just flipped (i.e. voltage areas are equal for inductors but current areas are the same for capacitors).

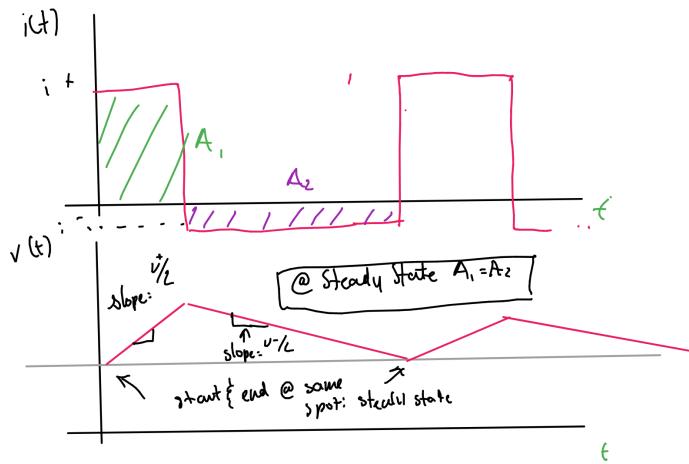


Figure 9. This $A_1 = A_2$ is the capacitor charge balance

As an aside it makes our lives easier to define the average current at steady state I_R so that instead of doing all the math relative to 0 we can take it to the center of the signal. What we really want to find is the average v_c and i_l

A_1 is adding charge to the capacitor, A_2 is removing charge

SUBSECTION 3.7

Lecture 10

1. Assume all switches are ideal
 2. Apply small ripple assumption
- $$\begin{aligned} i_l(t) &\approx I_l \\ v_l(t) &\approx V_c \end{aligned} \quad (3.36)$$
3. Sketch waveforms for $v_l(t)$ for inductors and $i_c(t)$ for capacitors. These should generally be square in wave shape.
 4. Apply V-s balance and charge balance
 5. Validate ripple sizes and small current assumption

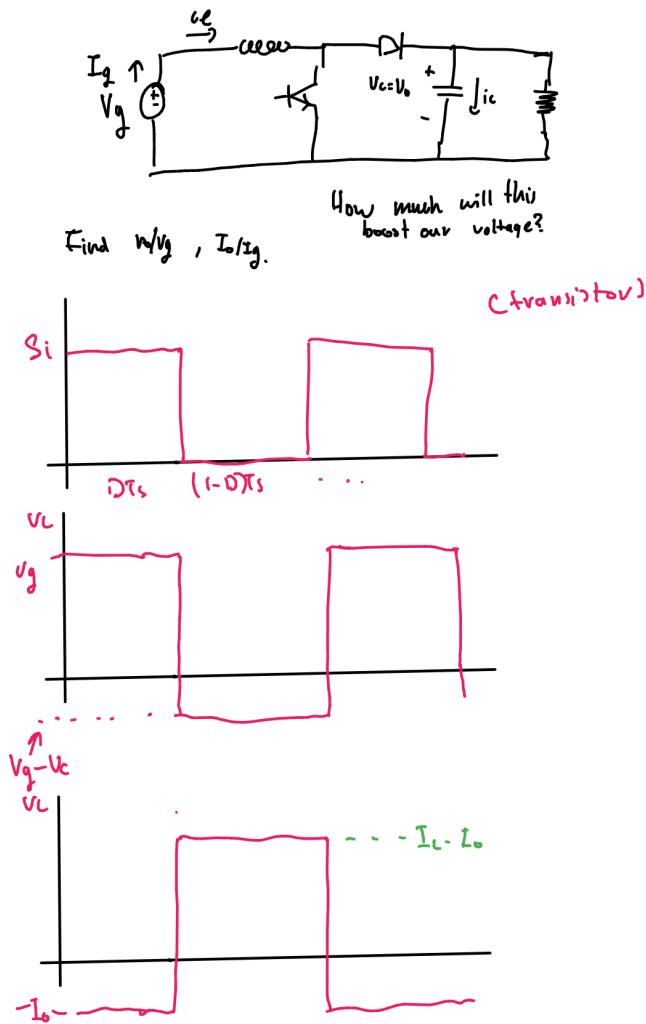


Figure 10. Example with a boost converter. Note: 3rd plot should be i_c vs t

Applying the $V - s$ balance:

$$\begin{aligned}
 \int_0^{T_s} v(t) dt &= 0 \quad \text{or: } A_1 = A_2 \\
 V_g DT_s + (V_g - V_c)(1 - D)T_s &= 0 \\
 V_g &= V_c(1 - D) \\
 \frac{V_o}{V_g} &= \frac{1}{1 - D}
 \end{aligned} \tag{3.37}$$

What this means is that we can boost v strongly by modifying the duty cycle D . In theory we can do this infinitely far but in practice after a while physical effects tend to make it really difficult to boost further.

Applying a charge balance:

$$\int_0^{T_s} = i_c(t)dt$$

$$-I_o(DT_s) + (I_L - I_o)(1 - D)T_s = 0 \quad (3.38)$$

$$\frac{I_o}{(I_L = I_g)} = 1 - D$$

So the output current can similarly be modulated via the duty cycle.
Let's check whether or not if the small ripple assumption is valid.

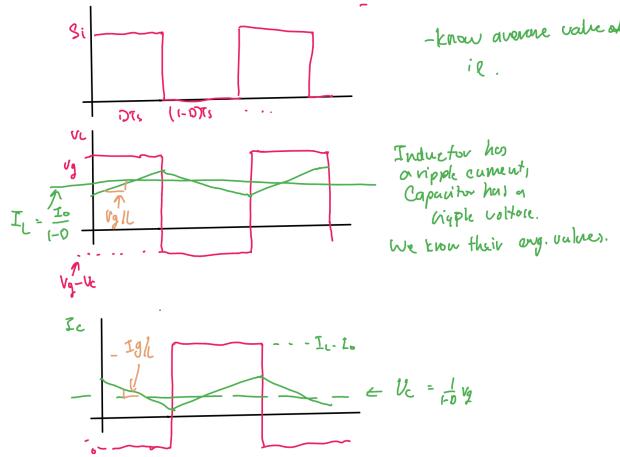


Figure 11. Inductor has a ripple current and the capacitor has a ripple voltage. The slope of the roughly saw-shaped waveform is inversely proportional to L, C . So we can pick L, C such that we get a small enough ripple.

So, what defines a 'small enough ripple'?

The amplitude of the ripple is defined from the peak to peak value of the waveform.³

$$\Delta I_{lp2p} = \frac{1}{L} \int_0^{DT_s} V_g dt = \frac{V_g}{L} DT_s \quad (3.39)$$

$$\Delta V_{op2p} = \left| \frac{1}{C} \int_0^{DT_s} -I_o dt \right| = \frac{I_o}{C} DT_s \quad (3.40)$$

Defining a small ripple is a little more hand-wavy but usually indicates the point at which our approximations start to break down. Generally it means within $< 30\%$ of the rated current and $< 30\%$ of the rate voltage. But this is highly dependent on the application; a computer PSU would want something a lot cleaner for example.

3.7.1 Design Process

1. The application will tell the load; I_o, V_o
2. We will have to pick the switch. It'll have to meet the current spec and maybe some thermal design
3. Look at switch data sheet and then look at the switching frequency that it can run at⁴
4. Compute required duty cycle D for the desired output characteristics V_o/V_g

These integrals can be found by looking at the graphs we drew with the square waves etc and the areas

³ The textbook defines it as center to peak, but prof likes peak2peak

$V_o = V_c$. Generally we don't care about I_{lp2p} ; it is an internal variable. User only cares about V_o , not V_c – they're the same in this case but not necessarily.

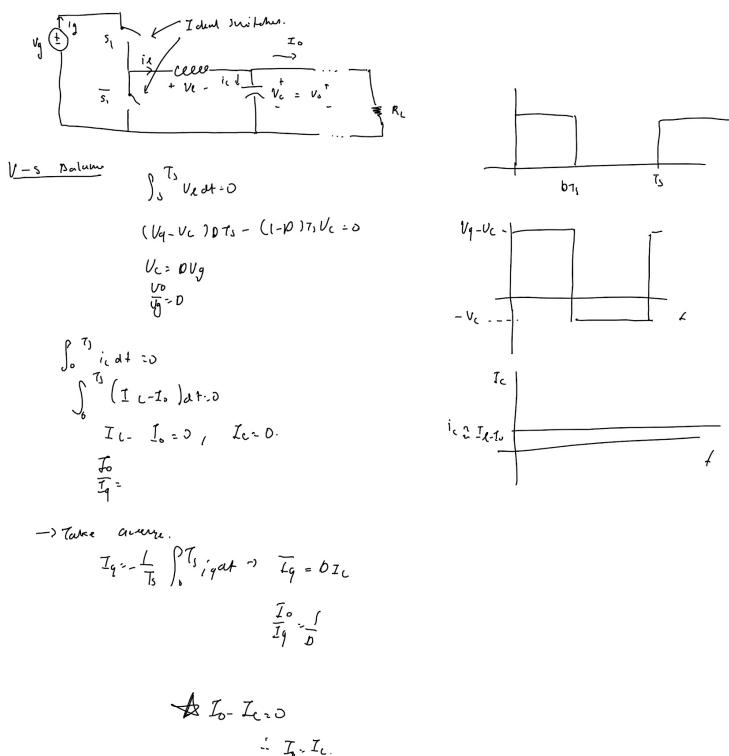
⁴Usually smaller switches can run faster than larger ones

5. Pick inductor current to achieve the small ripple assumption to get reasonable ΔI_{lp2p} . Inductors are more expensive than capacitors so most try to get it to within $< 20\%$ of rated and do the rest using capacitors
6. Pick C to meet ΔV_{cp2p} specification

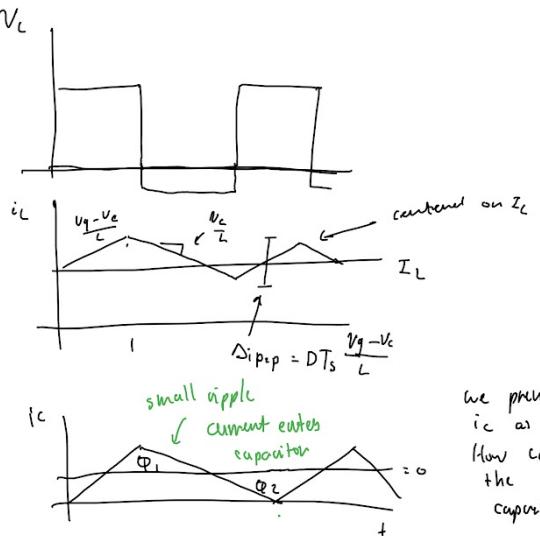
SUBSECTION 3.8

Lecture 11

I slept in and was behind on notes so here are handwritten ones.

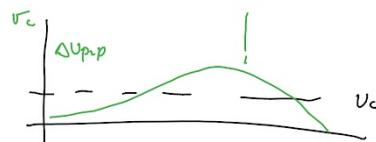


Steady ripple to size L, C



we previously treated
i_C as a straight line.
How can we approximate
the value of the
capacitor current i_C?

The difference between i_L, i_C should
actually cause a small ripple
in real world currents.



$$\Delta i_{p,p} = D T_s \frac{V_q - V_0}{L} \quad \text{Design eq #1}$$

$$\Delta V_{p,p} = \frac{Q_1}{C} =$$

$$\begin{aligned} Q_1 &= \frac{1}{2} \frac{T_s}{2} \frac{\Delta i_{p,p}}{2} \\ &= \frac{T_s}{8} D T_s \frac{V_q - V_0}{L} \end{aligned}$$

$$\Delta V_{p,p} = \frac{T_s^2}{8} D \frac{V_q - V_0}{L C} \quad \text{Design eq #2}$$

$$V_0 = DVq \quad \therefore V_q = \frac{1}{D} V_0$$

$$(V_q - V_0) = \left(\frac{1}{D} - 1 \right) V_0$$

$$= \frac{1}{D} (1 - D) V_0$$

Can multiply
in terms
of
others...

$$\frac{\Delta V_{p,p}}{V_0} = \frac{T_s^2}{8} \frac{(1 - D)}{L C} \quad \text{Alternative
design eq}$$

Lecture 12

PART

II

ECE352: Computer Organization**Admin stuff**

Taught by Prof. Andreas Moshovos

Lecture 1

- Lecture recordings on [YouTube](#)
- Online notes: <https://www.eecg.utoronto.ca/~moshovos/ECE352-2022/>. I "stole" a lot of figures from here but I think Prof. Moshovos is cool with it.

(*) Permission is given to reproduce these notes provided that a notice of their origin is clearly given. All rights reserved just in case :)

Figure 12. Excerpt from online course notes

- Course will cover the following:
 - C to assembly
 - How to build a processor that works
 - Intro to processor optimizations
 - Peripherals
 - OS support (Maybe)
 - (Maybe) Arithmetic circuits
 - Use NIOS II and cover a little bit of RISC-V

4.1.1 Mark breakdown

- Labs 15%
- project 5%
- midterm 30%
- Final 50%
- All exams will be open notes/book/whatever except another person/service helping you.

Preliminary**Lecture 2: Using binary quantities to represent other things**

Computers can represent information in bits; 0/1. Though they don't necessarily know or care what bits are, we may assign our own arbitrary meaning to them – usually numbers with the help of positioning; the LSB represents 2^0 and so forth.

C types

- int: 32b (word)
- char: 8b (byte)
- short: 16b (half word)
- long: 32b (word)
- long long: 64b

Signed numbers may be represented in a number of ways.

- Sign bit (make MSB represent positive or negative numbers and then the remaining $n-1$ bits represent the number. Con: hardware impl sucks because requires if/else)
- Two's complement⁵. Pro: only need to implement adders on hardware and then negative numbers will work just like any other except must be interpreted differently. Positive numbers would always start with a 0 and negatives would start with 1. So the range of possible values becomes $-(2^{n-1} - 1), +2^{n-1} - 1$

Adding together binary numbers can also cause overflow; $(A + B) \geq A, (A + B) \geq B$ may not always be true. Also, when we work with these types we always use all the bits. This has implications when working with values of different lengths.

- `char b = -1 (1111 1111)`
- `short int c = -1 (0000 0000 0000 0001)`
- `a = b + c 0000 0001 0000 0000`
- In order to deal with this we must `cast` the `char` to a `short int`. This is done via sign extension which prepends 0s or 1s⁶ to the `char` so that math can be done on it.

Or, just `;;include <stdint.h>...`

⁵ Flip bits, add one. Intuition; in 3 bit system, adding 7 to 1 would result in 8 which would get truncated to 0.

⁶ two's complement

5.1.1 Floating Point Numbers

Whereas fixed point numbers i.e. \$5.25 can be represented just as how an integer would be represented but with the understanding that the user would interpret it as having a decimal point somewhere that indicates the position of 2^0 . This decimal point would be the same for all numbers of that type, i.e. we could have a six bit number that has places $2^{21}2^02^{-1}2^{-2}$. This is common in embedded systems and how it is formatted isn't super clearly standardized.

Lemma 1 | Reference: [What Every Computer Scientist Should Know About Floats](#)

Definition 12

IEEE 754 Floating Point

This is a single precision 32 bit float

`S EEEEEEEE MMMM MMMM MMMM MMMM MMMM MMM` (5.1)

The most significant S bit is the sign bit, bits 30 through 23 E form the exponent which is an unsigned integer, and 22 through 0 form the (M)antissa. The number being represented can be found using the following:

$$(-1)^S \times 2^{(E-127)} \times 1.\text{Mantissa} \quad (5.2)$$

Example

For example, given the following float:

1 10000001 10000000000000000000000000000000

So S = 1, E = 10000001 = 129 and Mantissa = 10000000000000000000000000000000. The number is therefore

$$(-1^1) \times 2^{(129-127)} \times 1.1000000000000000000000000000000 = -6.0 \quad (5.3)$$

IEEE754 also defines 64 bit floating-point numbers. They behave the same except for now having an 11 bit exponent, the bias being 2047⁷, and the mantissa having 52 bits. A few special cases are also available to represent other quantities

- If E=0, M non-zero, value= $(-1)^S \times 2^{-126} \times 0.M$ (denormals)
- If E=0, M zero and S=1, value=-0
- If E=0, M zero and S=0, value=0
- If E=1...1, M non-zero, value=NaN 'not a number'
- If E=1...1, M zero and S=1, value=-infinity
- If E=1...1, M zero and S=0, value=infinity

Floating-point numbers are inherently imprecise. Addition and subtract are inherently lossy; the mantissa window may not be large enough to capture the decimal points. Multiplication and division just creates a ton of numbers.

Converting real numbers to IEEE754 floats, here using 37.64 as an example, can be done as follows

- Repeatedly divide the part of the number > 0 by 2 and get the remainders, i.e. $37/2 = 18$, rem = 1 $\rightarrow 18/2 = 0$, rem = 0 $\rightarrow 4/2 = 2$, rem = 0, $2/2 = 1$, rem = 0, $1/2 = 0$, rem = 1. As a 2 bit number E is 100101. But we need to convert it to IEEE754 format with the exponent; E - 127 = 5, E = 132 = 1000 0100.
- Do the same for the part of the number past the decimal, but multiplying by two and checking if > 1 : $0.64 * 2 = 1.28 \rightarrow 1$, $0.28 * 2 = 0.56 \rightarrow 0$, $0.56 * 2 = 1.12 \rightarrow 1$... and so forth. At some point we will hit a cycle but we'll just take the N_{mantissa} of digits.

So the full number is 01000010000101101000111101011111

`float` is a 32 bit float and `double` is 64

There are more floating point formats introduced by nvidia and google such as a half-precision or 8-bit float designed to reduce memory use for machine learning

SECTION 6

Memory and NIOS II

SUBSECTION 6.1

Lecture 3: Behavioural Model of Memory

Computers can be described as a set of units, each of which interact with each other and the outside world in a specified way. For example, modern computers tend to have memory units, processing units, display units, and so forth. Each unit has a set of inputs and outputs, and a set of rules that govern how the unit behaves. This gives the manufacturer flexibility in how they want to implement a unit, as long as the unit behaves as specified. When designing these operational units it is important to strike a balance between functionality and specificity; if

the unit is too specific it will be difficult to implement, but if it is too general it will be difficult to use.

6.1.1 Memory

Memory is a unit that stores information and is usually represented as a vector of elements, usually a byte (8 bites). Each element, or memory location, contains a binary quantity and has an associated *address*. The address is a number that uniquely identifies the location of the element in the vector, and is **permanently fixed** at time of manufacture. Most systems are byte-addressable, meaning that there is an unique address for each byte in the memory. The collection of all addresses is called the **address space** of the memory, which is typically a power of two. Modern systems tend to be 32 or 64 bit, meaning that the address space is 2^{32} or 2^{64} elements long.

For each memory location there are two operations available

- **Read:** Read the value stored at a given address
- **Write:** Write a value to a given address

Typical memory behaviour models define that the order of memory operations matters, i.e.

1. **store** 0x10, 0xf0
2. **store** 0x20, 0xf0

We would see that 0xf0 contains the value 0x20 and not 0x10 due to the sequential execution model. Memory that adheres to the sequential model offers operations that are **atomic**; the operations are performed on its own with no interaction or overlap with anything else.

In this case it is convenient to draw memory as an array where each row comprises four consecutive bytes.

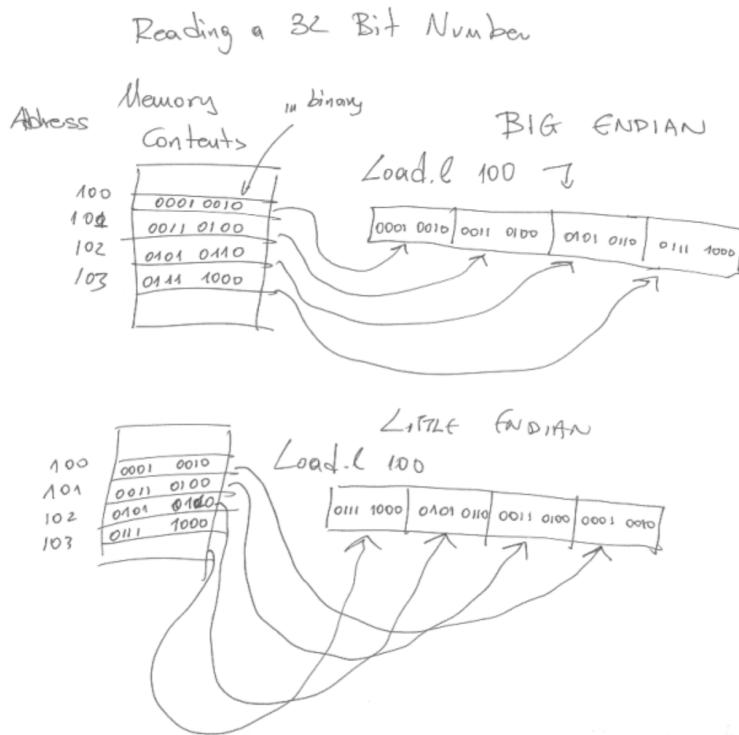
0x00 0000	0x11	0x22	0x33	0x44
0x00 0004	0xff	0x88	0x62	0x51
...				
0xff ffff				

Systems are generally also addressable by words, halfwords, and bytes. Different architectures have different constraints on allowing unaligned access⁸

Endianness refers to the order in which bytes are stored in memory. Though some processors are big-endian, most modern processors are little-endian. The NIOS II used for this course is little-endian.

Specification is the description of what an unit should do, and **implementation** is how it actually does it. For example, an OR gate can be specified as a truth table and then implemented via transistors or a person in a box.

⁸ Aligned access only means to allow [only] reads or writes for a data size i.e. halfword to an address divisible by the size of said data type. For example an longword access on our development board would be at an address divisible by 4



6.1.2 Physical Interface

What physical interfaces would be necessary to implement this behavioural model?

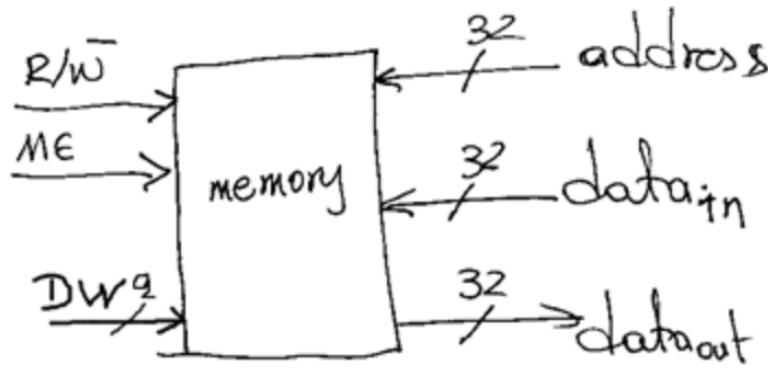
Given a summary of requirements as follows:

1. Read and write operations
2. Addressable by byte, word, longword
3. 24 bit address
4. 32 bit for writing
5. 32 bit for reading
6. signal for do nothing

A single bit signal can be used to indicate whether the memory is reading or writing, and a two bit signal can be used to specify if we're interested in addressing by byte, word, or longword. The address is 24 bits, so we need 24 address lines. As for reading/writing data, we have the option of having two 32 bit data lines, or multiplexing a single 32 bit line. A single bit signal can be used to indicate to do nothing or not.

One way of multiplexing the data lines is to use a tri-state buffer, which is a buffer that can be enabled or disabled. When enabled, the buffer acts as a normal buffer, but when disabled, the output is disconnected from the input. On the other hand this means that our memory chip would not support simultaneous reads or writes.

The use of a single bit signal to indicate 'do nothing' is necessary because a physical device won't be able to change all signals instantaneously, so we use it to tell the memory to wait until these transient effects die off



SUBSECTION 6.2

Lecture 4: NIOS II Programming Model

The NIOS II assumes a 32-bit address space where each address holds a single byte. Each byte is addressable, and three data types are supported. Halfword and word accesses must be aligned.

- **Byte:** 8 bits
- **Halfword:** 16 bits
- **Word:** 32 bits

The NIOS II also has a set of registers

- 32 general purpose 32 bit registers
 - $r0$ is always zero
- 6 control registers, 32 bits each
- Program counter (PC), 32 bits

There are certain conventions for the use of registers, which are as follows:

Many operations can be synthesized using another operation involving zero, i.e. assignment $A=B$ can be implemented as $A = B + 0$

Register	Name	Function	Register	Name	Function
r0	zero	0x00000000	r16		
r1	at	Assembler Temporary	r17		
r2		Return Value	r18		
r3		Return Value	r19		
r4		Register Arguments	r20		
r5		Register Arguments	r21		
r6		Register Arguments	r22		
r7		Register Arguments	r23		
r8		Caller-Saved Register	r24	et	Exception Temporary
r9		Caller-Saved Register	r25	bt	Breakpoint Temporary (1)
r10		Caller-Saved Register	r26	gp	Global Pointer
r11		Caller-Saved Register	r27	sp	Stack Pointer
r12		Caller-Saved Register	r28	fp	Frame Pointer
r13		Caller-Saved Register	r29	ea	Exception Return Address
r14		Caller-Saved Register	r30	ba	Breakpoint Return Address (1)
r15		Caller-Saved Register	r31	ra	Return Address

Notes to Table 3-1:

(1) This register is used exclusively by the JTAG debug module.

6.2.1 Adding Two Numbers

As an exercise, let's see how we can implement the following piece of code in NIOS II assembly

```

1 unsigned int a = 0x00000000;
2 unsigned int b = 0x00000001;
3 unsigned int c = 0x00000002;
4
5 a = b + c;

```

Register-only version

```

1 addi r9, r0, 0x1
2 addi r10, r0, 0x2
3 add r9, r10, r11

```

addi stands for 'add intermediate', the only difference being that the second operand is a number. It is used to set a constant

In general, most instructions take the form of **operation destination, source1, source2**.

Breaking it down even further we can see that these assembly instructions actually perform a number of steps

```

1 addi r9, r0, 0x1
2     ; 1. read r0
3     ; 2. Add value read in step 1 with 0x1
4     ; 3. Write result of step 2 to r9
5     ; 4. increment PC to next instruction
6 addi r10, r0, 0x2
7     ; 1. read r0
8     ; 2. Add value read in step 1 with 0x2
9     ; 3. Write result of step 2 to r10
10    ; 4. increment PC to next instruction
11 add r9, r10, r11
12     ; 1. read r10
13     ; 2. read r11
14     ; 3. Add values read in steps 1 and 2
15     ; 4. Write result of step 3 to r8
16     ; 5. increment PC to next instruction

```

What about 32 bit constants? An unfortunate quirk is that `addi` only supports 16 bit constants, so we need to use `ori` to set the upper 16 bits of the register.

```

1 movhi r9, 0x1122
2     ; Sets the upper 16 bits of r9 to 0x1122
3     ; and the lower 16 bits to zero
4 ori r9, r9, 0x3344
5     ; bitwise OR the value in r9 with 0x3344
6     ; which will set the lower 16 bits to 0x3344

```

This is a PITA so NIOS II offers a few pseudo-instructions to make this easier

```

1 movi rX, Imm16
2     ; sets rX to the sign-extended (signed) 16 bit immediate
3 movui rX, Imm16
4     ; sets rX to a zero-extended unsigned 16 bit immediate
5 movia rX, Imm32
6     ; sets rX to a 32 bit immediate

```

- Footnote1: `movia` does not use the `movhi` and `ori` instructions to create a 32-bit immediate but rather a `movhi` and a `addi`. `addi` will sign extend its 16-bit field so some adjustment might be needed for whatever is being passed to `movhi`.
- Footnote2: `movhi r9, %hi(0x11223344)` is equivalent to `movhi r9, 0x1122`. `Ori r9, %lo(0x11223344)` is equivalent to `ori r9, 0x3344`. That is, `%hi(Imm32)` returns the upper 16-bits of `Imm32` and `%lo(Imm32)` the lower 16 bits.
- Footnote3: `movhi r9, %hiadj(0x11223344)` followed by `addi r1, %lo(0x11223344)` is the correct way of creating a 32-bit immediate using `movhi` and `addi`. `%hiadj(Imm32)` returns the upper 16 bits of the immediate as-is or incremented by 1 if bit 15 is 1. Think why this is necessary based on footnote 1.
- Footnote4: `%hi()`, `%lo()`, and `%hiadj()` are macros supported by the assembler. They are not NIOS II instructions. They get parsed during compile time.

6.2.2 Adding two numbers using memory

NIOS II is a load/store architecture which means that all data manipulation happens only in registers.

```

1 ; read b from memory into r9
2 movhi r11, 0x0020
3 ori r11, r11, 0x0004
4 ldw r9, 0x0(r11)
5
6 ; read c from memory into r10
7 movhi r11, 0x0020
8 ori r11, r11, 0x0008
9 ldw r10, 0x0(r11)
10
11 ; add, then store into r8
12 add r8, r9, r10
13
14 ; store r8 into memory
15 movhi r11, 0x0020
16 ori r11, r11, 0x0000
17 stw r8, 0x0(r11)

```

The new instructions introduced here are

```

1 ldw rX, Imm16(rY) ;; 'load word' from memory
2 ;; rX, rY registers, Imm16 is a 16 bit immediate
3 ;; TLDR; Rx = mem[rY + sign-extended(Imm16)]
4 ; 1. read rY
5 ; 2. sign-extend Imm16 to 32bits
6 ; 3. adds the result of step 1 and 2
7 ; 4. reads from memory a word (32 bit) using the result of step 3 as
   → the address
8 ; 5. write the result of step 4 to rX

```

```

1 stw rX, Imm16(rY) ;; 'store word' to memory
2 ;; rX, rY registers, Imm16 is a 16 bit immediate
3 ;; TLDR; mem[rY + sign-extended(Imm16)] = rX
4 ; 1. read rY
5 ; 2. sign-extend Imm16 to 32bits
6 ; 3. adds the result of step 1 and 2
7 ; 4. write to memory rX using the result of step 3 as the address

```

This can be simplified using the `movia` macro

```

1 movia r11, 0x200004
2 ldw r9, 0x0(r11)
3 movia r11, 0x200008
4 ldw r10, 0x0(r11)
5 add r8, r9, r10
6 movia r11, 0x200000
7 stw r8, 0x0(r11)

```

In this lecture so far we have seen three addressing modes

1. Register addressing, i.e. rX
2. Immediate addressing, i.e. $Imm16$
3. Register indirect addressing with displacement, i.e. $Imm16(rY)$. This is how we calculate the referenced memory address. Register indirect refers to using a register's value to refer to memory, and 'displacement' refers to adding a constant prior to using the register value to access memory. Register indirect addressing is where we use a displacement of 0.

We can exploit register indirect addressing with displacement.

```

1 movhi r11, 0x0020
2 ori r11, r11, 0x0004
3 ldw r9, 0x0(r11)
4 ; can be replaced with
5 movhi r11, 0x0020
6 ldw r9, 0x4(r11)

```

Note that the value of $r11$ does not change since the subsequent operations use an offset to that value.

Generally when we want to read memory from A we can use

```

1 movhi r11, (upper 16 bits of A)
2 ori r9, r11, (lower 16 bits of A)

```

Care must be taken when the 16th bit of A is 1 since the addition that ldw performs will sign extend it to be a negative number, i.e.

```

1 movhi r11, 0x0020
2 ldw r9, 0x8000(r11)
3 ; this is incorrect because
4 ; will extend to 0xFFFF8000, which would result
5 ; in a final address of 0x001F800

```

This is where the macros `%hiadj(Imm32)` and `%lo(Imm32)` come in handy, since they will add 1 to the values if bit 15 of Imm32 is 1. This results in code that looks like this:

```

1 movhi r11, %hiadj(0x208000)
2 ldw r9, %lo(0x2080000)(r11)
3 ; will extend to 0xFFFF8000, which would result
4 ; in a final address of 0x001F800

```

SECTION 7

Assembly Basics

SUBSECTION 7.1

Lecture 5: Simple Control Flow

We have prior worked with straight-line sequences. In this lecture we will look at how to add control flow to our programs, i.e if-then-else, etc.

A pseudo-c program will be rewritten in assembly to illustrate the concepts.

```

1 unsigned int a = 0x00000000;
2 unsigned int b = 0x11223344;
3 unsigned int c = 0x22334455;
4
5 if (b == 0)
6 then a = b + c;
7 else a = b - c;
8

```

The `data` section contains stuff that you want to be initialized for you before the entry point of the program is called, e.g. global variables. This segment as a fixed size. The `text` or `code` segment contains executable instructions (typically read-only, unless the architecture allows self-modifying code) and typically resides in the lower parts of memory. `bss` contains static and global variables which are zero-initialized; usually used for uninitialized data

```

1      .section .data
2  va:   .long 0x0
3  vb:   .long 0x11223344
4  vc:   .long 0x55667788
5
6
7  main:
8      movia r11, va
9      ldw   r9, 4(r11)
10     beq   r9, r0, then
11  else:
12     ldw   r10, 8(r11)
13     sub   r8, r9, r10
14     stw   r8, 0(r11)
15     beq   r0, r0, after
16
17 then:
18     ldw   r10, 8(r11)
19     add   r8, r9, r10
20     stwio r8, 0(r11)
21 after:

```

Definition 13

We encounter two new instructions in this snippet.

- `sub` is a subtraction instruction
- `beq`: a branch-if-equals instruction

The `beq` instruction takes the general form

`beq RX, rY, label.`

This instruction will compare the values of rX and rY and if the condition is true then the program counter will jump to the destination label. When the branch changes the program counter it is called a taken branch, otherwise it is non-taken. Non-taken branches fall through to the next instructions.

Comment

Note: whereas the assembly `beq` command is written as a comparison with a label to jump to, in the NIOSII instruction the destination is encoded relative to the instruction location with a 16-bit displacement constant. The displacement is therefore calculated as $PC + 4 + \text{displacement}$, meaning that we can jump at most $+32774, -32772$ bytes⁹ from the current program counter. The compiler will complain if the branch cannot be implemented.

Encoding is as follows:

AAAAA BBBBB IIIIIIIIIIIIIII 0x26.

where A, B are register names (hence 5 bit), I is a 16 bit immediate value, and $0x26$ is the branch type, in this case `beq`

Most branches tend to be branch backwards because of loops.

⁹ 16 bit signed constant+4, with 4-byte alignment since instructions are 4 bytes long.

Other branch instructions include

- `br`: always/unconditional branch
- `bne`: branch if not equal
- `blt`: branch if less than, w/ signed comparison

- `bltu`: branch if less than, w/ unsigned comparison
- `bgt`: branch if greater than, w/ signed comparison
- `bgtu`: branch if greater than, w/ unsigned comparison

```

1      .data
2          .align 4 ; Align to word size addresses which are faster to
→ access
3  a: .word 0
4  b: .word 0x11223344
5  c: .word 0x55667788
6      .text
7      movia r11, a ; moves the address of a into r11
8      ldw r9, 4(r11) ; loads the value at address a+4 into r9
9      ldw r10, 8(r11)
10     add r8, r9, r10
11     stw r8, 0(r11)

```

SUBSECTION 7.2

Lecture 6, 7: For loops and arrays

How can we implement the following in assembly?

```

1 short arr[5] = { 1, 2, 3, 4, 5 }; // an array of word values (16 bit)
2 short n = 5;                      // the number of elements in the array
3 short sum = 0;
4
5 for (i = 0; i < n; i++)
6     sum = sum + arr[i];

```

Arrays are typically implemented at the machine level as a contiguous block of memory.

```

1      .data
2          .align 1
3  arr .hword 1,2,3,4,5
4  n   .hword 5
5  sum .hword 0

```

Listing 1: Creating a static array in assembly

Multidimensional arrays are handled by having an array of pointers to arrays and dereferencing twice to arrive at the desired memory location. Alternatively we can allocate a $1 \times (m \cdot n)$ chunk of memory and then address it as $a[i][j] = a[i \cdot n + j]$.

As for the for loop, let's look at what C does first.

Generally element $a[i]$ is at address $\&a[0] + \text{sizeof}(\text{TYPE}) * i$

C and pascal and other modern programming languages store arrays in row-major order, i.e. consecutive elements in a row are placed together in memory. FORTRAN uses column-major.

```

1 for (init; cond; post)
2     body

```

Listing 2: General form of a c for loop

Breaking it down a little more into atomic steps

```

1 INIT
2 if (!COND), we are done
3 BODY
4 POST
5 GOTO line 2

```

Listing 3: General form of a c for loop

Let's now rewrite this in assembly

```

1           .text
2 forloop:
3         add r8, r0, r0 ;; INIT
4         movia r9, n ;; set COND, i.e. the n that we compare i to. See:
5         ← loop invariant :)
6         ldh r9, 0(r9)
7         ;; movhi r9, %hiadj(n)
8         ;; ldh r9, %lo(n)(r9) ;; this code block is a shorter piece with the
9         ← same effect
10        loop:
11        bge r8, r9, endloop ;; test condition for when we are done
12        ;; body goes here
13        ;; let's use r10 to store the running sum at in the end we can
14        ← load it to memory
15        movia r11, arr ;; load the address of arr[0] into r11
16        ;; add the index to the address of arr[0]; &arr + i
17        add r11, r11, r8
18        ;; add again; &arr + i + i = &arr + 2i
19        add r11, r11, r8
20        ldhio r12, 0(r11) ;; r12 = arr[i]
21        add r10, r10, r12 ;; r[10] += arr[i]
22        addi r8, r8, 1 ;; POST i.e. r8 += 1
23        br loop
24 endloop:
25         movia r11, sum
26         sth r12, 0(r11) ;; write the sum to memory

```

A similar approach can be taken for while loops.

```

1 .text
2
3 add      r8, r0, r0 ; zero out r8
4 movia    r9, n ; assumes n >= 1
5 ldh      r9, 0(r9) ; loads condition
6
7 dooloop:
8             movia r11, arr ; grab address
9             ← of arr[0]
10            add  r11, r11, r8
11            ← of arr[i]
12            add  r11, r11, r8 ; grab addr
13            ldh  r12, 0(r11) ; load
14            add  r10, r10, r12 ; add
15            ← arr[i] into r12
16            addi r8, r8, 1 ; increment i
17            blt  r8, r9, dooloop ; evaluate loop condition
18 endloop:
19             movia r11, sum
20             sth  r12, 0(r11) ; write the sum into memory

```

Comment

gcc is not a compiler by itself: it actually orchestrates and calls out a lot of other things

1. `cpp`: the C preprocessor, $f.c \rightarrow f.i$; dealing with `;; defines`, etc. Inspect via `gcc -E`
2. `gcc -s f.i`: the `cc`, the C compiler $f.i \rightarrow f.s$: parse pure C from preprocessor and spew out assembly code
3. `as`: the assembler, $f.s \rightarrow f.o$: to turn assembly to machine code, creating objects ready for linkage
4. `ld`: the linker, $f.o \rightarrow f$: link together all the objects into a single executable. Looks through `LDPATH` to look for symbols to link. ¹⁰

¹⁰There is static and dynamic linking; static meaning that everything is inside. These days most are dynamically linked which can grab symbols from shared libraries (.so) right before/during runtime

SUBSECTION 7.3

Lecture 8: Subroutines

A subroutine (or how we more commonly understand them, function calls), is a way to break up a program into smaller pieces and is a core part of structured programming.

Let's look at how we can implement the following in assembly

```

1 int add3(int a, int b, int c){
2     return a+b+c;
3 }
```

First, for C subroutines to work as intended they must:

1. Be callable from anywhere in the program

2. Be able to pass unique parameters across different subroutine invocations
3. Be able to return a value to the caller
4. Must be able to change control flow such that it goes back to the point where it was called

This leads to a number of questions:

1. How does the subroutine return to the caller?
2. How does it return a value?
3. How do we pass arguments?
4. What about local memory?
5. What happens to registers when we call a subroutine?

These issues are addressed through a set of implementation-specific (though largely universal) rules that all valid subroutines must follow. These are called **calling conventions**.

Definition 14

A key concept in subroutine calling is the **stack frame**. As for how stacks work/are implemented that is a google search away.

A stack has the following operations defined:

1. push: put a value on the top of the stack
2. pop: remove the top value from the stack
3. peek (distance): look at the value at a certain distance from the top of the stack

The NIOS II stores its stack pointer, or the address to the top of the stack, in `r27`, i.e. the stack pointer `sp`. The `sp` starts at the top of the stack (highest address value) and then the stack grows downwards to lower addresses.¹¹ There is no way to represent an empty stack in NIOS II; `r27` will always contain a valid pointer address.¹²

Definition 15

push

```

1      ;; push r9 -> stack
2      subi sp, sp, 4 ;; grow stack by a long word
3      ;; (recall: grow downwards)
4      stw r9, 0(sp) ;; save value of r9 to sp

```

pop

```

1      ;; pop top of stack and return to r9
2      ldw r9, 0(sp) ;; read top value
3      addi sp, sp, 4 ;; increment sp; remove top elem

```

top

¹¹ Conversely the **heap** (commonly used for dynamic memory allocation) grows upwards towards higher addresses. The stack and heap are commonly implemented such that they grow towards each other

¹² An element is in the stack if its address is greater than `sp`

```

1      // access ith elem, assuming index in r9
2      add r9, r9, r9 ;; assume r9 contains nedex
3      add r9, r9, r9 // r9 = 4 * i
4      add r9, sp, r9 ;; sp + 4*i
5      ldw r9, 0(r9) // return value of ith element into r9

```

Calling a subroutine and having it return to the caller

The stack is used to implement this behaviour. If a function will be calling another it has to save the `ra` value on the stack in the beginning and restore it from the stack prior to returning. Consider the following pseudo C code

```

1 boo_calls(){
2     coo();
3     doo();
4     return;
5 }
6
7 coo(){
8     doo();
9     return;
10 }
11 doo(){
12     return;
13 }

```

The NIOS II code would be as follows

`r31` is commonly used as the return address pointer and is aliased as `ra`

Note that `ret` transfers execution to the address in `ra`

```

1      .text
2  boo:  ;; boo will be making calls, so it first pushes the ra value on
3      ↳ the stack
4      subi sp, sp, 4
5      stw ra,0(sp)  ;; push the return address onto the stack
6
7
8      call coo  ;; resume execution at coo, ra = PC + 4 = boo_ret1
9  boo_ret1:
10     call doo  ;; continue execution at doo, ra = PC + 4 = boo_ret2
11
12 boo_ret2:
13     ldw ra, 0(sp)  ;; pop return address from the stack
14     addi sp, sp, 4
15     ret  ;; resume execution at ra, which would be the return
16     ↳ address from the
17     ;; stack we just popped
18 coo:
19     subi sp, sp, 4
20     stw ra,0(sp)  ;; push the return address onto the stack
21     call doo  ;; resume execution at coo, ra = PC + 4 = coo_ret
22
23 coo_ret:
24     ldw ra, 0(sp)  ;; pop return address of boo from the stack
25     addi sp, sp, 4
26     ret  ;; resume execution there
27
28 doo:  ;; doo will not be making any calls, no need to save ra on the
      ↳ stack
      ret  ;; just return to whoever called

```

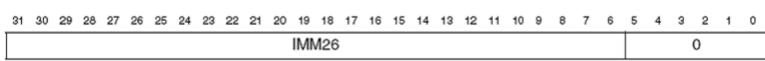


Figure 13. The call instruction accepts a label as the 2nd argument which is encoded as a 26 bit immediate. This limits the called function to be within 256 MB of the caller; the actual target is the immediate multiplied by 4 and then concatenating with the program counter, i.e. $PC_{31..28} = IMM26 * 4$

Note: we can look at executables with `objdump -d`.

Example:

```

1
2 #include <stdio.h>
3 int main () {
4     printf("hello world");
5 }
6

```

```

1  hello_world:      file format elf64-x86-64
2
3
4
5  Disassembly of section .init:
6
7  0000000000001000 <_init>:
8  1000: f3 0f 1e fa      endbr64
9  1004: 48 83 ec 08      sub    $0x8,%rsp
10 1008: 48 8b 05 c1 2f 00 00  mov    0x2fc1(%rip),%rax
    → # 3fd0 <__gmon_start__@Base>
11 100f: 48 85 c0      test   %rax,%rax
12 1012: 74 02      je    1016 <_init+0x16>
13 1014: ff d0      call   *%rax
14 1016: 48 83 c4 08      add    $0x8,%rsp
15 101a: c3      ret
16
17  Disassembly of section .plt:
18
19 0000000000001020 <printf@plt-0x10>:
20 1020: ff 35 ca 2f 00 00      push   0x2fca(%rip)
    → # 3ff0 <_GLOBAL_OFFSET_TABLE_+0x8>
21 1026: ff 25 cc 2f 00 00      jmp    *0x2fcc(%rip)
    → # 3ff8 <_GLOBAL_OFFSET_TABLE_+0x10>
22 102c: 0f 1f 40 00      nopl   0x0(%rax)
23
24 0000000000001030 <printf@plt>:
25 1030: ff 25 ca 2f 00 00      jmp    *0x2fca(%rip)
    → # 4000 <printf@GLIBC_2.2.5>
26 1036: 68 00 00 00 00      push   $0x0
27 103b: e9 e0 ff ff ff      jmp    1020 <_init+0x20>
28
29  Disassembly of section .text:
30
31 0000000000001040 <_start>:
32 1040: f3 0f 1e fa      endbr64
33 1044: 31 ed      xor    %ebp,%ebp
34 1046: 49 89 d1      mov    %rdx,%r9
35 1049: 5e      pop    %rsi
36 104a: 48 89 e2      mov    %rsp,%rdx
37 104d: 48 83 e4 f0      and
    → $0xfffffffffffff0,%rsp
38 1051: 50      push   %rax
39 1052: 54      push   %rsp
40 1053: 45 31 c0      xor    %r8d,%r8d
41 1056: 31 c9      xor    %ecx,%ecx
42 1058: 48 8d 3d da 00 00 00  lea    0xda(%rip),%rdi
    → # 1139 <main>
43 105f: ff 15 5b 2f 00 00      call   *0x2f5b(%rip)
    → # 3fc0 <_libc_start_main@GLIBC_2.34>
44 1065: f4      hlt
45 1066: 66 2e 0f 1f 84 00 00      cs    nopw 0x0(%rax,%rax,1)
46 106d: 00 00 00      ; and this goes on for a while loading in the .so &
    → printf

```

And more assembly later ...

```

1 00000000000000001139 <main>:
2 1139: 55
3 113a: 48 89 e5
4 113d: 48 8d 05 c0 0e 00 00
5   # 2004 <_IO_stdin_used+0x4>
6 1144: 48 89 c7
7 1147: b8 00 00 00 00
8 114c: e8 df fe ff ff
9 1151: b8 00 00 00 00
10 1156: 5d
11 1157: c3

push  %rbp
mov   %rsp,%rbp
lea   0xec0(%rip),%rax
mov   %rax,%rdi
mov   $0x0,%eax
call  1030 <printf@plt>
mov   $0x0,%eax
pop   %rbp
ret

12
13 Disassembly of section .fini:
14
15 00000000000000001158 <_fini>:
16 1158: f3 0f 1e fa
17 115c: 48 83 ec 08
18 1160: 48 83 c4 08
19 1164: c3

endbr64
sub   $0x8,%rsp
add   $0x8,%rsp
ret

```

SUBSECTION 7.4

Lecture 9

We've seen how to call and return from subroutines – but what about passing arguments and returning values? For this lecture we'll assume that only words are returned from and passed to subroutines, but other data types and structs can be used as well.

- Return value is passed in `r2`¹³
- First four parameters are passed in `r4`, `r5`, `r6`, `r7`
- Additional parameters are pushed onto the stack *in order*
 - We push the last argument to the stack first and so forth such that the first non-register argument (the 5th one) is the first to be popped off once we enter the function.

Consider a function which takes seven integers and adds them together.

The following calling convention is the one used by gcc for the NIOS II family

¹³only one return value can be given; multiple can be encoded via structures or ptrs etc

```

1      .data
2  sum: .word 0
3      .text
4  main:
5
6  addi sp, sp, -4
7  stw ra, 0(sp);
8  ;; return address ptr (ra) is pushed onto the stack
9
10;; fill first 4 args
11 movi r4, 1
12 movi r5, 2
13 movi r6, 3
14 movi r7, 4
15
16;; allocated for arguments 5-7
17 addi sp, sp, -12 ;; 3 words * 4 byte
18
19
20;; push arguments 5-7 onto the stack
21;; note order; <top> 5, 6, 7
22 movi r2, 7
23 stw r2, 8(sp)
24
25 movi r2, 6
26 stw r2, 4(sp)
27
28 movi r3, 5
29 stw r2, 0(sp)
30
31 call add7
32
33 add7:
34     add r2, r4, r5      # add the first two arguments and place the
  ↳ sum into r2
35     add r2, r2, r6      # add the third argument to r2
36     add r2, r2, r7      # add the fourth argument to r2
37     ldw r7, 0(sp)       # read the fifth argument from the stack
38     add r2, r2, r7      # add to r2
39     ldw r7, 4(sp)       # read the sixth argument from the stack
40     add r2, r2, r7      # add to r2
41     ldw r7, 8(sp)       # read the seventh argument from the stack
42     add r2, r2, r7      # add to r2 (return value in r2 by
  ↳ convention)
  ↳ ret
43
44;; and then do the cleanup etc

```

Note that in reality gcc will actually preallocate all the memory needed by the function before the function is called. So instead of allocating 12 bytes (3x 1 word arguments) as it does on line 17, it will actually allocate 16 bytes because we need to store the return address.

The allocated stack frame for the function will be the largest of any function being called from it.

For example,

```

1 int main(){
2     foo(1,2,3);
3     boo(1,2,3,4,5,6,7,8);
4 }
```

Boo has the maximum number of arguments, so we'll need to allocate $8 - 4 = 4$ words on the stack for the arguments and the return address $- 5 * 4 = 20$ bytes. This results in

```

1 ;; prologue
2 addi sp, sp, -20
3 stw, ra, 16($0) ;; note little-endian and remaining 16 bytes for 4
   ↳ argument words
4
5 ;; epilogue
6 ldw ra, 16(sp) // pop return addr from stack
7 addi sp, sp, 20;
```

SUBSECTION 7.5

Lecture 10: Recursive Subroutines

Consider this code block that computes the Ackerman function

```

1 int Ackerman(unsigned int x, unsigned int y)
2 {
3     if (x==0) return y+1;
4     if (y==0) return Ackerman(x-1, 1);
5     return Ackerman(x-1, Ackerman(x, y-1));
6 }
```

This function is interesting to implement because of its recursive nature.

Breaking it down a little more,

```

1 int Ackerman(unsigned int x, unsigned int y)
2 {
3     if (x==0) return y+1;
4     if (y==0) return Ackerman(x-1, 1);
5     int tmp = Ackerman(x, y-1)
6     return Ackerman(x-1, tmp);
7 }
```

The return address and the value of x must be stored on the stack, so we will need space for 2 words.

```

1      .text
2 Ackerman:
3      addi sp, sp, -8
4      stw ra, 4(sp)
5
6      bne r4, r0, Xnot0;
7 Xis0:
8      addi r2, r5, 1
9      br epilogue
10 Xnot0:
11     bne r5, r0, Ynot0
12 Yis0:
13     ; pass arguments
14     addi r4, r4, -1 ;First one is x-1, i.e. r4-1
15     addi r5, r0, 1 ; second is 1
16     call Ackerman
17     br epilogue
18 Ynot0:
19     stw r4, 0(sp) ; preserve value of x on the stack
20     ; x is already at the right place for fn call
21     addi r5, r5, -1 ; decrement y
22     add r5, r5, 0
23     call Ackerman
24 epilogue:
25     ldw ra, 4(s0)
26     addi sp, sp, 8
27     ret

```

SUBSECTION 7.6

Lecture 11: Structs and recursive structures

Consider a binary tree with a left and right child and a value. We can represent this as a struct in C.

```

1 struct node{
2     int value;
3     struct node *left;
4     struct node *right;
5 };

```

The memory layout of the struct is word-aligned and *exactly* like that of the struct definition. In this case value lives at addr+0, left at addr+4, and right at addr+8.

Now, consider the following recursive function to perform binary search on a BST

When this gets compiled note that in memory the struct will be word-aligned, i.e. on a 4-byte word machine the beginning of the struct will be at an address which is a multiple of the word size

```
1 int findv(int d, struct node) {
2     struct node * m;
3     m = root;
4     if (root == NULL) {
5         return 0;
6     }
7     if (root->val == d) {
8         return 1;
9     }
10    if (d <= root->v) {
11        return findv(d, root->left);
12    }
13    else {
14        return findv(d, root->right);
15    }
16    // Note that this is a tail-recursive function, i.e.
17    // nothing is done after the recursive call
18    // some compilers may unfurl this into a loop
19 }
```

How can we represent this in assembly?

```

1 // int d is in r4, root is in r5 (by convention)
2     .text
3
4 findv:
5     addi sp, sp, -4
6     stw R4, 0(sp);
7 isrootnull:
8     bne r5, r0, isdv
9 rootisnull:
10    movi r2, 0
11    br epilogue
12 isdv:
13    ldw r2, 0(r5) ;; r2 <- root->val
14    bne r2, r4, tryagain
15 found:
16    mov r2, 1
17    br epilogue
18 tryagain:
19    ble r4, r2, goleft
20 goleft:
21    ldw r5, 4(r5) ;; offset struct addr pointer to right member
22    call findv
23    br end
24 goright:
25    ldw r5, 8(r5) ;; offset struct addr pointer to right member
26    call findv
27    br end
28 notfound:
29    add r2, r0, r0
30 epilogue:
31    ldw ra, 0(sp)
32    addi sp, sp, +4
33    ret

```

Another example of a recursive function is the following routine which detects whether a string is a palindrome.

Note that strings in C are represented as null-terminated array of characters

```

1 char s[] = "abba";

```

Which in assembly looks like:

.datas : .byte 'a', 'b', 'c', 'a', 0;; Or, can write with double quotes i.e. string literals : .string "abba"
(7.1)

The C implementation of this function is as follows:

```
1 int palindrome(char *a)
2 {
3     char * e;
4     if (a == NULL) return 1; // empty string is a palindrome
5     e = a;
6     while (*e != 0) e++; // find end of string
7     if (e == a) return 1; // string of length 1 is a palindrome
8     e--;
9     while ((*a != 0) && (*a == *e)) {
10         a++;
11         e--;
12     }
13     if (*a == 0) return 1; // is a palindrome
14     return 0;
15 }
16 }
```

In assembly this would be as follows:

```

1      .text
2  palindrome:
3      beq r4, r0, ret1 ;; null string is a palindrome
4      add r8, r4, r0 ;; a in r4 (function call convention), e in r8
5
6  findzero:
7      ldb r9, 0(r8) ;; r9 <- *e
8      beq r9, r0, foundzero
9      addi r8, r8, 1
10     br findzero
11
12 foundzero:
13     beq r8, r4, ret1 ;; string of length 1 is a palindrome
14     subi r8, r8, 1 ;; move back to last nonzero char
15
16 cmploop:
17     ldb r10, 0(r4) ;; r9 <- *a
18     beq r10, r0, after ;; end loop when a gets to the end of the
→   string
19     ldb r9, 0(r8) ;; r9 <- *e
20     bne r9, r10, after ;; if *a != *e, then not a palindrome
21     addi r4, r4, 1 ;; increment a
22     subi r8, r8, 1 ;; decrement e
23     br cmploop
24
25 after:
26     ldb r9, 0(r4) ;; r9 <- *a
27     beq r9, r0, ret1 ;; if *a == 0, then is a palindrome
28
29 ret0: ;; not palindrome
30
31 ret1: ;; is palindrome
32     addi r2, r0, 1
33

```

As an example this palindrome function can be called as follows

```

1      .data
2  s: .string "abba"
3
4      .text
5      .global main
6
7 main:
8      movia r4, s
9      call palindrome

```

SUBSECTION 8.1

Lecture 12: Devices

A computer is great and all, but for it to be useful it must communicate with the outside world via devices. A simple device interface is the Parallel Port Interface, a common implementation of which is the GPIO (General Purpose Input/Output) interface¹⁴

There are two registers used to communicate with the NIOSII GPIO interface which live in memory;

- `dr` (data register) at `0xFF200060`
- `dir` (data register) at `0xFF200064`

`dr` is the data word to communicate in and out of the GPIO interface, and `dir` denotes if the signal is an input or an output. The PIT provides 32 bits of data which can be independently mapped to be input or output depending on the value in DIR. If bit 0 of DIR is 1, then D0 becomes an output. And so forth for the rest of the bits. Note that the `stwio` and `ldwio` variants of the `stw` and `ldw` instructions must be used when interfacing with IO devices

¹⁴ Of which the NIOS II board has two

The NIOSII has two GPIO interfaces. The other, GPIO2, is at `0xFF200070`

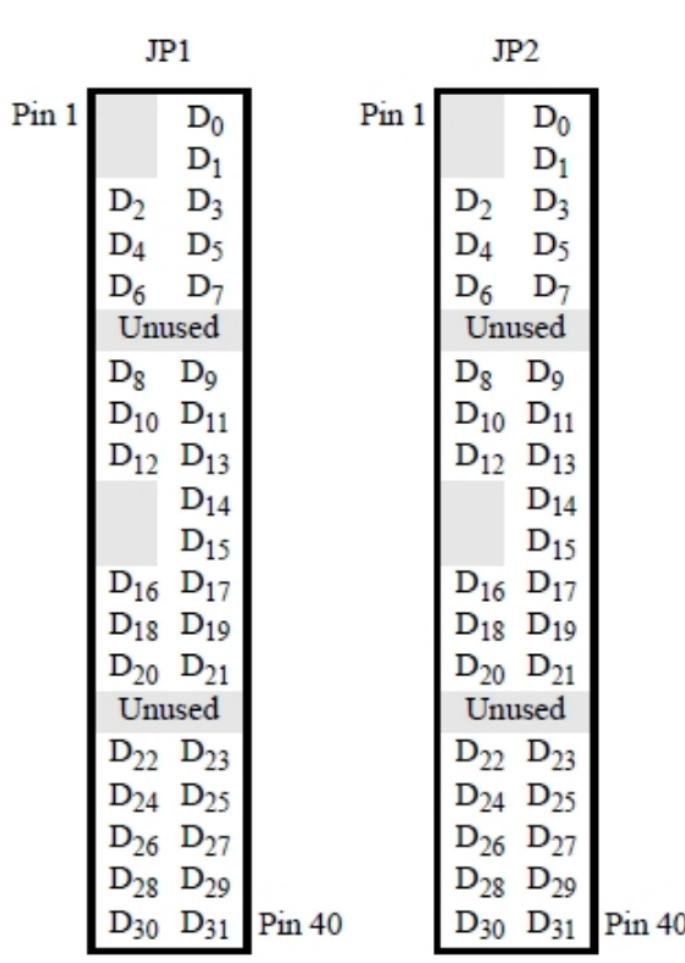


Figure 14. GPIO pinout

Two exercises we will cover include building a thermostat with the GPIO pins and building a keyboard.

```

1  ;; configure all pins as outputs
2  addi r8, r0, 0xFFFFFFFF
3  movia r9, 0xFF200064
4  stwio r8, 0(r9)
5
6  ;; configure pins D0-D3 as outputs, D4-D31 as inputs
7  addi r8, r0, 0x00000000F
8  movia r9, 0xFF200064
9  stwio r8, 0(r9)

```

Given that D_0 is connected to the thermostat's output and D_3 is connected to the heat fan motor,

A write to an input pin will be silently ignored.

```

1  PIT1_BASE equ, 0xFF200060
2  PIT1_DR_OFFSET equ, 0x0
3  PIT1_DIR_OFFSET equ, 0x4
4
5      .text
6  heat:
7      movia r9, PIT1_BASE
8      addi r8, r0, 8
9      stwio PIT1_DIR_OFFSET(r9) ;; configure all pins but d3 as
→   inputs (0x00000008)
10     stwio r0, PIT1_DR_OFFSET(r9) ;; turn off the motor
11  fever:
12      ldwio r8, PIT1_DR_OFFSET(r9) ;; read all pins
13      andi r8, r8, 0x1 ;; mask off all but d0
14      beq r8, r0, fanon ;; turn fan on because it's cold (d0 == 9)
15  fanoff:
16      stwio r0, PIT1_DR_OFFSET(r9) ;; turn off the motor
17      br fever
18  fanon:
19      addi r8, 8
20      stwio r8, PIT1_DR_OFFSET(r9) ;; turn on the motor
21      br fever

```

Something similar can be implemented in c

In the lecture notes he uses 0x00000004 but I think that's a typo.

```

1 #define DR ((unsigned int*) 0xFF200060)
2 #define DIR ((unsigned int*) 0xFF200064)
3
4
5 void heat(void)
6 {
7     unsigned int t;
8     *DIR = 0x8;
9     *DR = 0x0;
10    while (1)
11    {
12        t = *DR;
13        if (t & 0x1)
14            *DR = 0x0;
15        else
16            *DR = 0x8;
17    }
18 }
```

A nicer way that this can be implemented is to use a struct to represent the PIT

```

1 struct PIT_t {
2     unsigned int DR;
3     unsigned int DIR;
4 };
5
6
7 struct PIT_t *pitp = (struct PIT_t *) 0xFF200060;
8
9 void heat(void)
10 {
11     unsigned int t;
12     pitp->DIR = 0x8;
13     pitp->DR = 0x0;
14     while (1)
15     {
16         t = pitp->DR;
17         if (t & 0x1)
18             pitp->DR = 0x0;
19         else
20             pitp->DR = 0x8;
21     }
22 }
```

This chunk of code exploits the fact that the DR element appears first and is immediately followed by DIR. This allows us to elegantly map the fields of the struct to the actual memory layout (line 7)

Another example let's consider debouncing the metal switch used inside the thermostat. Whenever metal contacts are used there is a short period of time where the switch will rapidly change state. This may not be what the user intends and as such we'd like to debounce the values to make it more sane. To avoid this we take a number of samples and assume that the value read is 1 only if the number of samples that was 1 is greater than $\frac{N}{2}$

```

1  PIT1_BASE equ, 0xFF200060
2  PIT1_DR_OFFSET equ, 0x0
3  PIT1_DIR_OFFSET equ, 0x4
4
5      .text
6  heat:
7      movia r9, PIT1_BASE
8      addi r8, r0, 8
9      stwio PIT1_DIR_OFFSET(r9) ;; configure all pins but d3 as
→   inputs (0x00000008)
10     stwio r0, PIT1_DR_OFFSET(r9) ;; turn off the motor
11
12 fever:
13     addi r10, r0, 1000 ;; take 1000 samples
14     add r11, r0, r0 ;; count of samples that are 1
15
16 sampling:
17     ldwio r8, PIT1_DR_OFFSET(r9) ;; read all pins
18     andi r8, 0x1 ;; mask out r0 (thermostat)
19     add r11, r11, r8 ;; increment count if r8 is 1
20     subi r10, r10, 1 ;; decrement sample count
21     bne r10, r0, sampling ;; loop if we still have samples to take
22
23     addi r10, r10, 500 ;; check if we have more than 500 samples
→   that are 1
24     bgt r11, r10, fanon
25
26 fanoff:
27     stwio r0, PIT1_DR_OFFSET(r9) ;; turn off the motor
28     br fever
29
30 fanon:
31     addi r8, 8
32     stwio r8, PIT1_DR_OFFSET(r9) ;; turn on the motor
33     br fever

```

8.1.1 The PIT implementation

This example is a PIT with 8 connections with *DR* at 0xFF1110 and *DIR* at 0xFF1114. A 32 bit PIT should follow.

At a physical level the PIT presents the following signals

1. Data: 8 wires that supply the value to be written to or read from the PIT
2. DIR write signal: when this is 1 the value in the data wires is written to the DIR
3. DR write signal: when this is 1 the value in the data wires is written to the DR
4. DR read signal: when this is 1 the data signals are placed into output mode and take on values that currently appear on external port connections
5. External port: 8 wires that connect to the external world

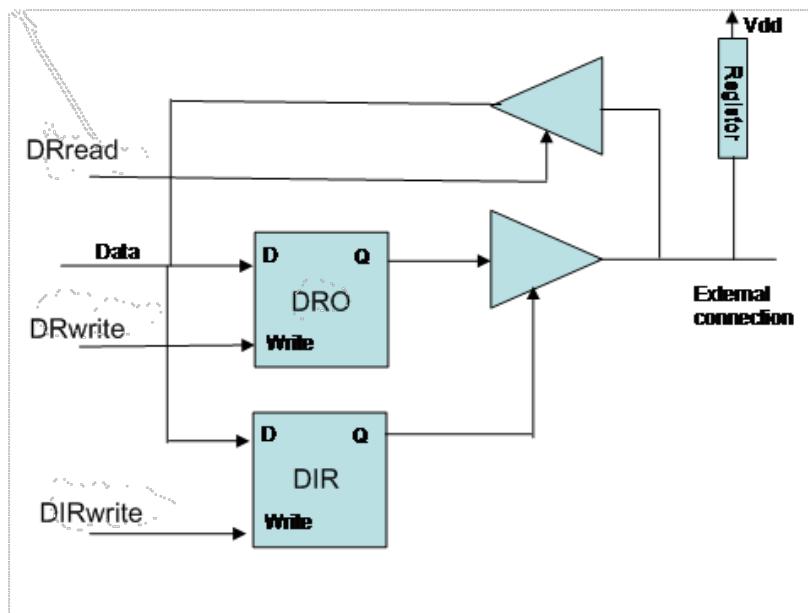


Figure 15. One PIT bit

The DRO and DIR boxes are latches whose values change using the Data and Write signals. The triangle looking things are **tri-state buffers**.

- DIR = 1 -> tri-state latch is closed -> passes value of DRO to external connection
- DIR = 0 -> tri-state latch is open -> external source determines value on the connection
 - resistor helps drive high voltage from external source
- DRread = 1 -> opens upper tristate buffer, puts external value into Data line

Eight of these can then be chained together to form a 8 bit PIT.

8.1.2 Memory

Our memory needs to meet these requirements

1. Load/store operations
2. Byte, halfword, or word types
3. Addr of 32 bits
4. data value of at most 32 bits for read/write values
5. A do-nothing signal

1) can be implemented with a single read-write signal. 2) can be done with 2 signals. Let's use 00 for byte, 01 for half-word, and 11 for long-word. 10 is not used. 3) and 4) need 32 signals each, and 5) requires a master-enable signal. This is useful to prevent the memory from taking on transient values.

Now comes the task of connecting the PIT to the memory interface. First we connect the PIT data lines to the lower 8 bits of the memory interface. So DIR, DRO, DRI can be accessed via load/store operations

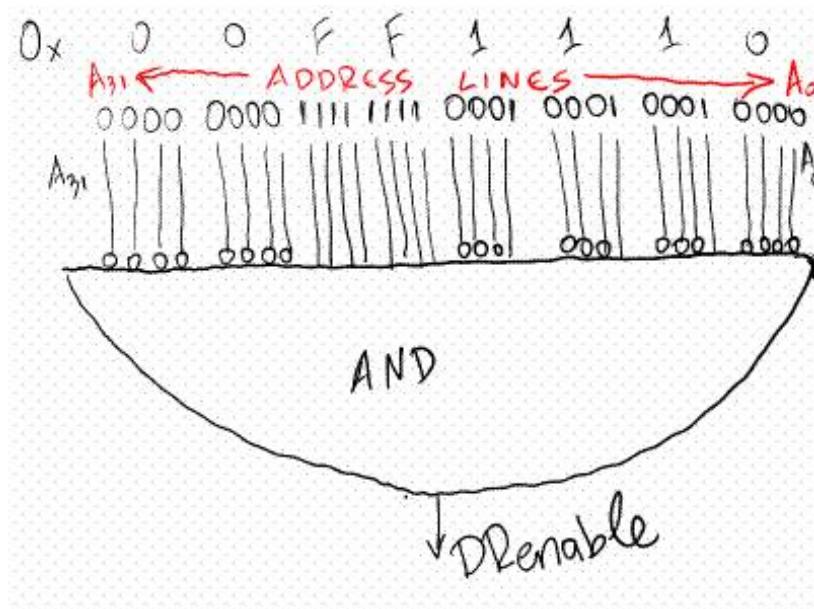
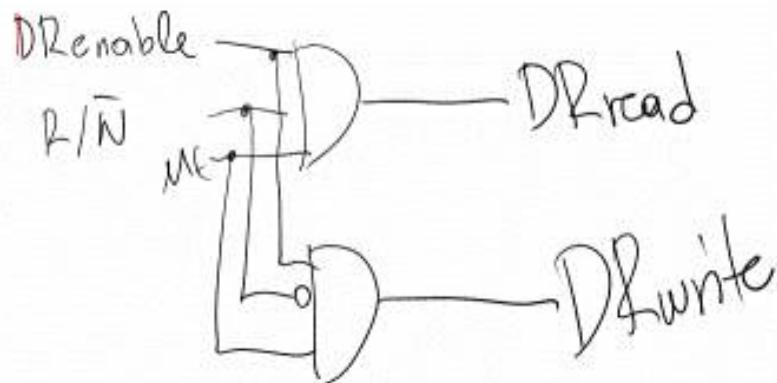


Figure 16. We'll also need to check if the address being accessed is the address of the PIT. The easiest way to do this is to just use an AND across the address lines and pass that signal to DREnable

DREnable can now be combined with R/W to form a single signal that generates the DR-read and DRwrite signals



Something similar can be done for DIRwrite.

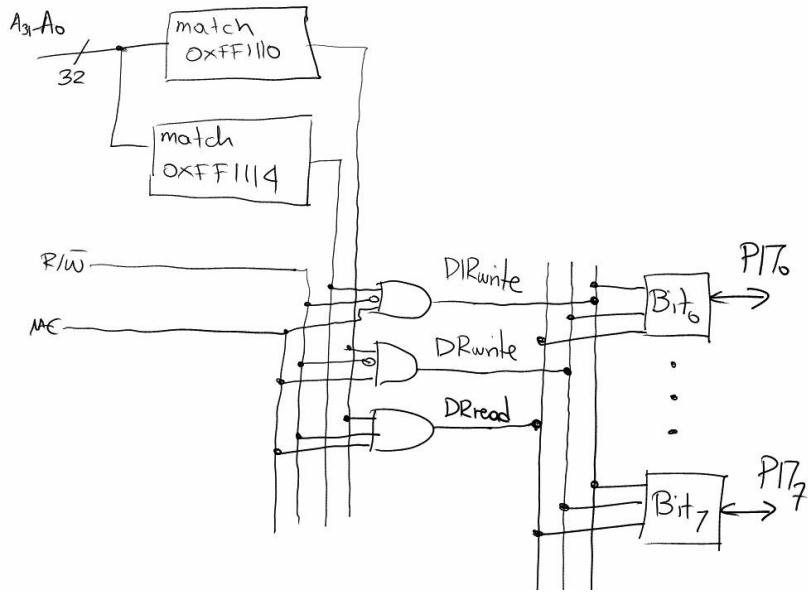


Figure 17. The complete design

SUBSECTION 8.2

Lecture 12: UART

The (JTAG) UART device is a UART implemented over the JTAG interface, which for the NIOSII is implemented over USB. UART stands for Universal Asynchronous Receiver Transmitter. It is a communication protocol which allows for non-instantaneous communication between two devices.

- Starts at address 0xFF201000
- Receiver Register RR at an offset of 0
- Transmitter Register at an offset of 0 (same as RR)
- Control and status register, CSR at an offset of 4

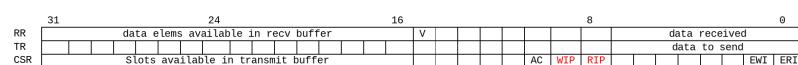


Figure 18. Register format

There are two transactions that can be asked of the UART: (1) send a character and (2) receive a character. They are implemented by placing bits into the data to send/receive into the transmitter/receiver register. However, in real life communication is more difficult than just that; sending characters takes considerably longer than the CPU's speed so we must be able to wait while the UART is sending a character before an attempt to send another one is made. Also, we should only read a character from the RR only if one has been received from the external source. This reading may also need to be buffered due to the variable nature of RR receive speed.

sending a character

Here's a subroutine for sending the character in r4.

```

1 .equ JTAG_UART_BASE, 0xFF201000
2 .equ JTAG_UART_RR, 0
3 .equ JTAG_UART_TR, 0
4 .equ JTAG_UART_CSR, 4
5
6     .text
7 putchar:
8     movia r8, JTAG_UART_BASE
9 wait:
10    ldwio r2, JTAG_UART_CSR(r8) # read csr
11    srlr r2, r2, 16 # keep only upper 16 (shift right logical
→   immediate)
12    ;; The csr's bits 16-31 contains a number which
13    ;; is the number of slots in the queue available
14    ;; to accept requests. If the number is 0, then
15    ;; the queue is full and we must wait.
16    beq r2, r0, wait # wait until upper 16 bits are nonzero
17
18    stwio r4, JTAG_UART_TR(r8) # place it into the FIFO
19    ret
20 wait:
21
22
23

```

receiving a character

To receive a character we must wait until a character is received. This information is in RR

- 0-7 contain the character, if one has been received
- 15 is 1 if a character has been received. If it is 0, then no character has been received and 0-7 are meaningless
- 31-16 contain additional information e.g. number of additional characters received and are waiting in the incoming queue.

```
1 .equ JTAG_UART_BASE, 0xFF201000
2 .equ JTAG_UART_RR, 0
3 .equ JTAG_UART_TR, 0
4 .equ JTAG_UART_CSR, 4
5
6     .text
7 getchar:
8     movia r8, JTAG_UART_BASE
9
10 wait:
11     ldwio r2, JTAG_UART_RR(r8) ;; read RR in r2
12     andi r10, r2, 0x8000 ;; extract bit 15
13     beq r10, r0, wait ;; wait until bit 15 is 1
14     andi r2, r2, 0xff ;; keep only character (bit 0-7) in r2
15     ret
16
17
18
```

Both of these subroutines busy waits are used. This is slow and inefficient. A better way is to use interrupts, which will be covered in a later lecture

```

1  ;; An echo routine
2  .equ JTAG_UART_BASE, 0xFF201000
3  .equ JTAG_UART_RR, 0
4  .equ JTAG_UART_TR, 0
5  .equ JTAG_UART_CSR, 4
6
7  .text
8 echo:
9  movia r8, JTAG_UART_BASE
10 waitr:
11  ldwio r2, JTAG_UART_RR(r8)    # read RR in r2
12  # extract bit 15 in register r10 / keep a copy of r9 since it
13  # contains the character if any
14  andi r9, r2, 0x8000
15  # if bit 15 was zero, there was no character, keep
16  # waiting/trying
17  beq r9, r0, waitr
18  # a character was received, copy the lower 8 bits to r2 and
19  # return
20  andi r2, r2, 0xff
21
22 waitt:
23  # read CSR in r9
24  ldwio r9, JTAG_UART_CSR(r8)
25  # keep only the upper 16 bits
26  srlt r9, r9, 16
27  # as long as the upper 16 bits were zero keep trying
28  beq r9, r0, waitt
29  # place it in the FIFO
30  stwio r2, JTAG_UART_TR(r8)
31  # life is interesting, keep doing what you do
32  br waitr
33  # never reaches here, this is for show
34 ret
35

```

At the communication link level the device represents the characters as a stream of bits. Each bit is communicated by setting the line to the corresponding voltage level for a pre-specified duration; this is the bit-cell shown.

Definition 16

baud rate : number of bit cells within one second. For example 9600 baud means a bit time of $\frac{1}{9600} = 104.16$ microseconds, so it would take 1.0416 milliseconds at least to send a full byte. Note that each byte has a start and stop bit so a byte takes 10 bit cells to transmit.

Ideally the transmitter and receiver would use identical time references for the bit cells, which would make life easy because it would mean the receiver can simply just take a single sample at the center of each bit cell to reconstruct the byte. However, life is not so easy and real-life devices have phase differences.

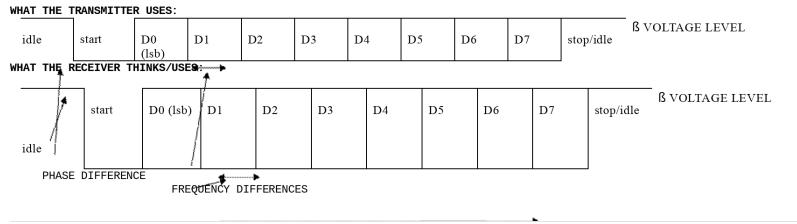
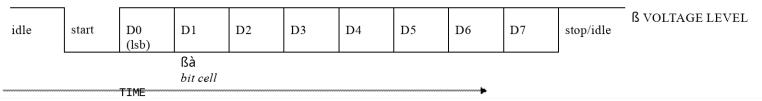


Figure 19. Phase differences

15

The common solution for this problem is to use over-sampling, i.e. taking several samples to detect the 0->1 transition for the stop/idle to start bits, then making carefully chosen single samples to fall at the center of bit time. For a receiver that takes 16 samples per bit cell, generally the i th sample should be taken at $(24 + i * 16)$ cycles, i.e. pass 16 samples to get past start bit and then another 8 to go to the middle of the first bit cell and so forth.

Even with these measures communication errors¹⁶ can still occur. Further reduction of errors can be achieved by using an additional parity bit which can be used to detect single errors.



¹⁵ In practice there shouldn't be a difference of more than 20% between the stop bit and where the receiver thinks the stop bit is. This is because of RS-232C standard imposing a requirement imposing a maximum of 2% difference between device baud rates, and the fact that there are start/stop bits which allow for the receiver to sync to sender time.

¹⁶ FRAME Errors

SUBSECTION 8.3

Lecture 13: Interrupts & UART

In the previous lecture we used busy polling to communicate with the serial device, i.e. probing the status register until the desired condition was met. Not only is this a wasteful use of device resources, it is also difficult to scale to multitasking scenarios.

At a high level, interrupts work as follows:

1. The processor is executing instructions
2. A device requests an interrupt
3. The processor decides to grant that request, i.e. finishing current instruction and then inducing a call to a special interrupt handler routine
4. Program flow is restored to previous instruction

Therefore there are a few requirements that must be met for interrupts to work

1. Devices must be able to request interrupts
2. Must be able to signal to device that its interrupt has been handled
3. Must have a way to map interrupts to the correct interrupt handler
4. Need to know where the interrupt handler lives in memory
5. Must be able to save and restore state affected by an interrupt call

Requirements 1 and 2 are handled by having two directed wires between the CPI and the device through which request and response 'handshake's can be implemented. The NIOS II provides 32 interrupt request wires IRQ0 to IRQ31, and notifying a device is done in software

i.e. writing to a specific location in memory. Exact process is device-specific. Requirements 3 and 4 i.e. inducing an interrupt handler. One way that this can be done is to have a subroutine at a hard-coded address for all interrupts which then queries which device requested the interrupt and then inducing the call to the appropriate handler. This is generally implemented as some sort of lookup table. The NIOS II implements this by having all interrupt requests start executing code at `0x00000020`; the `.section exceptions, "ax"` directive can be used to assign code to that part of memory. That code must interrogate all devices to figure out which devices are requesting interrupts and then fire off the appropriate ISRs. Requirement 5 can be met by just pushing everything relevant onto the stack.

Definition 17

- **ctl0**: Master interrupt enable switch. Bit 0 is 1 means interrupts will be serviced; they will be ignored otherwise.
- **ctl1**: Preserves current state of ctl0 upon accepting an interrupt
- **ctl3**: Enable/disable specific IRQ lines. Interrupts are accepted for IRQ_i if bit i of ctl3 is 1 and bit 0 of ctl0 is 1

Reading and writing from control registers is done through two specialized instructions; `rdctl ctrlx, ry`: reads value of a control register $ctrlx$ and stores in ry , and `wrctl ctrlx, ry` which writes the value of ry into $ctrlx$. Also, $r29$ also participates by storing the PC of the instruction following that of the interrupt and commonly is aliased to be ea , or exception address.

The more detailed sequences is therefore as follows

- Device assert an IRQ line to request an interrupt
- Processor finishes current instruction
- Processor saves the current interrupt enable bit (storing $ctl0$ into $ctl1$), then disables $ctl0$ to disable further interrupts
- Set $ea = PC + 4$; the PC immediately following the instruction that had just been executed
- Set $PC = 0x20$ so that execution continues in interrupt handler
- Handler terminates with `eret` which will restore $ctl0$ from $ctl1$ and PC from ea

SUBSECTION 8.4

Lecture 14: Timer

In lieu of lecture notes here is a big chunk of code which uses the NIOSII timer device to wait a second.

```

1  .equ  TIMERO_BASE,      0xFF202000
2  .equ  TIMERO_STATUS,    0
3  .equ  TIMERO_CONTROL,   4
4  .equ  TIMERO_PERIODL,   8
5  .equ  TIMERO_PERIODH,   12
6  .equ  TIMERO_SNAPL,    16
7  .equ  TIMERO_SNAPH,    20
8  .equ  TICKSPERSEC,     500000000
9    .text
10  waitasec:
11    movia r8, TIMERO_BASE
12    addi  r9, r0, 0x8          ; stop the counter
13    stwio r9, TIMERO_CONTROL(r8)
14
15    ; Set the period registers to 50M
16    addi  r9, r0, %lo(TICKSPERSEC)
17    stwio r9, TIMERO_PERIODL(r8)
18    addi  r9, r0, %hi(TICKSPERSEC)
19    stwio r9, TIMERO_PERIODH(r8)
20
21    ; tell the counter to start over automatically and start counting
22    addi  r9, r0, 0x6          ; 0x6 = 0110 so we write 1
23    ↳ to START and to CONT
24    stwio r9, TIMERO_CONTROL(r8)
25
26  onesec:
27    ldwio    r9, TIMERO_STATUS(r8) ; check if the T0 bit of
28    ↳ the status register is 1
29    andi    r9, r9, 0x1
30    beq     r9, r0, onesec
31
32    addi    r9, r9, 0x0          ; clear the T0 bit
33
34    ; decrement the number of remaining seconds and repeat until
35    ↳ this becomes zero
36    subi    r4, r4, 1
37
38    bne     r4, r0, onesec
39
40    ; stop the counter before exiting
41
42    addi    r9, r9, 8
43    stwio    r9, TIMERO_CONTROL(r8)
44    ret
45
46  ; And here's a function that calls waitasec asking it to wait for 20
47  ↳ seconds:
48    .text
49    .global main
50
51  main:
52    addi  r4, r0, 20
53    call   waitasec
54    ret

```

Checking the current value of the period registers cannot happen directly and instead must be done through snap registers.

```

1  movia r8, TIMER0_BASE
2  stwio r0, TIMER0_SNAPL(r7)  # Take a snapshot of the period
   ↳ registers
3  ldwio r9, TIMER0_SNAPL(r7)  # Read snapshot bits 0..15
4  ldwio r10, TIMER0_SNAPH(r7) # Read snapshot bits 16...31
5   slli r10, r10, 16          # Shift the upper bits to positions 16
   ↳ through 31 in register r10
6   or   r9, r9, r10          # Combine bits 0..15 and 16...31 into
   ↳ one register

```

SUBSECTION 8.5

Lecture 15: Code Races

Consider the following c code:

```

1  int cnt = 9;
2  interrupt:
3      cnt++;
4  program:
5      while (1) cnt--;

```

At the machine code level this can look like

```

1  .data
2  cnt: .word 9
3  .section .exceptions, "ax"
4  iHdlr:
5  H1:      ldw      r9, 0(r8)
6  H2:      addi     r9, r9, 1
7  H3:      stw      r9, 0(r8)
8  eret
9  .text
10 main:
11      movia    r8, cnt
12      movi     r9, 1
13      wrctl   ctl0, r1
14 wait:
15 I1:      ldw      r11, 0(r8)
16 I2:      addi     r11, r11, -1
17 I3:      stw      r11, 0(r8)
18      br       wait

```

We can possibly observe the sequence 10, 11, 10, 9, and so forth because of a *race* between the interrupt and the while loop. This is because NIOS II does not implement atomic operations, therefore the same value can be loaded a register for two operations that change it

without coordinating between the two programs, and therefore cause undesirable behaviour. Other architectures may implement atomic operations, which are operations that cannot be by another instruction (usually at the expense of performance).

8.5.1 Buffer Overflow Stack Attacks

Consider the following `c` method which takes a `char*` input over the network

```

1 void packet_parse_string(char* packet){
2     char buffer[100];
3     // some code
4     // copy packet (null-terminated string) into buffer
5     while (packet[i]){
6         buff[j++] = packet[i++];
7     }
8     // some other code
9 }
```

One may immediately notice that there is no check to see if the packet is larger than 100 bytes, and therefore the buffer may overflow. Though this may seem like an innocuous invalid write, this function can be exploited to run arbitrary code. From this course we know that when returning it will try to restore RA from the stack, perform RET, and start executing code from PC=RA. This means that if the malicious caller can identify where the top of the stack was when calling, they may change the value of RA via the buffer overflow to the address of `buffer` and therefore execute arbitrary code that they sent inside of `packet`.

Ways to defend against such attacks include writing non-sloppy code or forbidding execution of code on the stack¹⁷

SUBSECTION 8.6

Lecture 16: Emulating instructions in software

Interrupts can be used for more than just communicating with I/O devices, i.e. for detecting erroneous conditions during execution, misaligned memory access, or OS calls. They can also be used to emulate an instruction in software.

For example, on the NIOSII there is a `mulxuu rC, rA, rB` instruction to multiply two 32 bit registers and then store the upper 32 bits of the result into a `rC`. Some NIOSII implementations may include a hardware unit which implements this instruction, but others may rely on emulation to execute this instruction instead. Likewise, the 80386 processor does not have a hardware unit to implement floating point instructions (but one could do them in hardware with the 80387 co-processor). If one did not have the 80387 they can use interrupts to emulate the instruction in software instead (though slower)

Instruction Type:	R				
Instruction Fields:	A = Register index of operand rA B = Register index of operand rB C = Register index of operand rC				
31 30 29 28 27 26 25 24 23 22 21 20 19 18 17 16 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0					
A	B	C	0x07	0	0x3a

Figure 20. Encoding of the `mulxuu` instruction

¹⁷ though there are valid reasons for doing this at times, i.e. JIT compilation.

Bits 31-27, 26-22, 21-17 encodes the source and destination registers. 16-11 contains 0x7, 10-6 the value 0, and 5-0 the value 0x31. So TLDR bit 16 is 0 and the lower 16 bits 0x383a.

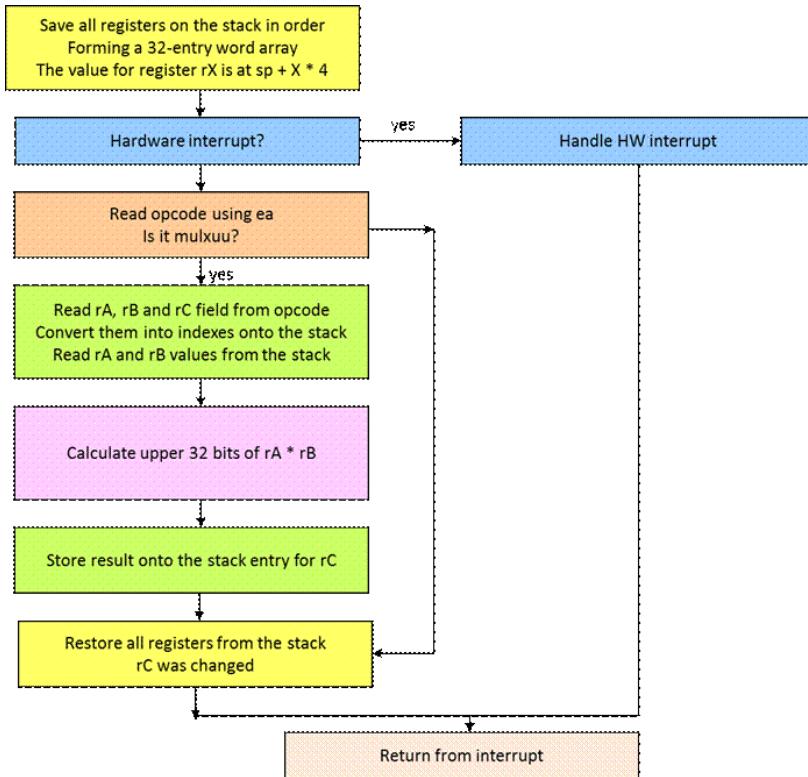


Figure 21. Interrupt handler flowchart

This chunk of code implements the first item in the flow chart, saving all the registers onto the stack. Note the commented `rdctl` checking for software interrupts

```

1  .section exceptions
2
3      ; tell the assembler to not introduce any additional
4      ; instructions overwriting registers
5      .set nobreak
6      .set noat
7
8  handler:
9      ;;;;;;;;;;;;;;;;;
10     ; store all registers on the stack
11     ; forming an array of words
12     ; the value for register X is at sp+X*4 where X a number
13     ; 0...32
14     ;;;;;;;;;;;;;;;;;
15     ; save all registers on the stack
16     subi  sp, sp, 32 * 4
17     stw   r0,0(sp)
18     stw   r1,4(sp)
19     stw   r2,8(sp)
20     stw   r3,12(sp)
21     stw   r4,16(sp)
22     stw   r5,20(sp)
23     stw   r6,24(sp)
24     stw   r7,28(sp)
25     stw   r8,32(sp)
26     stw   r9,36(sp)
27     stw   r10,40(sp)
28     stw   r11,44(sp)
29     stw   r12,48(sp)
30     stw   r13,52(sp)
31     stw   r14,56(sp)
32     stw   r15,60(sp)
33     stw   r16,64(sp)
34     stw   r17,68(sp)
35     stw   r18,72(sp)
36     stw   r19,76(sp)
37     stw   r20,80(sp)
38     stw   r21,84(sp)
39     stw   r22,88(sp)
40     stw   r23,92(sp)
41     stw   r24,96(sp)
42     stw   r25,100(sp)
43     stw   r26,104(sp)
44     stw   r27,108(sp)
45     stw   r28,112(sp)
46     stw   r29,116(sp)
47     stw   r30,120(sp)
48     stw   r31,124(sp)
49     rdctl et, ctl4          ; Check that interrupt was caused
50     by      software
51     beq et, r0, software    ; if not, it's a hardware
52     interrupt ignore
53
54     HANDLE HARDWARE INTERRUPTS HERE
55
56     br    iEpilogue

```

The software interrupt handler checks the opcode in order pass execution to the appropriate subroutine. For `mulxuu` this is just comparing the lower 16 bits of `ea` with `0x383a`

```

1      ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
2      ; read the instruction opcode to test whether it is a mulxuu
3      ; ea points to the instruction
4      ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
5 software:
6      ldw  r9, -4(ea) ;; lecture notes uses stw but this should be
7      ← ldw
8      add  r10, r9, r0 ; keep a copy of the opcode
9      andi r9, r9, 0xffff ; keep just the lower 16 bits
10     cmpeq r11, r9, 0x383a
11     beq  r11, r0, notmulxuu
12     srli r10, r10, 16 ; shift the upper 16 bits into the lower
13     ← 16
14     andi r11, r10, 0x1 ; test bit 0 which used to be bit 17
15     bne  r11, r0, notmulxuu ; if not zero this is not mulxuu
16

```

If we get to the `mulxuu` handler we need to load the registers from the stack

```

1      ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
2      ; Operand index calculations
3      ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
4 ismulxuu:
5      ; now calculate indexes into the stack for accessing
6      ; the input and output operands
7      ; treat the stack as a 32-entry array of words
8      ; we extract the 5 bit field for each operand
9      ; multiply by four because each entry is four bytes
10     ; and add the stack point which is the base of the array
11     ;;;;;;;;;;;;;;;;;;;
12     srli r10, r10, 1 ; keep just the upper 15 bits of the opcode
13     ; rC
14     andi r11, r10, 0x1f ; these are the 5 bits indicating rC the
15     ← destination register
16     slli r11, r11, 2 ; multiply by 4
17     add  r11, r11, sp ; add the base of the array
18     ; rB
19     srli r10, r10, 5
20     andi r12, r10, 0x1f ; keep the bits for rB
21     slli r12, r12, 2 ; multiply by 4
22     add  r12, r12, sp ; add the base of the array
23     ; rA
24     srli r10, r10, 5
25     andi r13, r10, 0x1f ; keep the bits for rA
26     slli r13, r13, 2 ; multiply by 4
27     add  r13, r13, sp ; add the base of the array

```

...and then multiply the two registers and store the upper 32 bits into the destination reg-

ister.

```

1      ;;;;;;;;;;;;;;;
2      ; Access input registers
3      ;;;;;;;;;;;;;;;
4      ; at this point:
5      ; r11 points to the entry for rC
6      ; r12 points to the entry for rB
7      ; r13 points to the entry for rA
8      ; read rA and rB into r9 and r10 respectively
9      ;;;;;;;;;;;;;;;
10     stw r9, 0(r13)
11     stw r10, 0(r12)
12
13
14     ;;;;;;;;;;;;;;;
15     ; Multiplication : No need to understand how this works
16     ; end result is in r10
17     ; I haven't tested it much :(
18     ;;;;;;;;;;;;;;;
19     srli r4, r9, 16 ; a = (v1 >> 16) & 0xffff;
20     andi r5, r9, 0xffff ; b = v1 & 0xffff;
21     srli r6, r10, 16 ; c = (v2 >> 16) & 0xffff
22     andi r7, r10, 0xffff ; d = v2 & 0xffff;
23
24     mul r9, r5, r7 ; LO = b * d;
25     srli r9, r9, 16 ; y = ((LO >> 16) & 0xffff)
26     mul r10, r4, r7 ; x= a * d
27     mul r12, r5, r6 ; x1 = c * b
28     add r10, r10, r12 ; x = x + x1
29     add r9, r9, r10 ; y = y + x
30     srli r9, r9, 16 ; y = (y >> 16) & 0xffff
31     mul r10, r4, r6 ; HI = a * c
32     add r10, r10, r9 ; HI = HI + y
33
34     ;;;;;;;;;;;;;;;
35     ; write result onto the corresponding stack entry
36     ;;;;;;;;;;;;;;;
37     ; store the result to the stack
38     stw r10, 0(r11)
39
40     ; declare this instruction as executed
41     addi ea, ea, 4
42
43

```

And clean up!

```

1
2
3
4
5 iEpilogue:
6 notmulxuu:
7     ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
8     ; restore all registers from the stack
9     ; one value has been changed
10    ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
11    stw r0,0(sp)
12    stw r1,4(sp)
13    stw r2,8(sp)
14    stw r3,12(sp)
15    stw r4,16(sp)
16    stw r5,20(sp)
17    stw r6,24(sp)
18    stw r7,28(sp)
19    stw r8,32(sp)
20    stw r9,36(sp)
21    stw r10,40(sp)
22    stw r11,44(sp)
23    stw r12,48(sp)
24    stw r13,52(sp)
25    stw r14,56(sp)
26    stw r15,60(sp)
27    stw r16,64(sp)
28    stw r17,68(sp)
29    stw r18,72(sp)
30    stw r19,76(sp)
31    stw r20,80(sp)
32    stw r21,84(sp)
33    stw r22,88(sp)
34    stw r23,92(sp)
35    stw r24,96(sp)
36    stw r25,100(sp)
37    stw r26,104(sp)
38    stw r27,108(sp)
39    stw r28,112(sp)
40    stw r29,116(sp)
41    stw r30,120(sp)
42    stw r31,124(sp)
43
44     ; restore the stack
45     addi sp, sp, 32 * 4
46     br idone
47
48     ; for hardware interrupts re-execute instruction that was
49     ; interrupted
50     eadec:
51     subi ea, ea, 4
52     idone:
53     eret

```

A piece of code that uses the new instruction is therefore

```

1 main:
2     movhi r9, 0xffff
3     ori   r9, r9, 0xffff
4     add   r10, r9, r0
5     mulxuu r11, r9, r10

```

SUBSECTION 8.7

Lecture 17: A single cycle processor

For the purposes of this class we will consider a simple CPU capable of executing these 10 instructions (which are encoded as follows)

1. load r1, r2

```

1     TMP = MEM [R2]
2     R1 = TMP
3     PC = PC + 1

```

7	6	5	4	3	2	1	0
R1	R1	R2	R2	0	0	0	0

2. store r1, r2

```

1     MEM [R2] = R1
2     PC = PC + 1

```

7	6	5	4	3	2	1	0
R1	R1	R2	R2	0	0	1	0

3. add r1, r2

```

1     TMP = R1 + R2
2     R1 = TMP
3             IF (TMP == 0) ZERO = 1; ELSE ZERO = 0;
4             IF (TMP < 0) NEGATIVE = 1; ELSE NEGATIVE = 0;
5     PC = PC + 1

```

7	6	5	4	3	2	1	0
R1	R1	R2	R2	0	1	0	0

4. **sub r1, r2**

```

1      TMP = R1 - R2
2      R1 = TMP
3      IF (TMP == 0) ZERO = 1; ELSE ZERO = 0;
4      IF (TMP < 0) NEGATIVE = 1; ELSE NEGATIVE = 0;
5      PC = PC + 1

```

7	6	5	4	3	2	1	0
R1	R1	R2	R2	0	1	1	0

5. **nand r1, r2**

```

1      TMP = R1 bitwise NAND R2
2      R1 = TMP
3      IF (TMP == 0) ZERO = 1; ELSE ZERO = 0;
4      IF (TMP < 0) NEGATIVE = 1; ELSE NEGATIVE = 0;
5      PC = PC + 1

```

7	6	5	4	3	2	1	0
R1	R1	R2	R2	1	0	0	0

6. **ori imm5**

```

1      TMP = K1 bitwise OR IMM5, where IMM5 is a 5 bit
2      ↳ constant
3      K1 = TMP
4      IF (TMP == 0) ZERO = 1; ELSE ZERO = 0;
5      IF (TMP < 0) NEGATIVE = 1; ELSE NEGATIVE = 0;
      PC = PC + 1

```

7	6	5	4	3	2	1	0
IMM5_4	IMM5_3	IMM5_2	IMM5_1	IMM5_0	1	1	1

7. **shift{left, right} r1, r2**

```

1   IF (L) THEN TMP = R1 << IMM2
2   ELSE TMP = R1 >> IMM2
3   R1 = TMP
4   IF (TMP == 0) ZERO = 1; ELSE ZERO = 0;
5   IF (TMP < 0) NEGATIVE = 1; ELSE NEGATIVE = 0;
6   PC = PC + 1
7   Alternative definition (equivalent to the previous one):
8       IMM3 is a 3 bit immediate in the sign - magnitude
9       → representation, i.e., bit 2 = sign, value = bits 1 and 0
10      IF (IMM3 > 0) THEN TMP = R1 << IMM3
11      ELSE TMP = R1 >> (-IMM3)
12      R1 = TMP
13      IF (TMP == 0) ZERO = 1; ELSE ZERO = 0;
14      IF (TMP < 0) NEGATIVE = 1; ELSE NEGATIVE = 0;
15      PC = PC + 1

```

7	6	5	4	3	2	1	0
R1	R1	L/R	IMM2_1	IMM2_0	0	1	1

8. bz imm4

```

1   where IMM4 is a 2's complement 4-bit immediate
2   IF (ZERO == 1) PC = PC + 1 + (SIGN-EXTEND8(IMM4))
3   ELSE PC = PC + 1

```

7	6	5	4	3	2	1	0
IMM4_3	IMM4_2	IMM4_1	IMM4_0	0	1	0	1

9. bnz imm4

```

1   IMM4 as in BZ
2   IF (ZERO == 0) PC = PC + 1 + (SIGN-EXTEND8(IMM4))
3   ELSE PC = PC + 1

```

7	6	5	4	3	2	1	0
IMM4_3	IMM4_2	IMM4_1	IMM4_0	1	0	0	1

10. **bpz imm4**

```

1  IMM4 as in BZ
2  IF (NEGATIVE == 0) PC = PC + 1 + (SIGN-EXTEND8(IMM4))
3  ELSE PC = PC + 1

```



A CPU can be broken down into two units: the datapath and the control path. The data path is where all the data manipulation occurs; it should be capable of computing all results of the instruction execution. The control path orchestrates signals sent to the data path (by means of a giant FSM).

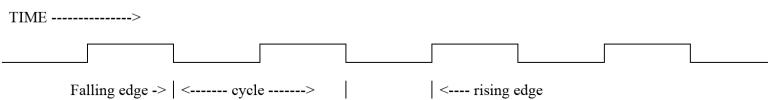


Figure 22. In the single-cycle processor each instruction will be executed entirely inside a single cycle. Code execution starts immediately after a falling edge and then all state changes will be complete right before the next falling edge.

CPU components include:

Definition 18

Register

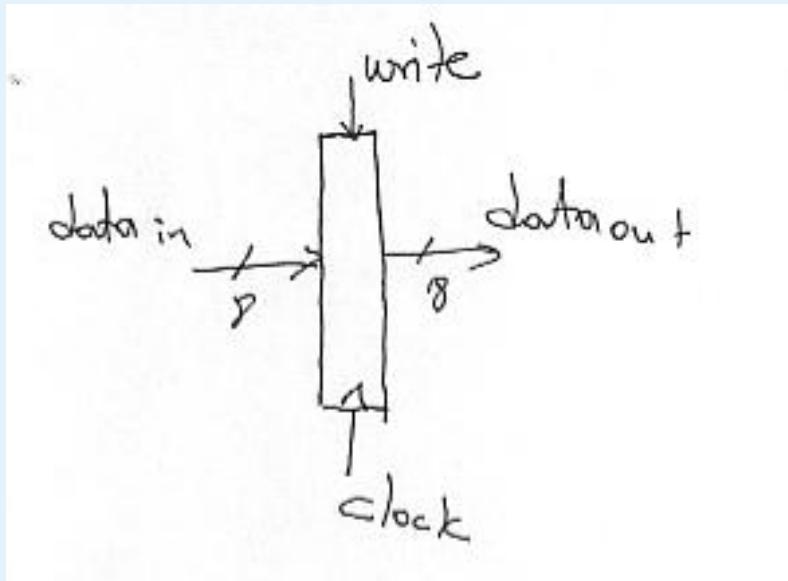


Figure 23. A collection of D flip-flops that can be read or written. Dataout gives values in the register, adding new values via datain. Write wire tells the register to latch onto the values in datain

Definition 19

Mux

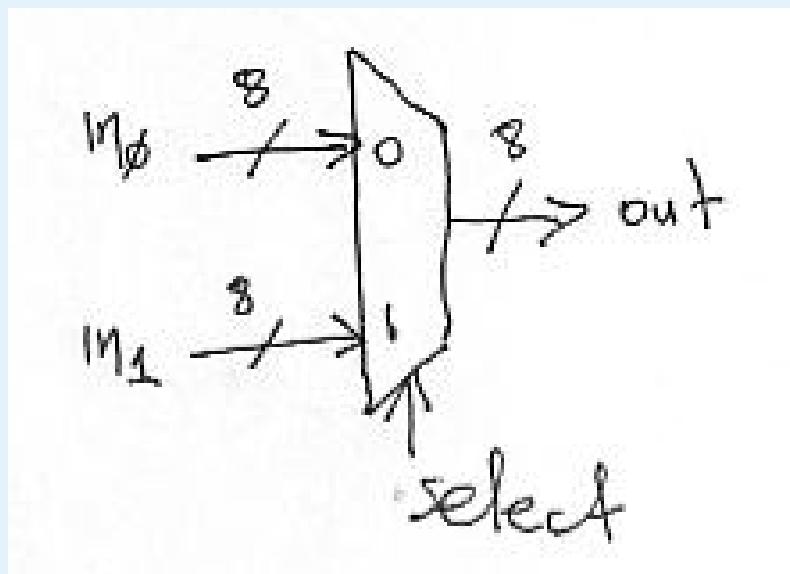


Figure 24. Multiplexes (Muxs) select a signal

Definition 20

Register File

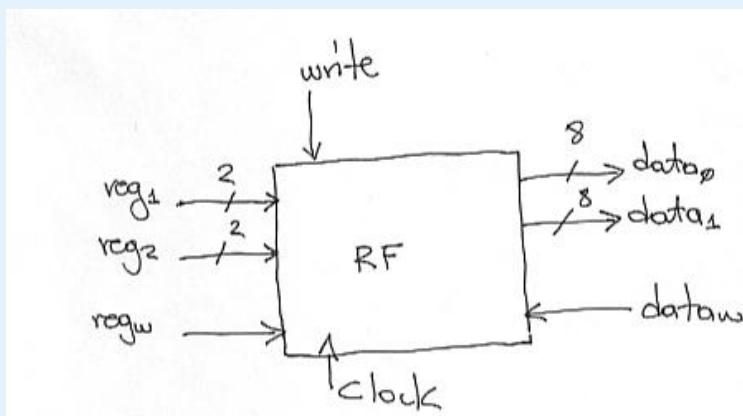


Figure 25. A register file contains a number of registers. Input signals to the register file cause it to output the data from the corresponding register to the data out lines. It also supports selecting and writing to a register by means of `regw` and `dataw`

. Note `write` signal as well.

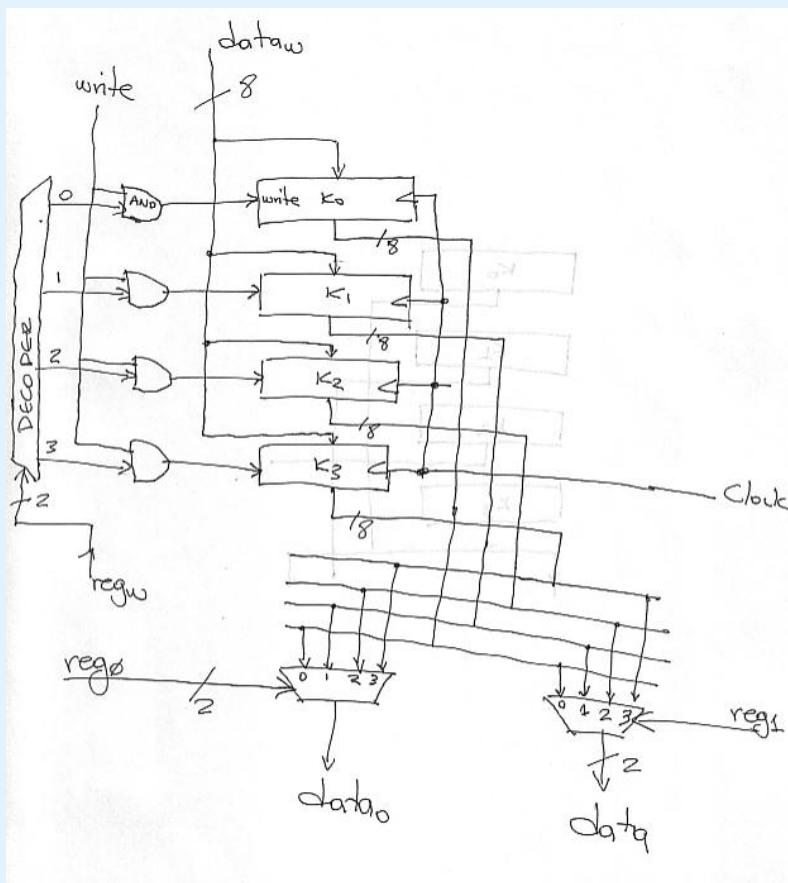


Figure 26. A register file implementation

Definition 21

ALU

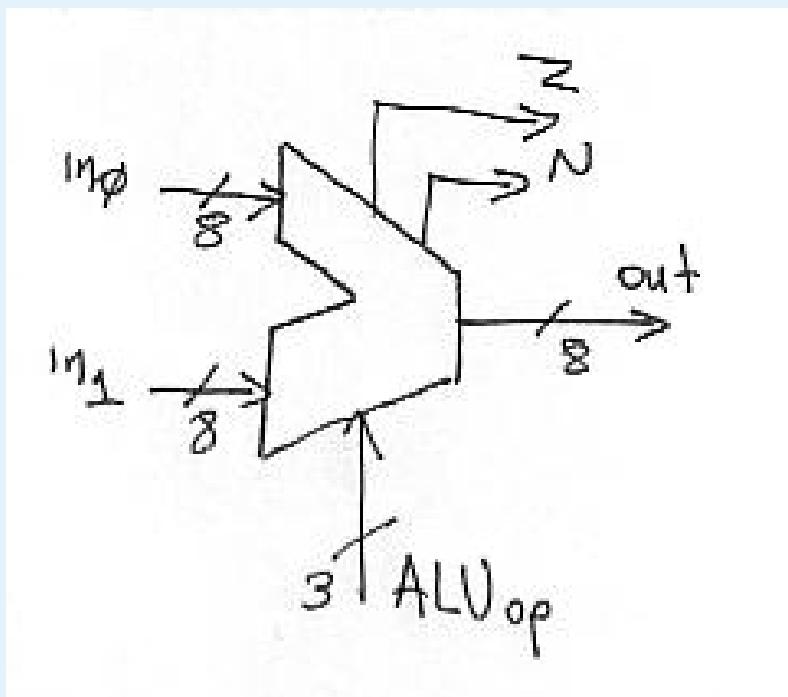


Figure 27. An ALU takes two inputs and then performs an operation as determined by the ALU_{Op} signal. In this generalized ALU block we have Z , N signals which indicate if output values are zero or negative. Simpler ALUs can be used at times if things such as only addition is necessary

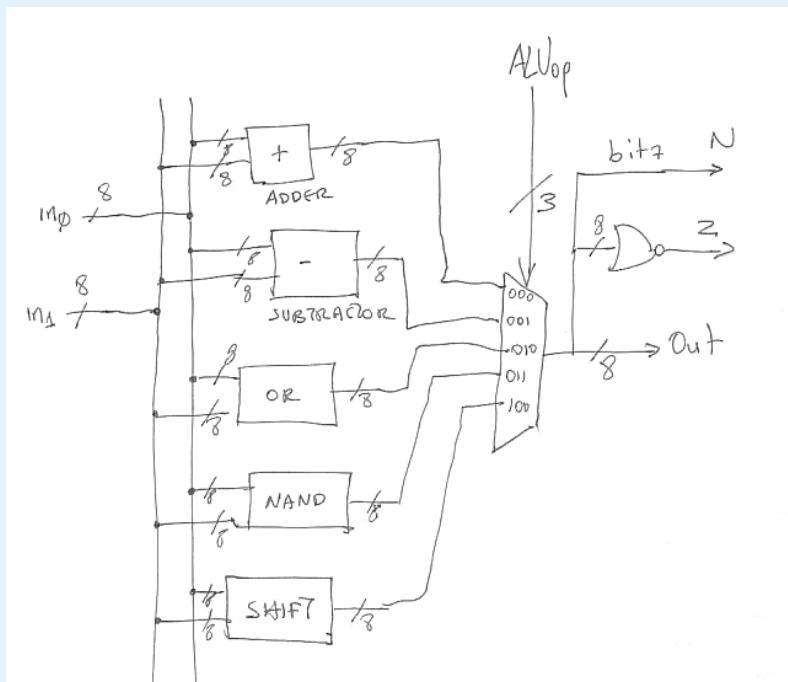


Figure 28. Possible impl

The and/sub boxes are fairly trivial, and so or OR/ANDs. Shift is a little more interesting

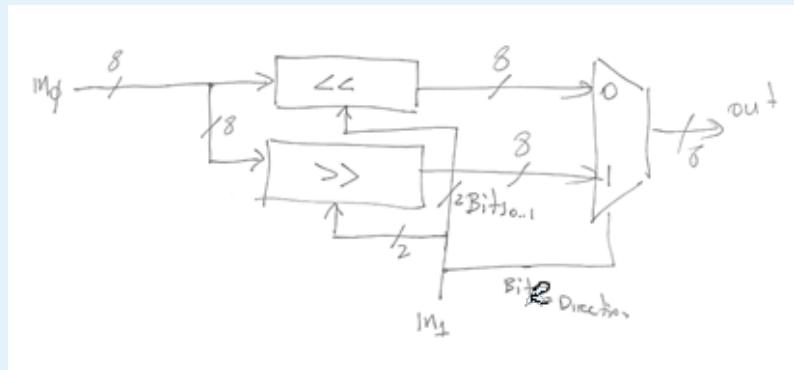


Figure 29. Shift is muxed on a shift left and a shift right

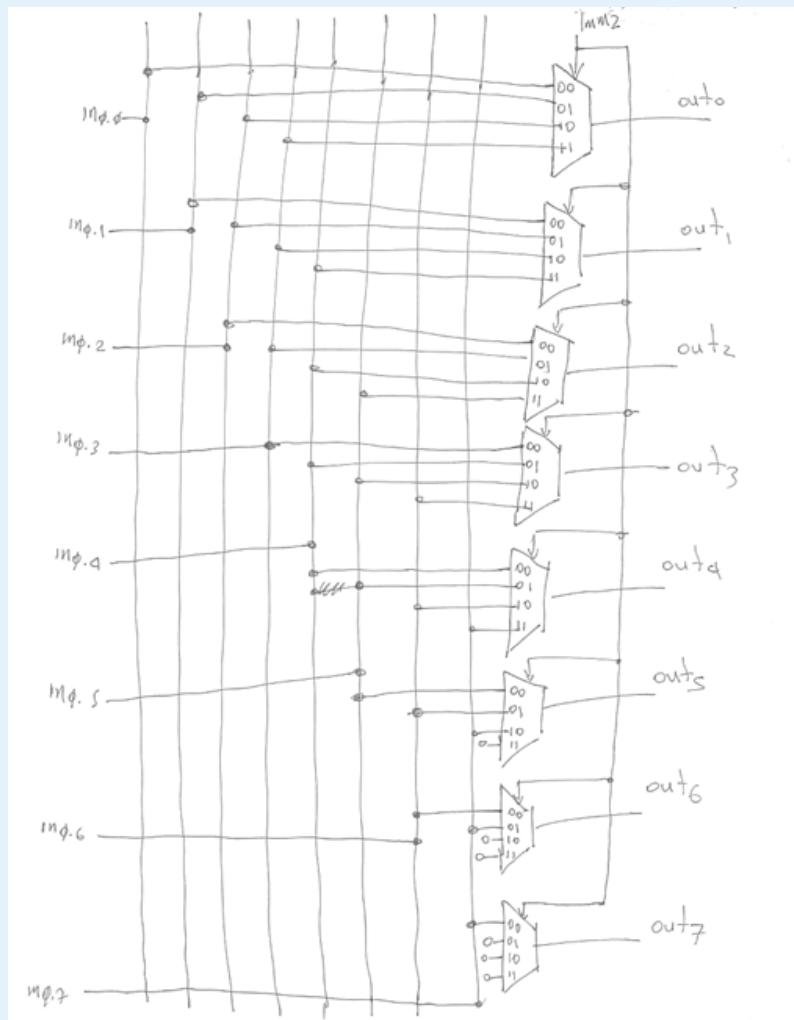
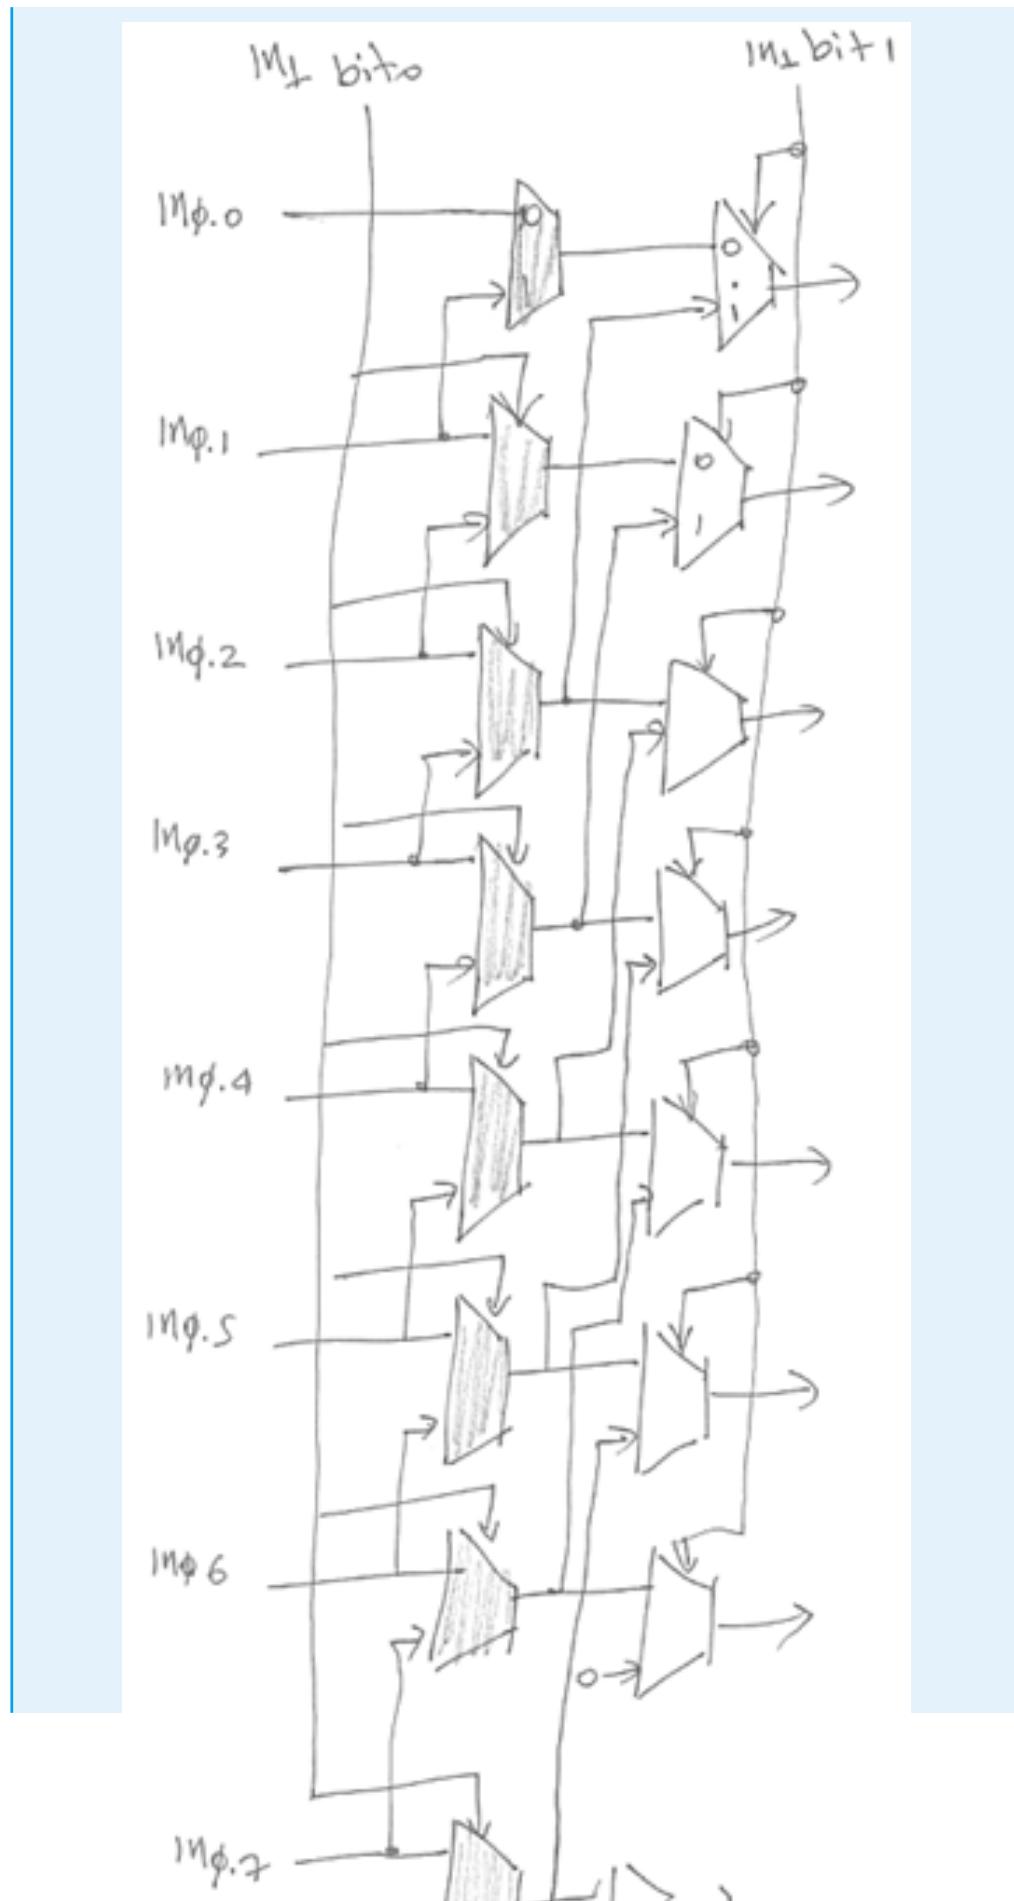


Figure 30. Shift left is a ton of muxes



Most modern processors use unified instruction and data memories. For the single-cycle processor we will use the *Harvard Architecture* which uses separate instruction and data memories.

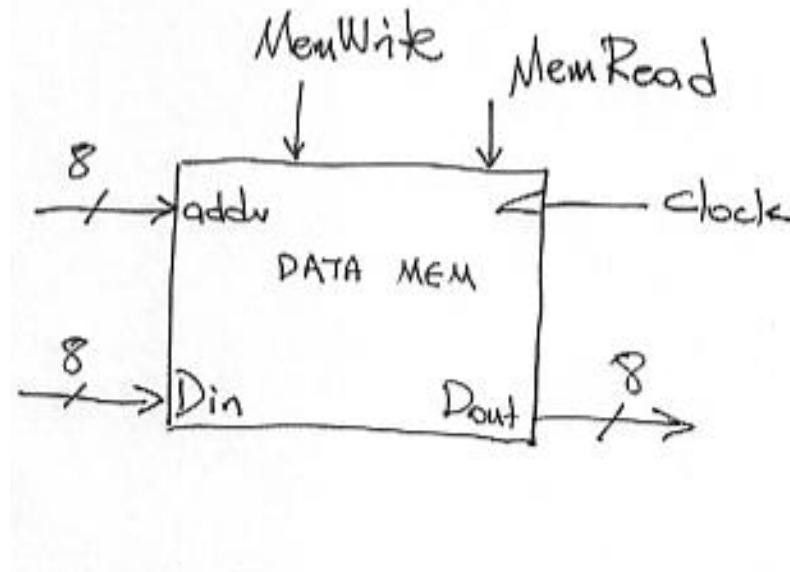


Figure 32. Memory

Data memory is capable of reading and writing (whereas instruction memory is read-only). Generally speaking it takes an address and depending on the value of MemWrite or MemRead it will either update the value at the address with **Din** or present the value at **Dout**

ADD							
7	6	5	4	3	2	1	0
R1	R1	R2	R2	0	1	0	0

Figure 33. Recall: format of ADD instruction

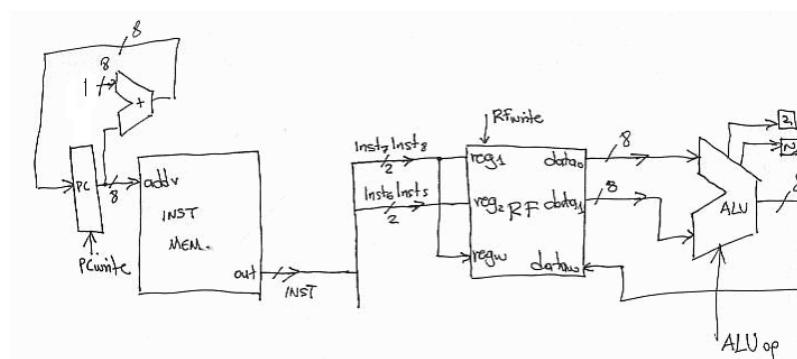


Figure 34. Datapath for ADD

Notice that bits 7 and 8 of the instruction are being piped directly into **reg1** and **regw** since our instruction definition gives $\text{reg1} = \text{reg1} + \text{reg2}$. The adder attached to the PC register is used to increment PC after each cycle.

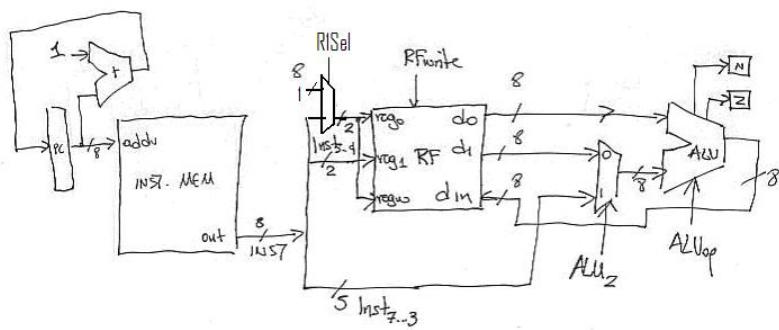


Figure 35. Augmented datapath to support ORI

ORI takes an immediate value from the instruction.

SUBSECTION 8.8

Lecture 18: Modifying the single cycle processor

SUBSECTION 8.9

Lecture 19: Multi-cycle processor

ADD R1, R2

1. Read mem[PC] into IR.

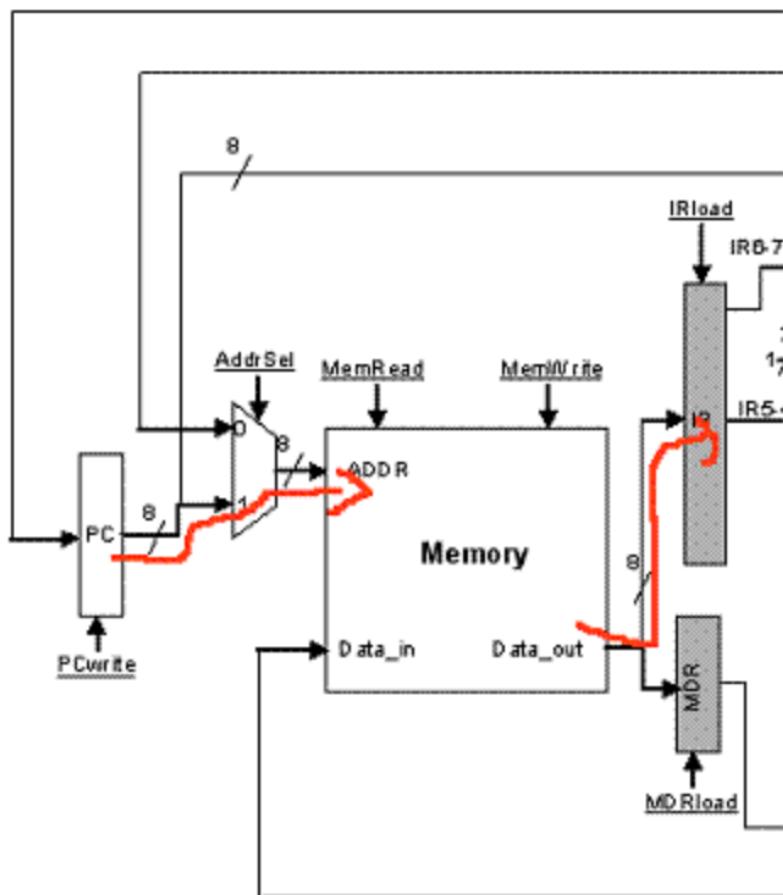


Figure 36. Cycle one for the ADD. This fetches the instruction by sending the address from PC register to MemRead and then storing the stored instruction into the Instruction Register IR

2. Decoding the instruction and reading from the register file

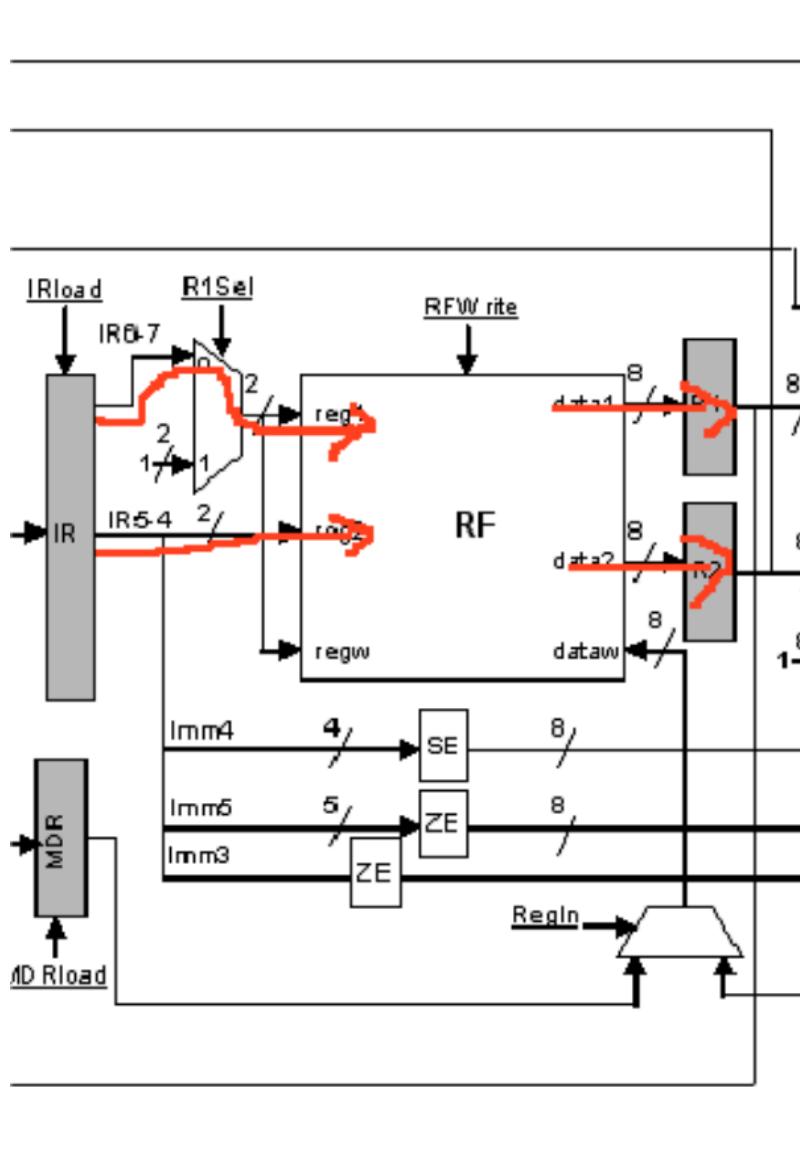


Figure 37. Control uses the instruction opcode to decide what happens during the next cycle. In this case, we read the two registers from the register file RF and then latch them into R1, R2

3. $R1 = RF[127, 6], R2 = RF[IR5, 4]$
4. $AluOut = R1 + R2$, update Z & N (Zero, Negative flags) This is done by muxing the latched values into the ALU (as well as the appropriate instruction as applicable) and then latching the ALU output and zero/negative flags. This step can be combined with the above step (discussed later).
5. $RF[127, 6] = AluOut$
6. Increment PC by 1. At this point the ALU is idle, so we can take a small detour and increment PC.

By the end of this, this just needs to ensure that the programmer sees the desired behaviour; that the values have been added and the PC has been incremented.

What are some optimizations?

- PC: reading at cycle 1, read at cycle 6. So we only need the current value at the beginning of cycles 1 and 6, so if we think about this we can probably speed things up by hoisting cycle 6 to cycle 1. So Cycle 1 can become $PC = PC + 1$. This way we can save a clock cycle. While this may sound good on paper, can it happen during actual execution? In this case nobody else is using the ALU during cycle 1, so this would only necessitate a change in control. This optimization is always ideal, all the time.
- Another optimization can happen during Cycle 2 and 3. It isn't certain that we can move Cycle 3 into Cycle 2; sometimes this will be helpful and sometimes it won't be helpful. But, it'll be helpful *always* and it won't destroy any state. So we can *speculate* that storing these results into the registers will *probably* be useful and combine steps 1 and 2. This gives us

1. $PC = PC + 1$
2. Decode & Read from RF, $R1 = RF[IR7, 6]$, $R2 = RF[IR5, 4]$
3. $AluOut = R1 + R2$, update Z & N (Zero, Negative flags)
4. $RF[IR7, 6] = AluOut$

Comment

Note that cycle 1 and cycle 2 must be the exact same for all instructions – this is because at that point we don't know what the instruction actually is. Only after instruction 2 do they change.

LHS happens at end of clock cycle, RHS happens during

There are situations where we speculate in a way that will destroy the state of the machine and then it turns out that the speculation is not that helpful.

ORI imm5

1. Read mem[PC]
2. Decode the instruction and read from the register file (R1, R2). We've read both but we only need one of them. But it doesn't hurt that we already have R2.
3. $R1 = RF[K1]$
4. $ALUout = R1 \text{ OR imm5}$ (zero-extended)

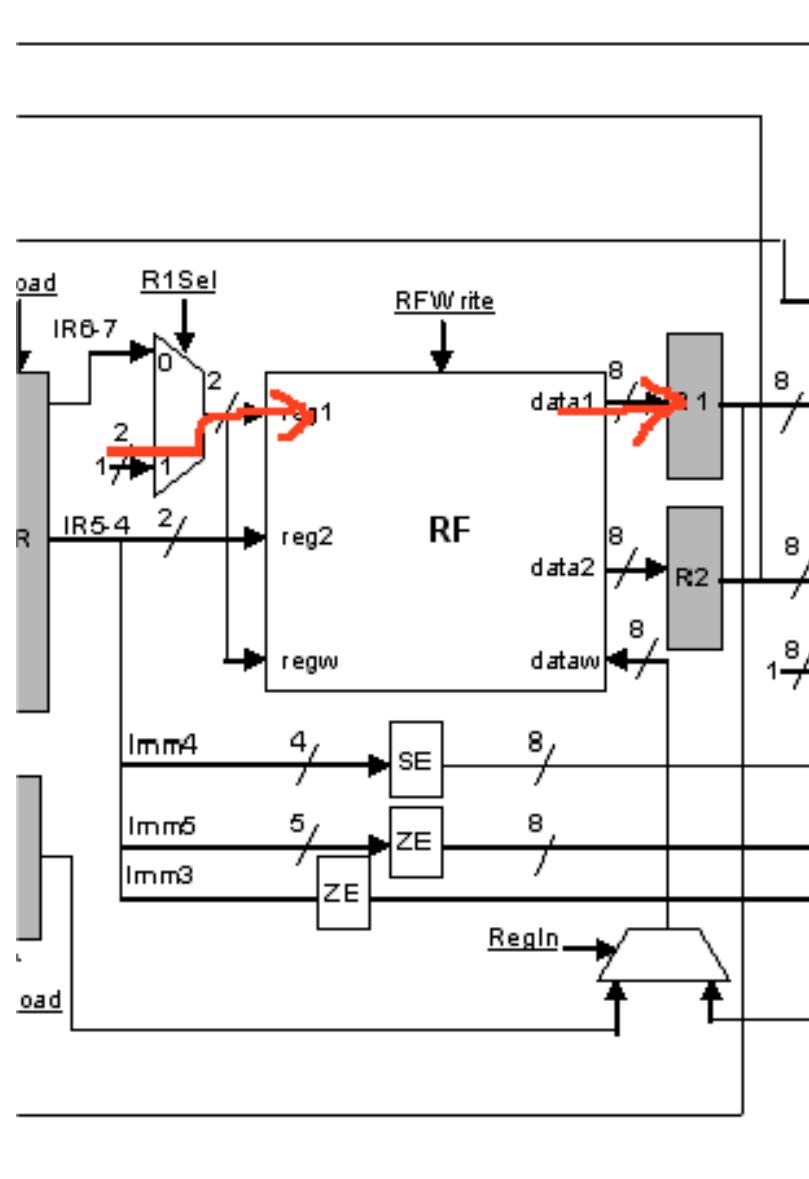


Figure 38. Note that imm5 is actually encoded into the instruction (see the control path)

5. $RF[K1] = ALUout$

LOAD

We could've actually speculated another way to favor the ORI instruction to make it take 4 cycles instead. But statistically speaking the NAND, ADD, and SUB instructions end up being used a lot more often and therefore more useful to optimize for

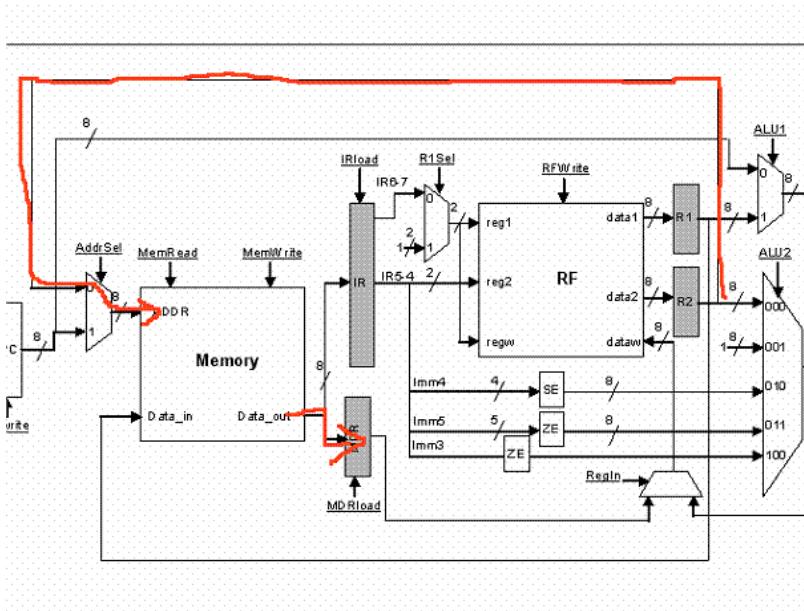


Figure 39. The LOAD instruction accesses memory during cycle 3 and then stores the returned value into MDR.

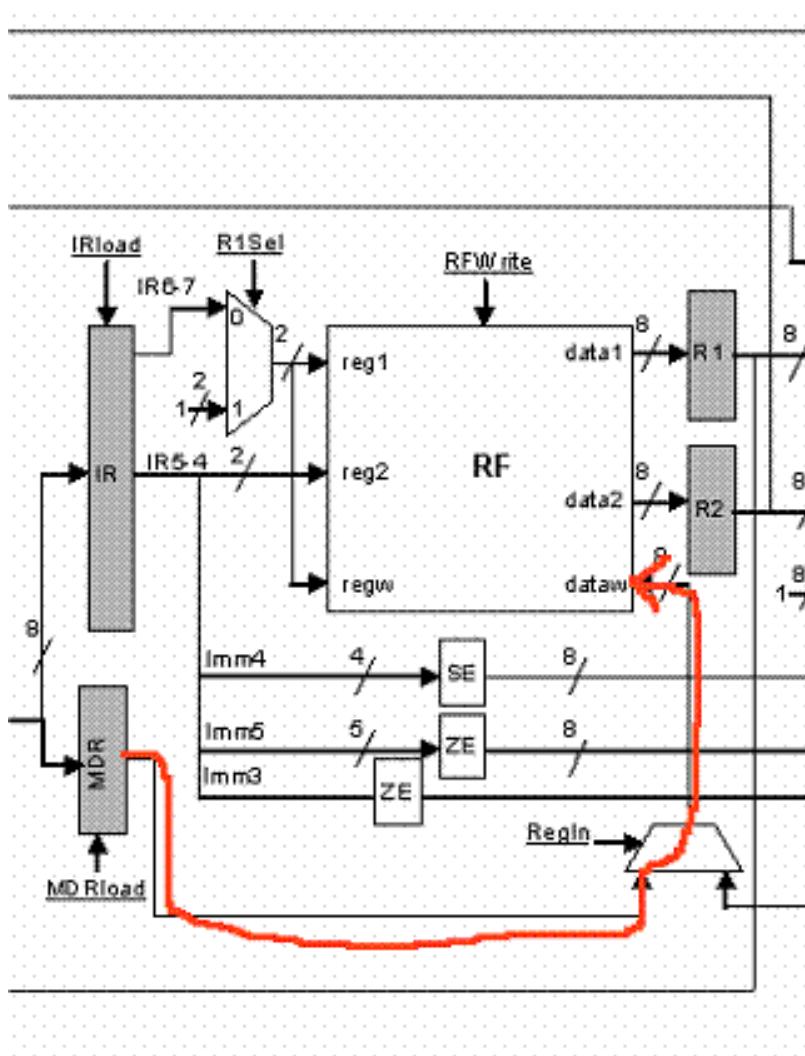


Figure 40. In cycle 4 the MDR value can be written to the register file.

STORE

1. PC = PC+1; grab value & +1 using ALU
2. Grab the two register values from the register file and latch them
- 3.

BZ Imm4

- Read mem[PC]
- Decode
- If (Z is true), PC = PC + 1 + SE(Imm4)¹⁸
-
-
- Decode
- Read

¹⁸ SE denotes sign extend

SUBSECTION 8.10

Lecture 20: Modifying the multi-cycle processor

SUBSECTION 8.11

Caches

Caches are small and fast memories which contain a subset of the current contents of main memory. Modern processors usually include three levels of caches; L1, L2, and L3, with L1 being the fastest. Caches are also invisible to the ISA and the programmer with the exception of methods such as `ldwio`, `stwio`¹⁹.

¹⁹and volatile in c, etc, I guess

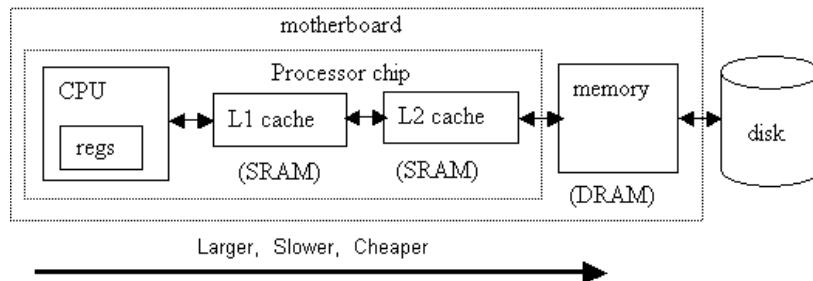


Figure 41. Cache and memory hierarchy

Caches exploit a common tendency of programs: spatial and temporal *locality*; recently referenced items are likely to be referenced again soon and nearby elements tend to be referenced close together in time.

For example, consider

```

1 for (int i = 0; i < n; i++){
2     sum += A[i];
3 }
```

This snippet displays spatial locality in the data access of the array and temporal locality in that the body of the loop composes a sequence of instructions that will be fetched and executed repeatedly. By storing a subset of memory into spatially local *blocks*, we can often reduce the amount of disk accesses by having a, well, cached copy available nearby.

Typical processor implementations have separate L1 caches for instructions and data, and a single (larger) L2 cache that is shared between the two.

Some keywords:

- cache capacity: the number of bytes the cache can hold
- placement: where the cache block should be placed
- identification: how to find a block in the cache
- cache hit: when the requested block is in the cache
- hit rate: total number of hits / total number of accesses
- cache miss: when the requested block is not in the cache

- miss rate: total number of misses / total number of accesses
- replacement: remove block to make room for new block
- replacement strategy: how to choose which block to remove
- write strategy: how to handle writes to the cache

Cache operation usually involves the following steps:

1. CPU loads from addr A
2. If cache hit \rightarrow return value at A stored in cache [DONE]
3. If cache miss \rightarrow load block containing A into cache
4. Place block in cache, replacing another block if necessary
5. Return value at A stored in cache [DONE]

Since cache is a subset of memory, we need to be able to identify which block of memory we're looking for. Cache lookups by convention break up the bits of a memory address into the *tag*, *set index*, and *offset*.

- *tag* (t): since multiple locations can map to each cache block, the tag is used to distinguish between them i.e. which memory location currently is cached
- *set index* (s): how to locate a block in the cache: index into the array of sets containing blocks in the cache
- *offset* (b): which byte within the block we're looking for

A *valid bit* is also associated for every block in the cache since a 0x0 tag is a valid address.

Tag: $2^{n_{rows}} \rightsquigarrow \text{tag} = 32 - N$. Note that the number of ways does not have to be a power of 2. We can have a 3-way cache, etc.

8.11.1 Implementations

Definition 22

Direct mapped cache: each memory location maps to a single specific location in the cache. For a byte-addressable address space of 2^n , then the size of a block is $B = 2^b$ bytes, the number of sets $S = 2^S$, and S is also equal to the total number of cache blocks in the cache; i.e. there is one block per set in cache. Capacity of a direct mapped cache is therefore $B \cdot S = 2^{b+s}$ bytes

Example

Consider a 16 bit address space with a direct-mapped cache with $t = 8, s = 4, b = 4$. The capacity of this cache is $2^{s+b} = 256$ bytes. The number of sets is $2^4 = 16$ and the block size $b = 2^4 = 16$.

Set Index	Valid?	Tag	Hex data values (for bytes 15..0)
0x0	0	0x00	15 14 13 12 11 10 09 08 07 06 05 04 03 02 01 00
0x1	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x2	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x3	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x4	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x5	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x6	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x7	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x8	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0x9	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0xa	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0xb	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0xc	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0xd	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0xe	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0xf	0	0x00	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00

Figure 42. Initial cache state

```
movia r8,0xface; ldb r8,0(r8)
```

1. Break address `0xface` to `tag = 0xfa`, `set = 0xc`, `offset = 0xe`.
2. The set-index is used to index the cache (`0xc`)
3. Check the valid bit if the block is valid. Initially this is invalid, so we will have a cache miss.
4. A request will be made to main memory to load the block containing the address `0xface`, i.e. the cache block that starts at `0xfac0` and contains 16 bytes of data between `0xfac0 ~ 0xfacf`
5. Memory returns the block (after a delay) and the cache set at `0xc` is updated with the new block; the contents are copied over, the valid bit set, and the tag is set equal to `0xfa`
6. Use the offset value `0xe` to index into the cache block and return the requested block

A cache hit occurs when the tag matches and the valid bit is set. A more interesting case is a cache miss involving an overwrite.

Consider

Out of consideration of typing time, the full example will not be typed out and instead I defer the reader to find it [here](#).

8.11.2 Handling Writes

What are ways to handle writes to the cache as to maximize performance i.e. minimize cache misses?

There are two things we can do:

1. write something
2. don't write anything

When do we pick one over the other? Some heuristics include

1. LRU cache: kick the least recently used block out of the cache²⁰
2. Random: kick a random block out of the cache

²⁰Hashmap to doubly linked list, or ring buffer for fixed size

PART

III

ECE355: Signal Analysis and Communication

SECTION 9

Taught by Prof. Sunila Akbar

Admin and Preliminary

SUBSECTION 9.1

Lecture 1

- CT and DT signals
- A ton of LTI (Linear time invariant) systems
- Processing of signals via LTI systems
- Fourier transforms
- Sampling

9.1.1 Mark Breakdown

Table 1. Mark Breakdown

Homework	20
MT1	20
MT2	20
Final	40

- Continuous enclose in (), independent is t
- Discrete: enclose in [], independent is n

In many systems we are interested in power and energy of signals over an infinite time interval; $-\infty < \{t, n\} < \infty$

Theorem 2**Energy for Complex Signals**

$$E_{[t_1, t_2]} = \int_{t_1}^{t_2} |x(t)|^2 dt \quad (9.1)$$

$$E_{[t_1, t_2]} = \sum_{n=n_1}^{n_2} |x(n)|^2 \quad (9.2)$$

Average Power for Complex Signals

$$P_{avg, [t_1, t_2]} = \frac{1}{t_2 - t_1} \int_{t_1}^{t_2} |x(t)|^2 dt \quad (9.3)$$

$$P_{avg, [t_1, t_2]} = \frac{1}{n_2 - n_1 + 1} \sum_{n=n_1}^{n_2} |x(n)|^2 \quad (9.4)$$

SECTION 10

Transformations

SUBSECTION 10.1

Lecture 2

Most of this lecture was review. When applying transforms just note to always scale, *then* shift, i.e.

$$1. y(t) = x(\alpha t)$$

$$2. y(t) = x(\alpha t + \frac{\beta}{\alpha})$$

Definition 23**Fundamental Period**

$$x_t = x(t + mT), m \in \mathbb{Z} \quad (10.1)$$

The fundamental period, T_o is the smallest positive value of T for which this holds true

Definition 24**Even signals**

$$x(t) = x(-t) \quad (10.2)$$

Definition 25**Odd signals**

$$x(t) = -x(-t) \quad (10.3)$$

Theorem 3

Any signal can be broken into an even and odd component

$$\begin{aligned} x(t) &= Ev \{x(t)\} = \frac{1}{2} [x(t) + x(-t)] \\ x(t) &= Od \{x(t)\} = \frac{1}{2} [x(t) - x(-t)] \end{aligned} \quad (10.4)$$

SUBSECTION 10.2

Lecture 3

Again, most of this lecture was review from the waves portion of PHY293 from last year or some other course prior.

A complex exponential and sinusoidal system can be represented as

$$x(t) = Ce^{at} \quad (10.5)$$

Where C, a are complex numbers.

Two cases may occur.

If a imaginary and C is real we have, depending on ω , either a constant signal or a periodic sinusoidal system.

$$x(t) = e^{j\omega_0 t} \quad (10.6)$$

- Important property: this is periodic, i.e. $Ce^{j\theta_0 t} = Ce^{j\theta(t+T)}$
- Implies that $e^{j\omega_0 T} = 1$
- Implies that for $\omega \neq 0 \rightarrow T_0 = \frac{2\pi}{|\omega_0|}$

On the other hand, if a imaginary and C complex, we have a periodic signal with $T = \frac{2\pi}{\omega_0}$

$$x(t) = Ce^{j\omega_0 t} = |C|e^{j\omega_0 t+\phi} = |C|\cos(\omega_0 t + \phi) + j|C|\sin(\omega_0 t + \phi) \quad (10.7)$$

The energy of the signal is given by (9.1), or

$$E_{period} = \int_0^{T_0} |e^{j\omega_0 t}|^2 dt = \int_0^{T_0} 1 dt = T_0 \quad (10.8)$$

$$P_{period} = \frac{E_{period}}{T_0} = 1 \quad (10.9)$$

Recall: implication that $e^{j\omega_0 T} = 1$, therefore the quantity inside the integral evaluates to 1

10.2.1 General Continuous Complex Exponential Signals

The most general case of a complex exponential can be represented as a combination of the real exponential and the periodic complex exponential;

$$C = Ce^{at} \quad (10.10)$$

and

$$a = r + j\omega_0 \quad (10.11)$$

can be combined to give

$$Ce^{at} = |C|e^{rt}e^{j\omega_0 t+\theta} \quad (10.12)$$

Euler's relation can be used to simplify this to

$$Ce^{at} = |C|e^{rt}(\cos(\omega_0 t + \theta) + j\sin(\omega_0 t + \theta)) \quad (10.13)$$

By inspection we can see that the signal has the following properties:

- $r = 0$: real and imaginary parts of sinusoidal

Definition 26

2. $r > 0$: sinusoidal signal with exponential growth
3. $r < 0$: sinusoidal signal with exponential decay

I will be skipping notes on the discrete case as it is essentially the same as the continuous case, but with the following differences

Table 2. Comparison of continuous and discrete complex exponential signals

$e^{j\omega_0 t}$	$e^{j\omega_0 n}$
Distinct signals for distinct ω_0	identical signals for distinct $\omega_0 \in \{\omega_0 \pm 2\pi i, i \in \mathbb{Z}\}$
Periodic for any ω_0	Periodic only if $\omega_0 = \frac{2\pi m}{N}$ for integers $N > 0, m$
Fundamental frequency ω_0	Fundamental frequency $\frac{\omega_0}{m}$
Fundamental period $\omega_0 = 0 \rightarrow$ undefined, otherwise $T_0 = \frac{2\pi}{\omega_0}$	Fundamental period $\omega_0 = 0 \rightarrow$ undefined, otherwise $T_0 = m \frac{2\pi}{\omega_0}$
	Since unique ω does not mean unique signal, pick $0 \leq \omega_0 \leq 2\pi$ or $-\pi \leq \omega_0 \leq \pi$

SECTION 11

Basic Signals

SUBSECTION 11.1

Lecture 4: Step and Impulse Functions

One of the simplest discrete-time signals is the **unit impulse**²¹ function, $\delta[n]$

²¹ or unit sample

Definition 27

$$\delta[n] = \begin{cases} 0 & n \neq 0 \\ 1 & n = 0 \end{cases} \quad (11.1)$$



Figure 1.28 Discrete-time unit impulse (sample).

Another basic signal is the **unit step** function, $u[n]$

Definition 28

$$u[n] = \begin{cases} 0 & n < 0 \\ 1 & n \geq 0 \end{cases} \quad (11.2)$$

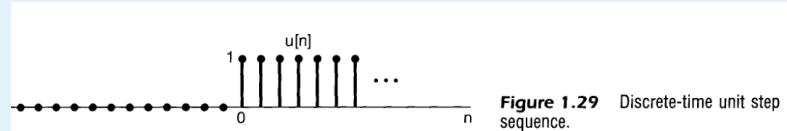


Figure 1.29 Discrete-time unit step sequence.

The unit impulse function is the first difference of the discrete time step function, i.e.

$$\delta[n] = u[n] - u[n - 1] \quad (11.3)$$

And the unit step function is the running sum of the unit impulse function, i.e.

$$u[n] = \sum_{m=-\infty}^n \delta[m] \quad (11.4)$$

This can be rewritten with $k = n - m$ to make a more convenient expression for moving the function along $-\infty \dots 0 \dots \infty$

$$u[n] = \sum_{k=0}^{\infty} \delta[n - k] \quad (11.5)$$

Theorem 4

The unit impulse function $\delta[n - n_0]$ can be used to sample a function at a specific $n = n_0$ since the impulse function will take on the value 0 for all values of $n \neq n_0$

$$x[n]\delta[n - n_0] = x[n_0]\delta[n - n_0] \quad (11.6)$$

The continuous equivalents of the unit impulse and unit step functions are defined similarly

Definition 29

$$u(t) = \begin{cases} 0 & t < 0 \\ 1 & t > 0 \end{cases} \quad (11.7)$$

Likewise, the continuous unit step function is a running sum integral of the continuous unit impulse function

$$u(t) = \int_{-\infty}^t \delta(\tau) d\tau \quad (11.8)$$

A relationship analogous to the discrete case can be found for the continuous case; the continuous unit impulse function can be thought of as the first derivative of the continuous-time unit step function

Definition 30

$$\delta(t) = \frac{du(t)}{dt} \quad (11.9)$$

(11.9) is discontinuous at $t = 0$ so it is non-differentiable. We can address this by considering an approximation of (11.9) for a Δ short enough to not matter for any practical purpose

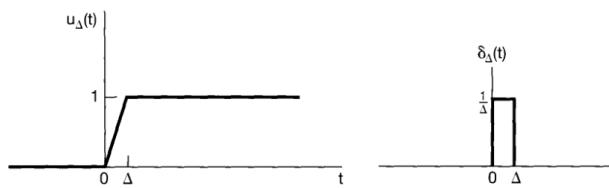


Figure 1.33 Continuous approximation to the unit step, $u_D(t)$.

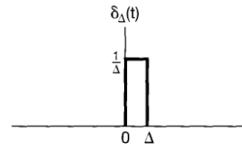


Figure 1.34 Derivative of $u_D(t)$.

(11.8) can be rewritten as follows to make it more convenient to use along $\sigma \in -\infty \dots 0 \dots \infty$.

$$u(t) = \int_0^\infty \delta(t - \sigma) d\sigma \quad (11.10)$$

Theorem 5

And by the same argument as for the discrete case, the continuous impulse function has an important sampling property.

For any arbitrary point t_0 ,

$$x(t)\delta(t - t_0) = x(t_0)\delta(t - t_0) \quad (11.11)$$

SUBSECTION 11.2

Lecture 5**Definition 31**

A **system** is a process that transforms a signal, i.e.

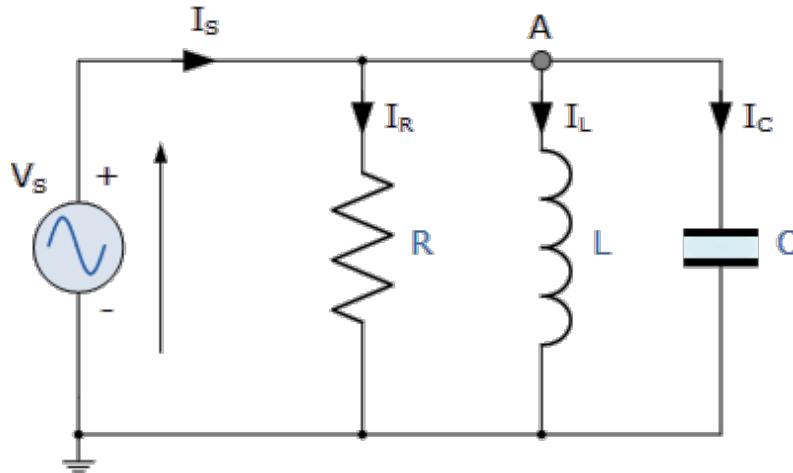
$$x(t) \xrightarrow{\text{sys}} y(t) \quad (11.12)$$

A similar definition can be applied for the discrete case

Some examples of the relationship between signals and physical systems include circuits, i.e. RLC circuits and spring-mass systems.

Example

For example a RLC circuit can be modeled as a system that transforms a voltage signal into a current signal;



$$v(t) \xrightarrow{\text{RLC}} i(t) \quad (11.13)$$

Solving and modelling this system was covered in the other circuit classes.

Systems can be combined, i.e. making a mobile call

$$a(t) \xrightarrow{\text{mic}} y(t) \xrightarrow{\text{antenna}} z(t) \xrightarrow{\text{tower}} u(t) \xrightarrow{\text{antenna}} w(t) \xrightarrow{\text{speaker}} b(t) \quad (11.14)$$

11.2.1 Types of systems

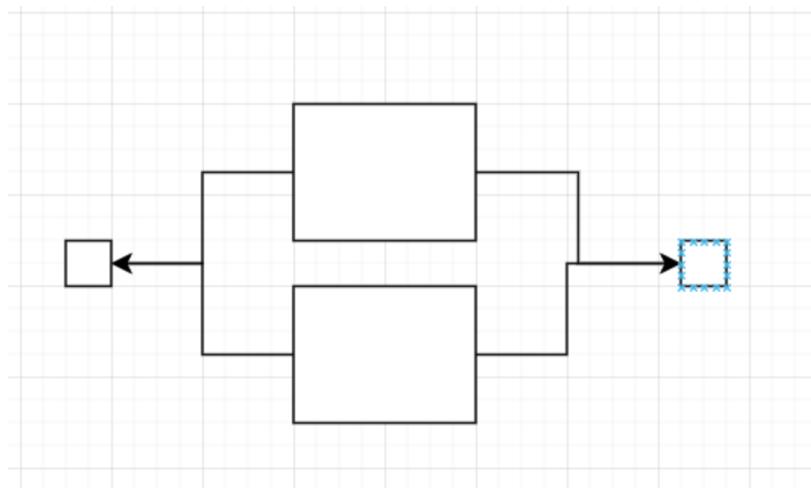


Figure 43. Parallel systems

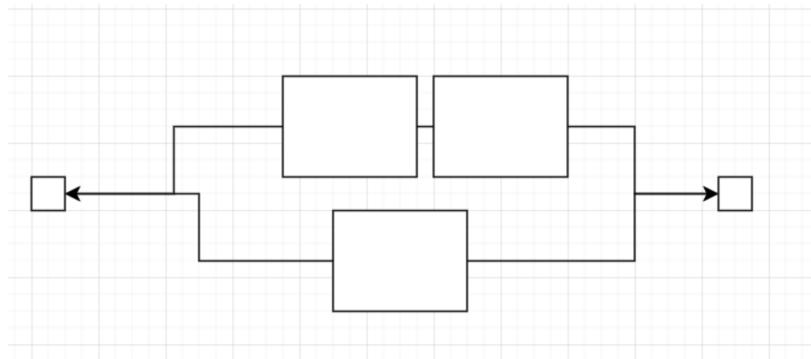
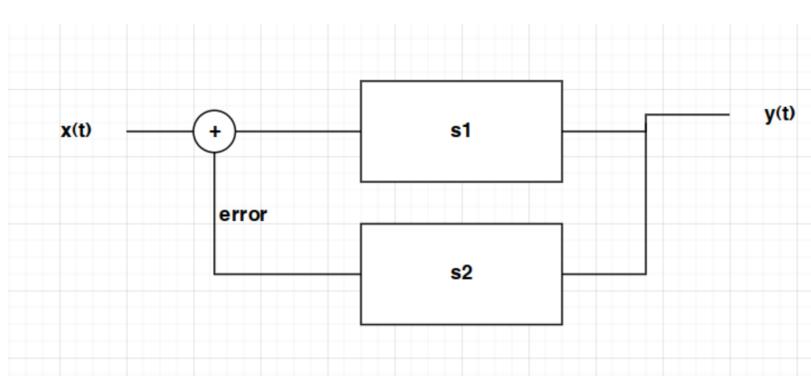


Figure 44. Series-parallel



The feedback control system will be discussed in depth in the control systems class

Figure 45. Feedback control system

11.2.2 System properties

Definition 32

Memoryless systems are systems where its output of the independent variable at a given time is dependent only on the input at the same time.

For example,

$$y[n] = 2x[n] - x^2[n] \quad (11.15)$$

is a memoryless system, but

$$y[n] = 2x[n] - x^2[n-1] \quad (11.16)$$

is not.

Other simple memoryless systems include the identity system $x(t) = y(t)$

A system with memory is the *accumulator*

$$y[n] = \sum_{k=-\infty}^n x[k] \quad (11.17)$$

A capacitor is an example of a continuous-time system with memory

$$v(t) = \frac{1}{C} \int_{-\infty}^t i(\tau) d\tau \quad (11.18)$$

There can also be systems that are dependent on future values of the input and output²².

²²PID go brr

Definition 33

A system is *invertible* if distinct inputs lead to distinct outputs²³ For example, the identity system is invertible, but the accumulator is not.

²³Recall: MAT185 and matrix invertability

Definition 34

A system is *casual* if its output at a given time is dependent only on the input at the current time and in the past.

Then it follows that

Lemma 2 | All memoryless systems are causal

Though causal systems are useful, non-causal systems are also of great utility in modelling systems in which the independent variable is not time, or in *anticipative* models that account for the future values of the input or output, i.e a controller.

Definition 35

A system is *stable* if the output is bounded if the input is bounded.

Definition 36

A system is *time invariant* if the behaviour and characteristics are fixed over time. For example, a *RC* circuit is time-invariant if the circuit *R*, *C* values are constant over time. More formally, a system is time invariant if a time shift in the input signal results in an identical time shift in the output.

If

$$x[n] \xrightarrow{\text{sys}} y[n] \quad (11.19)$$

Then, for a time invariant system,

$$y[n - n_0] \xrightarrow{\text{sys}} x[n - n_0] \quad (11.20)$$

Definition 37

A system is *linear* if it possess the property of superposition, i.e. it possess the additive property and the scaling, or homogeneity property.

More formally, a linear system is one such that

$$ax_1(t) + bx_2(t) \xrightarrow{\text{sys}} ay_1(t) + by_2(t) \quad (11.21)$$

Where y_1 is the output of a system with input x_1 and y_2 is the output of a system with input x_2 .

SECTION 12

LTI systems

SUBSECTION 12.1

Lecture 6: Discrete LTI systems

Many physical systems can be modelled as linear time-invariant (LTI) systems. This is useful because there exists a large body of theory that can be applied to LTI systems, in part due to LTI systems possessing the superposition property.

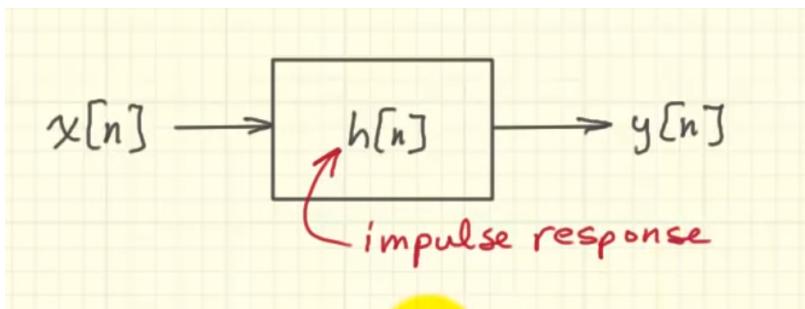


Figure 46. In this lecture the idea of LTI systems being a sum of $h[n]$ impulse responses will be

Theorem 6

Discrete signals can be represented as a summation of impulses.

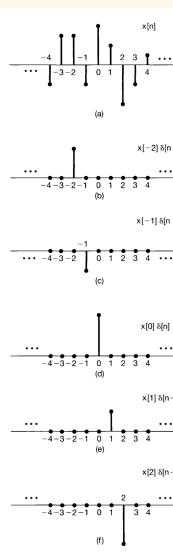


Figure 2.1 Decomposition of a discrete-time signal into a weighted sum of shifted impulses.

More generally, a discrete signal can be rewritten as a sum of shifted unit impulses with weights $x[k]$

$$x[n] = \sum_{n=-\infty}^{\infty} x[k]\delta[n-k] \quad (12.1)$$

Example As an example we can rewrite the unit step function as

$$u[n] = \sum_{n=0}^{\infty} \delta[n-k] \quad (12.2)$$

A useful result of the properties of LTI systems (i.e. the superposition property) is that the output of a system is the convolution of the input and the impulse response of the system, i.e.

Definition 38

Convolution sum

$$y[n] = \sum_{k=-\infty}^{\infty} x[k]h[n-k] \quad (12.3)$$

Theorem 7

An LTI system can be modelled as a **convolution** sum, or a sum of scaled and shifted impulse responses. So, if we know the response of the linear system to the set of shifted unit impulses – we can determine the response of the system to any input signal!

Symbolically convolution can be written as

$$y[n] = x[n] * h[n] \quad (12.4)$$

SUBSECTION 12.2

Lecture 7: Continuous LTI systems

The concept of a LTI system responding to unit impulses and taking advantage of the sifting property may be extended to the continuous systems by idealizing the pulse as one so short that its duration is inconsequential for any real system.

Definition 39

Sifting property of continuous-time impulse

$$x(t) = \int_{-\infty}^{\infty} x(\tau)\delta(t-\tau)d\tau \quad (12.5)$$

Can interpret $x(\tau)\delta(t-\tau)$ equals $x(t)\delta(t-\tau)$, i.e. that it is a *scaled impulse*

$$x(t) = \int_{-\infty}^{\infty} x(\tau)\delta(t-\tau)d\tau = x(\tau) \int_{-\infty}^{\infty} \delta(t-\tau)d\tau \quad (12.6)$$

If we let $h_{\tau}(t)$ denote the response at time t to an unit impulse $\delta(t-\tau)$ at time τ , then

$$y(t) = \lim_{\Delta \rightarrow 0} \sum_{k=-\infty}^{\infty} x(k\Delta)\hat{h}_{k\Delta}(t)\Delta \quad (12.7)$$

As $\Delta \rightarrow 0$ the summation becomes an integral, i.e.

The discrete unit impulse also possesses a *sifting* property; since $\delta[n-k]$ is nonzero only when $k = n$, this summation will sift through the values of $x[k]$ and preserves only the values corresponding to $k = n$

$$y(t) = \int_{-\infty}^{\infty} x(\tau)h_{\tau}(t)d\tau \quad (12.8)$$

For notational convenience the τ subscript can be dropped and generalizing to all t

$$y(t) = \int_{-\infty}^{\infty} x(\tau)h(t-\tau)d\tau \quad (12.9)$$

Symbolically the convolution is represented $y(t) = x(t) \cdot h(t)$

SUBSECTION 12.3

Lecture 12

A class of LTI systems that are really important are ones related through linear constant-coefficient differential equations. For example, consider

$$\frac{dy(t)}{dt} + 2y(t) = x(t) \quad (12.10)$$

This equation offers an implicit specification of the system, i.e. a relationship between input and output instead of an explicit expression; to find a explicit expression we must plug in initial conditions.

Given the exponential input

$$x(t) = Ke^{3t}u(T) \quad (12.11)$$

to (12.10), the complete solution to a differential equation of this type is given as the sum of a particular and homogeneous solution²⁴

²⁴ Recall ESC195

$$y(t) = y_p(t) + y_h(t) \quad (12.12)$$

A common method for finding the particular solution is to first look for the *forced response*; an output signal with the same form as the input.

Take

$$y_p(t) = Ye^{3t} \quad (12.13)$$

and then doing some math we find $Y = \frac{K}{5}$ to arise at $y_p(t) = \frac{K}{5}e^{3t}, t > 0$. To find the homogeneous solution we hypothesize a solution of form

$$y_h(t) = Ae^{st} \quad (12.14)$$

and then substitute into the differential equation to find

$$y(t) = Ae^{-2t} + \frac{K}{5}e^{3t}, \quad t > 0 \quad (12.15)$$

The constants A, K can then be found by substituting in the initial conditions.

More generally speaking, a N th order linear constant-coefficient differential equation is given by

$$\sum_{k=0}^N a_k \frac{d^k y(t)}{dt^k} = \sum_{k=0}^M b_k \frac{d^k x(t)}{dt^k} \quad (12.16)$$

The discrete time counterpart is

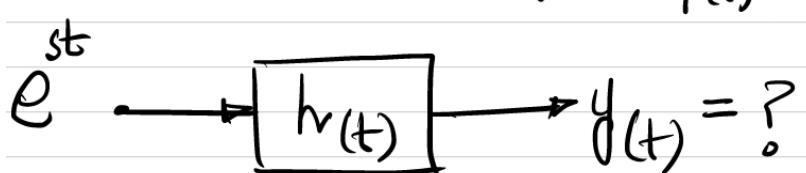
$$\sum_{k=0}^N a_k y[n-k] = \sum_{k=0}^M b_k x[n-k] \quad (12.17)$$

With the solution $y[n]$ being the sum of the particular solution to the problem above and a solution to the homogeneous equation

$$\sum_{k=0}^N a_k y[n-k] = 0 \quad (12.18)$$

A important special case of LCCDEs are LTIs.

Consider a response to an arbitrary exponential input signal



Definition 40

Laplace transform of h , or the transfer function of LTI systems

$$y(t) = \int_{-\infty}^{\infty} h(\tau) e^{-s\tau} d\tau = \int_0^1 e^{-s\tau} d\tau = \begin{cases} \frac{1}{s}(1 - e^{-s}) & s \neq 0 \\ 1 & s = 0 \end{cases} \quad (12.19)$$

Theorem 8

If $x(t) = e^{st}$ and H_s exists, $y(t) = H(s)e^{st}$

This can be interpreted in two ways

1. Same constant exponential as input e^{st} (eigenfunction of LTI system)
2. Scaled by constant $H(s)$ (eigenvalue)

Corollary

If the input to an LTI system with an impulse response $h(t)$ is

$$x(t) = \sum_k a_k e^{s_k t} \quad (12.20)$$

then it's output is given by

$$y(t) = \sum_k a_k H_{s_k} e^{s_k t} \quad (12.21)$$

Since almost all periodic signals can be expressed as a sum of harmonically related complex exponentials, i.e.

$$x(t) = \sum_{k=-\infty}^{\infty} a_k e^{jk\omega_t} = \sum_{k=-\infty}^{\infty} a_k e^{j\frac{2\pi k}{T}t} \quad (12.22)$$

Where the k denotes the $|k|$ th harmonic component and a_0 being the constant component²⁵

²⁵ In circuit analysis this is the constant component

SUBSECTION 12.4

Lecture 15: Which periodic signals have a Fourier representation?

We have the following expressions for synthesis and analysis, respectively

$$x(t) = \sum_{k=-\infty}^{\infty} a_k e^{kj\omega t} \quad (12.23)$$

$$a_k = \frac{1}{T} \int_T x(t) e^{-kj\omega t} dt \quad (12.24)$$

SUBSECTION 12.5

Fourier series representation of continuous periodic systems

$$x(t) = \cos \omega_o t \Leftrightarrow e^{j\omega_o t} \quad (12.25)$$

Both of the signals are periodic with fundamental frequency ω_o with associated harmonically related complex exponentials

$$\Phi_k(t) = e^{jk\omega_o t}, \quad k = 0, \pm 1, \pm 2 \quad (12.26)$$

TABLE 4.1 PROPERTIES OF THE FOURIER TRANSFORM

Section	Property	Aperiodic signal	Fourier transform
		$x(t)$	$X(j\omega)$
		$y(t)$	$Y(j\omega)$
4.3.1	Linearity	$ax(t) + by(t)$	$aX(j\omega) + bY(j\omega)$
4.3.2	Time Shifting	$x(t - t_0)$	$e^{-j\omega_0 t} X(j\omega)$
4.3.6	Frequency Shifting	$e^{j\omega_0 t} x(t)$	$X(j(\omega - \omega_0))$
4.3.3	Conjugation	$x^*(t)$	$X^*(-j\omega)$
4.3.5	Time Reversal	$x(-t)$	$X(-j\omega)$
4.3.5	Time and Frequency Scaling	$x(at)$	$\frac{1}{ a } X\left(\frac{j\omega}{a}\right)$
4.4	Convolution	$x(t) * y(t)$	$X(j\omega)Y(j\omega)$
4.5	Multiplication	$x(t)y(t)$	$\frac{1}{2\pi} \int_{-\infty}^{+\infty} X(j\theta)Y(j(\omega - \theta)) d\theta$
4.3.4	Differentiation in Time	$\frac{d}{dt} x(t)$	$j\omega X(j\omega)$
4.3.4	Integration	$\int_{-\infty}^t x(t) dt$	$\frac{1}{j\omega} X(j\omega) + \pi X(0)\delta(\omega)$
4.3.6	Differentiation in Frequency	$tx(t)$	$j \frac{d}{d\omega} X(j\omega)$
4.3.3	Conjugate Symmetry for Real Signals	$x(t)$ real	$\begin{cases} X(j\omega) = X^*(-j\omega) \\ \Re\{X(j\omega)\} = \Re\{X(-j\omega)\} \\ \Im\{X(j\omega)\} = -\Im\{X(-j\omega)\} \\ X(j\omega) = X(-j\omega) \\ \angle X(j\omega) = -\angle X(-j\omega) \end{cases}$
4.3.3	Symmetry for Real and Even Signals	$x(t)$ real and even	$X(j\omega)$ real and even
4.3.3	Symmetry for Real and Odd Signals	$x(t)$ real and odd	$X(j\omega)$ purely imaginary and odd
4.3.3	Even-Odd Decomposition for Real Signals	$x_e(t) = \Re\{x(t)\}$ [x(t) real] $x_o(t) = \Im\{x(t)\}$ [x(t) real]	$\Re\{X(j\omega)\}$ $j\Im\{X(j\omega)\}$
4.3.7	Parseval's Relation for Aperiodic Signals	$\int_{-\infty}^{+\infty} x(t) ^2 dt = \frac{1}{2\pi} \int_{-\infty}^{+\infty} X(j\omega) ^2 d\omega$	

Figure 47. a pretty good table

PART

IV

ECE358: Foundations of Computing

SECTION 13

Taught by Prof. Shurui Zhou

Admin and Preliminary Maths

SUBSECTION 13.1

Lecture 1

Topics covered will include:

- Graphs, trees
- Bunch of sorts
- Fancy search trees; red-black, splay, etc
- DP, Greedy
- Min span tree, single source shortest paths
- Maximum flow
- NP Completeness, theory of computation
- Blockchains??
- Θ

Solutions will be posted on the window of SF2001. Walk there and take a picture.

13.1.1 Mark Breakdown

Table 3. Mark Breakdown

Homework x 5	25
Midterm (Open book)	35
Final (Open book)	40

SUBSECTION 13.2

Complexities

This lecture we talked about big O notation. For notes on this refer to my tutorial notes for ESC180, ESC190: <https://github.com/ihasdapie/teaching/>

Definition 41 Big O notation (upper bound)

$g(n)$ is an asymptotic upper bound for $f(n)$ if:

$$O(g(n)) = \{f(n) : \exists c, n_0 \text{ s.t. } 0 \leq f(n) \leq c \cdot g(n), \forall n \geq n_0\} \quad (13.1)$$

PROOF | **What is the big-O of $n!$?**

$$n! \leq n \cdot n \cdot n \cdot n \dots n = n^n \Rightarrow n! \in O(n^n) \quad (13.2)$$

□

Definition 42

Big Ω notation (lower bound)

$h(n)$ is an asymptotic lower bound for $f(n)$ if:

$$\Omega(h(n)) = \{f(n) : \exists c, n_0 > 0 \text{ s.t. } 0 \leq c \cdot h(n) \leq f(n), \forall n \geq n_0\} \quad (13.3)$$

PROOF | **Find Θ for $f(n) \sum_i^n i$.**

For this we will employ a technique for the proof where we take the right half of the function, i.e. from $\frac{n}{2} \dots n$ and then find the bound

$$\begin{aligned} f(n) &= 1 + 2 + 3 \dots + n \\ &\geq \left\lceil \frac{n}{2} \right\rceil + \left(\left\lceil \frac{n}{2} \right\rceil + 1 \right) + \left(\left\lceil \frac{n}{2} \right\rceil + 2 \right) + \dots + n \quad n/2 \text{ times} \\ &\geq \left\lceil \frac{n}{2} \right\rceil + \left\lceil \frac{n}{2} \right\rceil + \left\lceil \frac{n}{2} \right\rceil + \dots + \left\lceil \frac{n}{2} \right\rceil \\ &\geq \frac{n}{2} \cdot \frac{n}{2} \\ &= \frac{n^2}{4} \end{aligned} \quad (13.4)$$

And therefore for $c = \frac{1}{4}$ and $n = 1$, $f(n) \in \Theta(n^2)$

□

Definition 43

Big Θ notation (asymptotically tight bound)

$$\Theta(g(n)) = \{f(n) : \exists c_1, c_2, n_0 \text{ s.t. } 0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n), \forall n \geq n_0\} \quad (13.5)$$

PROOF | Prove that

$$\sum_{j=1}^n j^k = \Theta(n^{k+1}) \quad (13.6)$$

First, prove $O(f(n)) = O(n^{k+1})$

$$\begin{aligned} f(n) &= \sum_{j=1}^n j^k = 1^k + 2^k + \dots + n^k \\ &\leq n^k + n^k + \dots + n^k \\ &= n^{k+1} \end{aligned} \quad (13.7)$$

Next, prove $\Omega(f(n)) = \Omega(n^{k+1})$

$$\begin{aligned}
 f(n) &= \sum_{j=1}^n i^k = 1^k + 2^k + \dots n^k \\
 &= n^k + (n_1)^k + \dots 2^k + 1^k = \sum_{i=1}^n (n - i + 1)^k \\
 &\geq \frac{n^k}{2} * n \geq \frac{n^{k+1}}{2^k} = \Omega(n^{k+1})
 \end{aligned} \tag{13.8}$$

Therefore $f(n) = \Theta(n^{k+1})$ □

Note that we may not always find both a tight upper and lower bound so not all functions have a tight asymptotic bound.

Theorem 9

Properties of asymptotes:

Note: \wedge means AND

Transitivity²⁶

$$(f(n) = \Theta(g(n)) \wedge g(n) = \Theta(h(n))) \Rightarrow f(n) = \Theta(h(n)) \tag{13.9}$$

Reflexivity²⁷

$$f(n) = \Theta(f(n)) \tag{13.10}$$

Symmetry

$$f(n) = \Theta(g(n)) \iff g(n) = \Theta(f(n)) \tag{13.11}$$

Transpose Symmetry

$$\begin{aligned}
 f(n) = O(g(n)) &\iff g(n) = \Omega(f(n)) \\
 f(n) = o(g(n)) &\iff g(n) = \omega(f(n))
 \end{aligned} \tag{13.12}$$

²⁶ The following applies to O, Θ, o, ω

²⁷ The following applies to O, Θ

Runtime complexity bounds can sometimes be used to compare functions. For example, $f(n) = O(g(n))$ is like $a \leq b$

- $O \approx \leq$
- $\Omega \approx \geq$
- $\Theta \approx \approx$
- $o \approx <;$ an upper bound that is **not** asymptotically tight
- $\omega >$ a lower bound that is **not** asymptotically tight

Note that there is no trichotomy; unlike real numbers where we can just do $a < b$, etc, we may not always be able to compare functions.

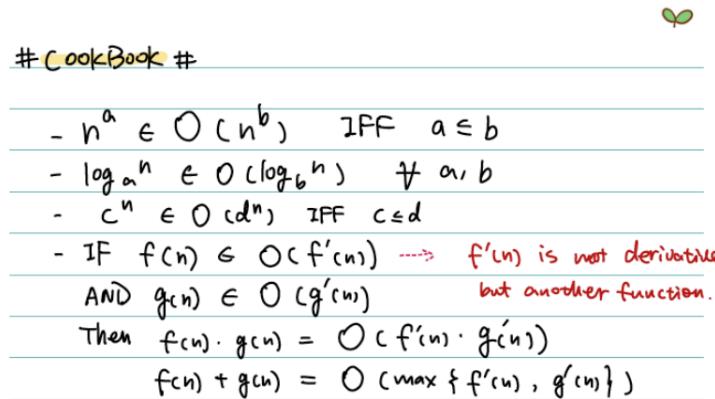


Figure 48. Complexity Cookbook

SUBSECTION 13.3

Lecture 3: Logs & Sums

13.3.1 Functional Iteration

$f^{(i)}(n)$ denotes a function iteratively applied i times to value n .
 For example, a function may be defined as:

$$f^{(i)}(n) = \begin{cases} f(n) & \text{if } i = 0 \\ f(f^{(i-1)}(n)) & \text{if } i > 0 \end{cases} \quad (13.14)$$

Given (13.14) we see that

1. $f(n) = 2n$
2. $f^{(2)}(n) = f(2n) = 2^2 n$
3. $f^{(3)}(n) = f(f^{(2)}(n)) = 2^3 n$
4. $f^{(i)}(n) = 2^i n$

As an exercise we may look at an iterated logarithm function, 'log star'

$$\lg^*(n) = \min\{i \geq 0 : \lg^{(i)} n \leq 1\} \quad (13.15)$$

This describes the number of times we can iterate $\log(n)$ until it gets to 1 or smaller.

- $\log^* 2 = 1$
- $\log^* 4 = 2 = \log^* 2^2 = 1 + \log^* 2 = 2$
- for practical reasons \log^* doesn't really get bigger than 5. This is one of the slowest growing functions around.

Summations & Series

PROOF | Proof for a finite geometric sum:

$$\begin{aligned}
 \sum_{k=0}^n x^k &= S \\
 S &= 1 + x + x^2 \dots x^n \\
 xS &= x + x^2 + x^3 \dots x^{n+1} \\
 S &= \frac{1 - x^{n+1}}{1 - x}
 \end{aligned} \tag{13.16}$$

□

$$\sum_{i=1}^{\infty} x^i = \frac{1}{1 - x} \quad \text{if } |x| < 1 \tag{13.17}$$

$$\sum_{k=0}^{\infty} kx^k = \frac{x}{(1 - x)^2} \quad \text{if } |x| < 1 \tag{13.18}$$

PROOF Begin by differentiating both sides over x

$$\sum_{k=0}^{\infty} x^k = \frac{1}{(1 - x)} \quad \text{if } |x| < 1 \tag{13.19}$$

$$\sum_{k=0}^{\infty} kx^{k-1} = \frac{1}{(1 - x)^2} \tag{13.20}$$

And then multiply both sides by x , therefore (13.18) follows. □

Telescoping Series

$$\sum_{k=1}^n a_k - a_{k-1} = a_n - a_0 \tag{13.21}$$

PROOF Write it out and cancel out terms

$$(a_1 - a_0) + (a_2 - a_1) \dots (a_n - a_{n-1}) = a_n - a_0 \tag{13.22}$$

Therefore the sum telescopes □

Another telescoping series may be proved similarly:

$$\sum_{k=1}^{n-1} \frac{1}{k(k+1)} \xrightarrow{\text{math}} \sum_{k=1}^{n-1} \left(\frac{1}{k} - \frac{1}{k+1} \right) = \left(1 - \frac{1}{2} \right) + \dots + \left(\frac{1}{n-1} - \frac{1}{n} \right) = a_o - a_n \tag{13.23}$$

SUBSECTION 13.4

Lecture 4: Induction & Contradiction

13.4.1 Induction

The general steps for proving a statement by induction are:

1. Basis
2. Hypothesis
3. Inductive step

I.e. if the basis holds for some i , i.e. $0, 1, 2, 3, 12, \dots$ AND if we assume that the hypothesis holds for an arbitrary number k , then we just need to prove that the inductive step follows, or that $P(n + 1)$ holds.

Example | Prove that $P(n) = 1 + 2 + 3 + \dots + n = \frac{n(n+1)}{2}$

PROOF | 1. Basis: $P(0) = 0 = \frac{0(0+1)}{2}$

2. Hypothesis (assume that it is true): $P(k) = \frac{k(k+1)}{2}$

3. Inductive step (need to prove $P(k + 1) = \frac{(k+1)(k+2)}{2}$): $P(k + 1) = \underbrace{1 + 2 + \dots + n}_{\frac{n(n+1)}{2}} + (n + 1) = \dots = \frac{n^2 + 3n + 2}{2} = \frac{(n+1)(n+2)}{2}$

□

Example | Show that for any finite set S , the power set 2^S has $2^{|S|}$ elements (that is, there are $2^{|S|}$ distinct subsets of S)

PROOF | 1. Basis:

$$n = 0, |S| = 0, |2^S| = 1 = 2^0 \quad (13.24)$$

$$n = 1, |S| = 1, |2^S| = 2 = 2^1 \quad (13.25)$$

2. Hypothesis: Assume that 2^S has 2^n elements when $|S| = n$

3. Inductive step: need to prove that when $|S| = n + 1$, $|2^S| = 2^{n+1}$

Let $B = S \setminus \{a\}$ for some $a \in S$. Now there are two types of subsets of S ; those that include a and those who do not include a

For subsets that do *not* include a , $|2^B| = 2^{|B|} = 2^n$, by the hypothesis.

For subsets that do include a , these sets are of size $2^B \cup \{a\}$, which is 2^n .

Therefore the total number of subsets of S is $2^n + 2^n = 2^{n+1}$, as desired.

□

The power set of a set S is the set of all subsets of S

The same kind of argument can be applied to problems such as the **Towers of Hanoi** and the tiling problem.

13.4.2 Contradiction

1. Assume the theorem is false
2. Show that the assumption is false (leads to a contradiction)
 - Therefore the theorem is true

Example | Prove that $\sqrt{2}$ is irrational

PROOF

Assume that $\sqrt{2}$ is rational.
Therefore we can write $\sqrt{2}$ as

$$\sqrt{2} = \frac{a}{b} \quad (13.26)$$

Where a, b **have no common factors**.

We can square both sides

$$2 = \frac{a^2}{b^2} \rightarrow a^2 = 2b^2 \quad (13.27)$$

Therefore a^2 is even.

Let $a = 2c$

$$2^2c^2 = 2b^2 \rightarrow b^2 = 2c^2 \quad (13.28)$$

Therefore b is even as well.

This results in a contradiction since we assumed that a, b have no common factors, but our analysis shows that both would have to be even (and share a common factor of 2). \square

SUBSECTION 13.5

Lecture 5: Recurrences

Many recursive algorithms can be thought of as a divide-and-conquer approach where we break the problem into subproblems that are similar to the original but smaller in size, solve them recursively, then combine them to create a solution to the original problem.

Definition 44

A recurrence is a function defined in terms of:

- 1+ base cases
- Itself, with smaller arguments

For example, finding a Fibonacci number is a recurrence;

$$T(n) = T(n-1) + T(n-2) \text{ with some base cases.}$$

Example

Mergesort

Sorting $[3, 1, 7, 5]$

1. Divide: break into partitions: $[[3, 1], [7, 5]]$
2. Sort partitions: $[[1, 3], [5, 7]]$
3. Create result array
4. Compare: have two pointers to front of array
 - compare 1, 5. 1 is smaller; $result = [] \leftarrow 1$
 - Move ptr to left array (1, 3) ahead one. Compare 3, 5. 3 is smaller, so $result = [3] \leftarrow 3$
 - One of the arrays is now empty so we can just append the rest
5. $result = [1, 3, 5, 7]$

Definition 45

Pseudocode for mergesort is given by:

```

1     mergesort(A, p, r)
2         if p < r
3             q = [(p+r)/2]
4             mergesort(A, p, q) // N/2
5             mergesort(A, q+1, r) // N/2
6             merge(A, p, q, r) // merge the sorted
    →    subarrays

```

This mergesort partitions in half each time²⁸.

In the worst case we will compare $N - 1$ times, so $O(N)$ worst case.

Proving that merge sort is $\Omega(N)$

²⁸binary partitioning(?)

Here we're discussing not time complexity but rather the number of times we call mergesort

How much time does MergeSort take?

The time of mergesort is defined recursively as:

$$T(N) = \begin{cases} O(1) & n = 1 \\ T(N) = 2T(\frac{N}{2}) + \Theta(N) & \text{otherwise} \end{cases} \quad (13.29)$$

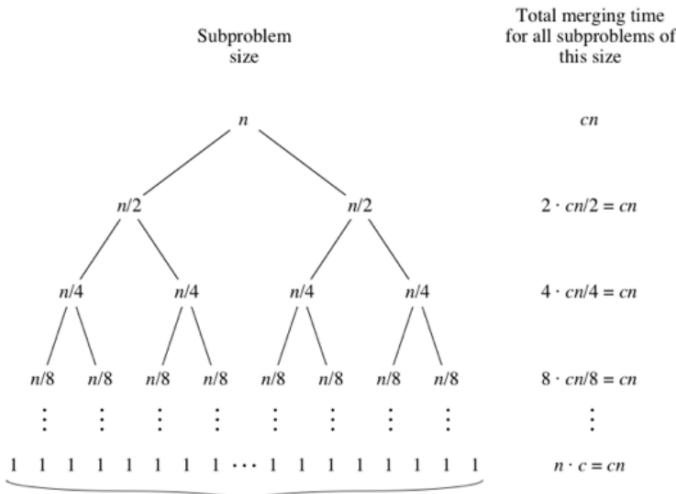
More generally we can find the runtime of a recurrence algorithm via

Note that $\Theta(N)$ is for the merge operation

- Recurrence trees
- Substitution
- Master theorem

13.5.1 Recurrence Trees

Example Recurrence trees can be used to find the time complexity of mergesort.



The height of the tree is $\log N$

The total cost is the total cost per level times the number of levels, which is

$$N \cdot \log N \quad (13.30)$$

So the complexity is $O(N \log N)$

13.5.2 Substitution

1. Guess the answer
2. Apply induction

Example Determine an asymptotic upper bound on $T(n) = 2T(\lfloor \frac{n}{2} \rfloor) + \Theta(n)$

PROOF This expression can be simplified since we don't care too much about floors or ceilings for asymptotic behaviour.

$$T(n) = 2T\left(\frac{n}{2}\right) + N \quad (13.31)$$

Let's guess that the upper bound is $O(n \log n)$

Then, we need to prove that $T(n) < C \cdot n \log n$ for some $C > 0$. Let's apply induction.

1. Basis: this is tricky since if $n = 1$ we end up with $T(1) \leq C \cdot 1 \cdot \log 1 = 0$ which cannot hold since that would just not make sense. Instead, observe that $T(1) = 1$, $T(2) = 2T(1) + 2 = 4$, $T(3) = 2T(1) + 3 = 5$, $T(4) = 2T(2) + 4 \dots$. So $T(n)$ is therefore independent of $T(1)$, so we can use two bases, $T(2), T(3)$. Since $T(2) \leq C * 2 \log 2 = 2C$, $T(3) \leq C \cdot \log 3$
2. Hypothesis: Assume that the upper bound holds for all possible $m < n$, let $m = \lfloor \frac{n}{2} \rfloor$. This yields $T(\lfloor \frac{n}{2} \rfloor) \leq C \cdot \lfloor \frac{n}{2} \rfloor \cdot \log \lfloor \frac{n}{2} \rfloor$
3. Inductive step: substitute hypothesis into recurrence yields

$$T(N) \leq C \cdot (C \cdot \left\lfloor \frac{N}{2} \right\rfloor \cdot \log \left\lfloor \frac{N}{2} \right\rfloor + N) = cN \log N - (1-c)N \leq Cn \log n \quad (13.32)$$

□

A few pitfalls to avoid is guessing $T(n) = O(n) = c \cdot n$ and so forth we would get

$$T(N) \leq 2C \cdot \left\lfloor \frac{n}{2} \right\rfloor + n = cn + n = (c+1)n \quad (13.33)$$

This would be wrong since we cannot change the constant to $c+1$; we have to prove it with exactly the hypothesis given.

13.5.3 Master Theorem

Definition 46

The master method applies to recurrences of the form

$$T(n) = aT\left(\frac{n}{b}\right) + f(n) \quad \text{where } a \geq 1, b > 1, f \text{ asymptotically positive} \quad (13.34)$$

It distinguishes 3 common cases b comparing $f(n)$ with $n^{\log_b a}$

1. If $f(n) = O(n^{\log_b a - \varepsilon})$ for some $\varepsilon > 0$, then $T(n) = \Theta(n^{\log_b a})$
2. If $f(n) = \Theta(n^{\log_b a})$, then $T(n) = \Theta(n^{\log_b a} \log n)$
3. If $f(n) = \Omega(n^{\log_b a + \varepsilon})$ for some $\varepsilon > 0$ and $af\left(\frac{n}{b}\right) \leq cf(n)$ for some $c < 1$, then then $T(n) = \Theta(f(n))$

Proof is out of scope for the course

There are a few technicalities to be aware of. In each example we compare $f(n)$ with $F = n^{\log_b a}$ and take the larger of each as the solution to the recurrence. For the first case we note that $f(n)$ must be *polynomially* smaller than F ; i.e. it must be asymptotically smaller than F by some factor of n^ε . In the third case $f(n)$ must be greater than F as well as being polynomially larger and satisfy the regularity condition $af(\frac{n}{b}) \leq cf(n)$. There are areas where the master theorem does not cover, for example a gap between cases 1, 2 where $f(n) > F$ but is not polynomially larger. If $f(n)$ falls in one of these gaps or the regularity condition does not hold, the master method cannot be used to solve the recurrence.

Example What is the closed form of $T(n) = T(\frac{2n}{3}) + 1$?

PROOF $a = 1, b = 2/3, f(n) = 1$.

$$\log_b a = \log_{\frac{2}{3}} 1 = 0 \quad (13.35)$$

$$f(n) = \Theta(n^0) \quad (13.36)$$

So

$$T(n) = O \log(n) \quad (13.37)$$

□

SECTION 14

Graphs

Definition 47

A **graph** is a data structure comprised from set of vertices V and a set of edges E , where each edge connects a pair of vertices. A **directed graph (digraph)** is a graph where edges E have a *direction*, i.e. an edge (u, v) is different from (v, u) . Conversely, an **undirected graph** is a graph where edges E do not have a direction.

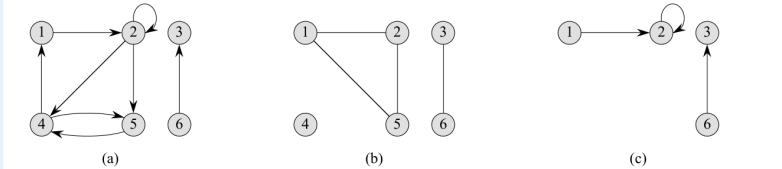


Figure 49. (a) directed graph, (b) undirected graph, (c) a subgraph of (a)

Some conventions:

- Edges are denoted by (u, v) ²⁹ where $u, v \in V$
- If (u, v) is a edge in a directed graph, then (u, v) is incident from or leaves u , and is incident to or enters v
- If (u, v) is a edge a graph, then u, v are adjacent
- The **degree** of a vertex is the number of edges incident to it
- A **path** is a sequence of vertices (v_0, v_1, \dots, v_k) from vertex to another such that each vertex is incident³⁰ to the ones prior and after.
 - If there exists a path from a to b then b is **reachable** from a and a
 - A path is **simple** if no vertex is repeated

²⁹not $\{u, v\}$

³⁰with the exception of start/end vertices

- A path forms a cycle if the first and last vertices are the same
- A directed graph with no self-loops is **simple**
- A graph with no simple cycles is acyclic
- An undirected graph is **connected** if there exists a path between any two vertices
- A directed graph is **strongly connected** if every vertex is reachable from every other vertex
- Two graphs V, V' are **isomorphic** if there exists a bijection³¹ between the vertices of the two graphs such that the edges are preserved
- Given graph $G, G' = (V', E')$ is a **subgraph** of G if $V' \subseteq V$ and $E' \subseteq E$
- Given a set $V' \subseteq V$, the subgraph of G **induced** by V' is the graph $G' = (V', E')$ where $E' = \{(u, v) \in E \mid u, v \in V'\}$
- Given an undirected graph, the **directed version** of G is $G = (V, E')$ where $(u, v) \in E'$ if and only if $(u, v) \in E$. In other words we replace all undirected edges and replace them with their directed counterpart.
- The corollary can be applied to a directed graph to get the **undirected version** of G .
- A **neighbor** of u in a directed graph is any vertex v such that $(u, v) \in E$ where E is the set of edges for the undirected counterpart of the graph
- A **complete** graph is a graph where every pair of vertices are connected by an edge
- A **bipartite graph** is an undirected graph G in which it's V can be partitioned into two disjoint sets V_1, V_2 such that every edge $(u, v) \in E$ connects a vertex in V_1 to a vertex in V_2 or vice-versa
- An acyclic undirected graph is a **forest**
- A connected acyclic undirected graph is a **free tree**.
 - A directed acyclic graph is often termed a DAG
- A multi-graph is a graph where edges can be repeated and self-loops are allowed
- A hyper-graph is a graph where edges can connect more than two vertices
- The **contraction** of an undirected graph G by an edge $e = (u, v)$ is a graph G' where $V' = V - \{u, v\} \cup \{x\}$, where x is a new vertex. Then, for each edge connected to u, v the edges are deleted and then reconstructed with x , effectively 'contracting' u, v into a single vertex

³¹ $f : V \rightarrow V'$, i.e. we can relabel the vertices of V to be those of V' and the two graphs would be identical

In code graphs are commonly represented as adjacency lists or adjacency matrices. This was covered in ESC190, but for reference:

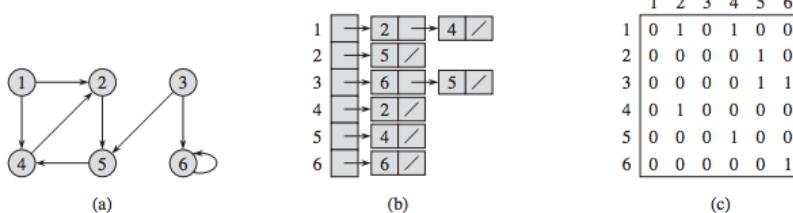


Table 4. Time complexities of graph representations

	adjacency list	matrix
Time	$O(n)$	$O(1)$
Memory	$O(E)$	$O(n^2)$

SECTION 15

Trees

A **tree** is a common and useful subset of graphs

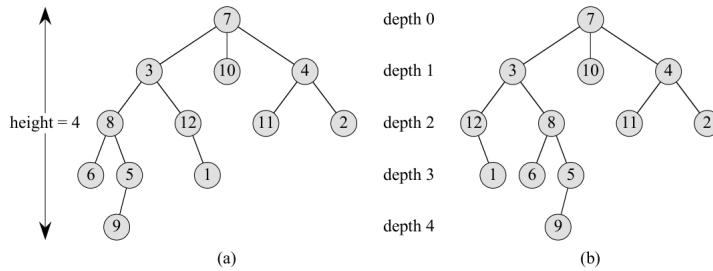


Figure B.6 Rooted and ordered trees. (a) A rooted tree with height 4. The tree is drawn in a standard way: the root (node 7) is at the top, its children (nodes with depth 1) are beneath it, their children (nodes with depth 2) are beneath them, and so forth. If the tree is ordered, the relative left-to-right order of the children of a node matters; otherwise it doesn't. (b) Another rooted tree. As a rooted tree, it is identical to the tree in (a), but as an ordered tree it is different, since the children of node 3 appear in a different order.

Definition 48

A tree is a common subset of graphs, i.e. ones that are **connected**, **acyclic**, and **undirected**. This gives a few useful properties, i.e. the existence of a *root* node, the parent-child relationship, and the existence of a unique path between any two nodes.

Some conventions for trees

- Depth of node: length from root to node
- Height length of longest path from node to leaf
- Degree of node: number of children. Binary trees have degree 2, n-ary tree has degree n

Theorem 10

All of the following statements are equivalent for a tree $T = (V, E)$:

1. $\forall v \in V, v$ is a tree; all nodes in a tree are trees unto themselves.
2. Every two nodes are connected by a unique path
3. T is connected by if any edge is removed the resulting graph is disconnected T is connected and $|E| = |V| - 1$ T is acyclic and $|E| = |V| - 1$ T is connected but if a edge is added the resulting graph has a cycle

SECTION 16

Lecture 7, 8: Probability and Counting

Most of the probability stuff is review from ECE286 so I'll be omitting most notes.

Definition 49

Probability distribution $Pr \{ \cdot \}$: mapping from events of S to real numbers where

1. $Pr \{ \emptyset \} = 0$
2. $Pr \{ A \} \geq 0 \forall A \in S$
3. $Pr \{ S \} = 1$
4. $Pr \{ A \cup B \} = Pr \{ A \} + Pr \{ B \} \forall A, B \in S$

For any two events A, B we can define the triangle inequality

Definition 50

$$Pr \{ A \cup B \} \leq Pr \{ A \} + Pr \{ B \}$$

The complement of an event A is $\bar{A} = S - A$, and the probability is $Pr \{ \bar{A} \} = 1 - Pr \{ A \}$

Definition 51

Baye's theorem

$$Pr \{ A \mid B \} = \frac{Pr \{ B \mid A \} Pr \{ A \}}{Pr \{ B \}} = \frac{Pr \{ A \} Pr \{ B \mid A \}}{Pr \{ A \} Pr \{ B \mid A \} + Pr \{ \bar{A} \} Pr \{ B \mid \bar{A} \}} \quad (16.1)$$

The expected value of a random variable is

$$E[X] = \int_{-\infty}^{\infty} x Pr \{ X = x \} dx \quad (16.2)$$

And in the discrete case,

$$E[X] = \sum_{x \in \mathbb{Z}} x Pr \{ X = x \} \quad (16.3)$$

The variance is

$$Var[X] = E \{ (X - E[X])^2 \} = E[X^2] - E^2[X] \quad (16.4)$$

The **standard deviation**, σ , is the positive square root of the variance.

SECTION 17

Sorting

SUBSECTION 17.1

Heapsort

Heapsort is a sorting algorithm that runs in $O(n \log(n))$ and sorts its elements in place by leveraging a *heap*.

This should be review from ESC190

Definition 52

A heap is a data structure that satisfies the *heap property*. For a max heap, the value of a mode is at most the value of its parent and vice-versa for a min-heap. Heaps are commonly used to implement a priority queue because it offers a `pop{min, max}` operation in $O(1)$ time and insertion in $O(\log(n))$ time. Another reason for its use is its ability to be represented as an array with `Parent(i) = floor(i/2)`, `Left(i) = 2i`, `Right(i) = 2i+1`

The `MAXHEAPIFY` method is used to coerce heaps back into fulfilling the heap property by 'bubbling down' nodes until the heap property is satisfied. The intuition for this is that as

long as the heap property is not satisfied³², node i is exchanged for the largest of its' children. Since the heap property is now satisfied for i but not necessarily its' children, this heapify operation is then recursively called on the larger of i 's children.

³² which we can tell by looking at the node's children

```

1  MAXHEAPIFY(A, i)
2      l = LEFT(i)
3      r = RIGHT(i)
4      if l <= A.size and A[l] > A[i]
5          largest = l
6      else largest = i
7      if r <= A.size and A[r] > A[largest]
8          largest = r
9      if largest != i
10         swap A[i], A[largest]
11         MAXHEAPIFY(A, largest)

```

The child sub-trees can have size of at most $\frac{2n}{3}$ and so we can describe MAXHEAPIFY by the recurrence relation

$$T(n) = \leq T\left(\frac{2n}{3}\right) + \Theta(1) \quad (17.1)$$

This heapify operation can then be used to build a max heap from an array by calling heapify on all nodes that are not leaves. This is done by starting at the last node that has children and calling heapify on it. This is repeated until the root is reached. It can likewise be used to sort an array.

```

1  HEAPSORT(A)
2  for i = A.length -> 2,
3      swap A[1], A[i]
4      A.heapsize = A.heapsize - 1
5      HEAPIFY(A, 1)

```

SUBSECTION 17.2

Placeholder for Quick, counting, selection, and radix sort

I'm going to skip these since these are fairly trivial and easily found online or in the textbook. I also didn't go to class.

SUBSECTION 17.3

BSTs

TODO: Notes on red-black trees, etc.

SECTION 18

Amortized Analysis

- **Average cost:** mean across all inputs

$$\frac{1}{T-0} \int_0^T c(x) dx \quad (18.1)$$

- **Expected cost:** expectation over all inputs w.r.t a probability distribution

$$\int_0^T p(x)c(d)dx \quad (18.2)$$

- **Amortized Cost:** average cost over a *sequence* of operations

$$\frac{1}{n} \sum_{i=1}^n c\left(\frac{1}{n}\right) \quad (18.3)$$

The aggregate analysis method is simple and is as follows

1. Given a function or operation $f(x)$ and a sequence $\{x_1, x_2, \dots, x_n\}$
2. Determine the total cost of the sequence $\sum_{i=1}^n c(x_i) \equiv T(n)$
3. The amortized cost is $\frac{T(n)}{n}$

The complicated but more powerful way to deal with this is the accounting method.

1. Declare a cost \hat{c} for each operation/function call
2. Describe a procedure for how the cost will be allocated
3. Assert a credit invariant, i.e. that the credit in an object must be greater than or equal to 0
4. Prove the credit invariant
5. Using the credit invariant argue why the credit never goes negative
6. The amortized cost is $O(\hat{c})$

Let's consider a contrived example.

Example

Consider a data structure that has a cost of

$$c(x) = \begin{cases} x & \text{if } x \text{ is a power of 2} \\ 1 & \text{otherwise} \end{cases} \quad (18.4)$$

Determine the amortized cost over the sequence $\{1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$

The aggregate method for solving this problem is trivial; simply compute the summation to find

$$\sum_{x=1}^n c(x) = \dots (2n-1) \leq 3n \quad (18.5)$$

What's more interesting is how we may apply the accounting method

1. Let cost $\hat{c} = 3$

2. If x is not a power of 2, uses 1 to pay for the operation and store the remaining 2
3. If x is a power of 2, store 2 in credit and use x from previous iterations to pay for the operation
4. (Steps 3, 4, 5) Prove the credit invariant. When $x = 2^m$ all elements $\in (2^{m-1}, 2^m]$ have 2 stored. The math can be a little handwavy but just draw out the ranges and show base cases etc.
5. And the amortized cost is $O(\hat{c}) = O(3) = O(1)$

What amortized cost is really telling us is that the cost of all operations as a whole is never larger than some value, though on the individual level it may fluctuate a little bit.

$$\frac{1}{n} \sum_{i=1}^n \hat{c} \geq \frac{1}{n} \sum_{i=1}^n c(x_i) \quad \forall n \quad (18.6)$$

We can think of this is that the cumulative work stored in the data structure (amortized) being an upper bound on the actual cost.

$$\frac{1}{n} \sum_{i=1}^n (\hat{c} - c(x_i)) \geq 0 \quad \forall n \quad (18.7)$$

The credit invariant is used to prove the above inequality.

Example Consider a FIFO queue implemented using two stacks S_{in}, S_{out} . The stacks have the operations $\text{Push}(Q, x)$ (pushes x onto S_{in}), $\text{copy}(Q)$ (pops all elements from S_{in} and pushes onto S_{out}), and $\text{pop}(Q)$ (calls copy and then pops off S_{out})

If $c(x_i) \leq \hat{c}$ then we good and can store anything left over. If $c(x_i) > \hat{c}$ then we need to use credit from previous iterations. And the credit invariant is used to argue that we never use too much credit.

operation	actual cost	amortized cost
push	$O(1)$	$O(1)$
copy	$O(n)$	$O(1)$
pop	$O(n)$	$O(1)$

- Push: 1 pays for push onto S_{in}
- 2 pays for the eventual copy
- 1 pays for the pop from S_{out}

The credit invariant is therefore

1. Elements in S_{in} have 3 in stored credit
2. Elements in S_{out} have 1 in stored credit

The amortized cost of each operation is $O(1)$ for all

Example Consider a sequence of stack operations on a stack which never exceeds k elements. After k operations, we copy the stack

PROOF

1. Let cost = 2. Push: pay 1, store 1. Pop: pay 1, store 1, Copy: pay k
2. credit invariant holds in stack because the number of credit stored on the stack equals the number of operations since the last copy. Since after each operation +1 is stored, then since the stack never has more than k elements after k operations we will have k stored and can therefore pay for the copy

□

SUBSECTION 18.1

Skewed Heaps

Consider a *skewed heap* data structure which implements a `SkewHeapMerge` operation for merging heap-ordered trees by merging the rightmost paths of two trees without keeping any explicit balance conditions. Good performance of skewed heaps is guaranteed by a rebalancing step which swaps all the left and right children of nodes along the rightmost path (except for the last one). Since there is no balance condition the trees won't have $O(\log n)$ worst case performance but they do have good amortized performance.

`SKEWHEAPMERGE(h, h')`

```

1
2  if  $h == NIL$ 
3    return  $NIL$ 
4  if  $h' == NIL$ 
5    return  $NIL$ 
6  if  $h.key \leq h'.key$ 
7    return SKEWHEAPMERGE( $h', h$ )
8   $h.right = \text{SKEWHEAPMERGE}(h.right, h')$ 
9  SWAP( $h.left, h.right$ )
10 return  $h$ 
```

Example

Show that a `SkewedMeapMerge`, `Insert`, and `DeleteMin` uses amortized $O(\log_2 n)$ time.

1. Allocate $\hat{c} = c \log(n)$ per operation (since we want to prove that it's $O(\log n)$ amortized)
2. Define $WT(x)$ as the number of nodes in the subtree rooted at x. Define a heavy node as a node with $WT(x) \geq WT(x.parent)$. Note that this doesn't include the root. A light node is exactly the opposite.
 - If all nodes on the rightmost path are light, then the cost of `SKEWHEAPMERGE` is minimized
 - Define a misplaced node as a heavy node on the rightmost path

Find the number of light nodes on any path from the root to the leaf.

SECTION 19

Minimum Spanning Trees

Definition 54

A minimum spanning tree is an acyclic subset of edges in a graph which connects all the vertices and has the minimum total weight³³

³³i.e. sum of edge cost

Two common algorithms for solving the minimum spanning tree problem are Kruskal's algorithm and Prim's algorithm. They are similar in that they both attempt to greedily add edges to the minimum spanning tree, but whereas Kruskal's algorithm grows a forest by adding the least weight edge in the graph that connects two distinct components, Prim's algorithm grows a single tree by adding the least-weight edge connecting the tree to a vertex outside of the tree.

Definition 55

Terminology

- **Cut** $(S, V - S)$ is a partition of V ³⁴
- A cut **respects** a set A if no edge crosses the cut
- A **light edge** is the edge crossing a cut with the minimum weight
- A **safe edge** is any light edge crossing a cut $S, V - S$ of G that respects the MST A

³⁴ see it as dividing V into two new sets $S, S - V$

MST-KRUSKAL(G, w)

```

1   $A = \{0\}$ 
2  for each vertex  $v \in G.V$ 
3    MAKE-SET( $v$ )
4    sort  $G.E$  by weight  $w$  in nondecreasing order
5    for each edge  $(u, v) \in G.E$ 
6      if  $\text{FIND-SET}(u) \neq \text{FIND-SET}(v)$ 
7         $A = A \cup \{(u, v)\}$ 
8        UNION( $u, v$ )
9  return  $A$ 
```

In summary Kruskal's algorithm adds edges (u, v) to the MST in non-decreasing order of weight, *if* (u, v) do not belong to the same tree (as to prevent creating a cycle). Note: The runtime of these algorithms can be improved using a priority queue on the edge weights. However if the graph is dense ($E = V^2$) an array can be better an array can be better

MST-PRIM(G, w, r)

```

1  for each vertex  $u \in G.V$ 
2     $u.\text{key} = \infty$ 
3     $u.\pi = \text{NIL}$ 
4   $Q = G.V$ 
5  while  $Q \neq \{0\}$ 
6     $u = \text{EXTRACT-MIN}(Q)$ 
7    for each vertex  $v \in G.\text{Adj}[u]$ 
8       $u.\text{key} = w(u, v)$ 
9       $u.\pi = u$ 
```

Prim's algorithm builds a MST set A that always forms a single tree. Each step adds a light edge to A that connects it to an isolated vertex. Since this only adds edges that are safe for A and A will eventually contain all vertices in V , A forms a minimum spanning tree.

SECTION 20

Shortest Paths

Common algorithms for solving the shortest path problem are Dijkstra's algorithm, Bellman-Ford's algorithm, and A*.

Dijkstra's algorithm and A* were discussed in ESC190 and such will not be discussed here. Bellman-Ford's algorithm, unlike Dijkstra, can handle negative weight edges³⁵

Definition 56

Relaxation

For each vertex $v \in V$, we maintain a *shortest path estimate* $v.d$ describing the upper bound on the weight of a shortest path from source to v , i.e. $s \rightsquigarrow v$. The initial state of $v.d$ will just be ∞

³⁵ though still fail to produce a shortest path as long as no negative-weight cycles are reachable, but can detect and report its existence.

INITIALIZE-SINGLE-SOURCE(G, s)

```

1  for each  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = NIL$ 
4       $s.d = 0$ 

```

relaxing an edge (u, v) means to check if the shortest path to v so far can be improved by going through u , and updating $v.d, v.\pi$ if so.

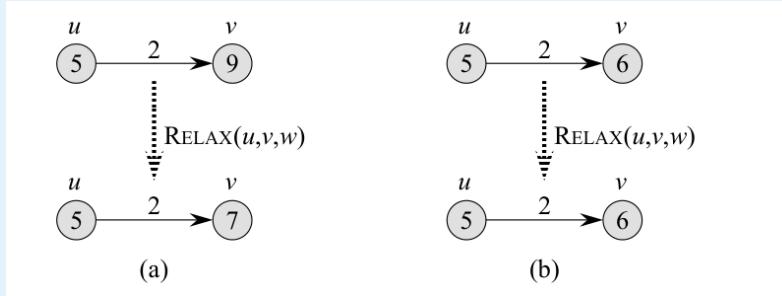


Figure 50. Relaxing an edge

RELAX(u, v, w)

```

1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 

```

An issue with relaxation-based algorithms is that order matters; continuing to relax in a graph with negative weight cycles can cause an infinite loop, and a poor choice of order can lead to many unnecessary relaxations. Bellman-Ford proposes to label edges E in a topological order $e_1 \dots e_n$ and then relax them in that order $|V| - 1$ times.

Definition 57

Bellman-Ford Algorithm

BELLMAN-FORD(G, w, s)

```

1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $1 \leq i \leq |G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE

```

All of the presented algorithms use only RELAX to change shortest-path estimates and predecessors. I'll also include the Dijkstra algo using RELAX later as well.

The algorithm makes $|V| - 1$ passes over the graph, attempting to relax each edge $|V| - 1$ times. Briefly we can see that the algorithm is correct by observing that the longest acyclic path in a graph with $|V|$ vertices has length $|V| - 1$. And the algorithm maintains the invariant that the i -th pass $v.d$ is at most the weight of every path from $s \rightsquigarrow v$ using at most i edges. So by the $|V| - 1$ -th pass each edge is maximally relaxed, i.e. $v.d = \delta(s, v)$ for all $v \in V$, where $\delta(u, v)$ denotes the actual shortest path weight

It then checks for a negative weight cycle reachable from the source, the proof for how this works can be found in CLRS Pg. 653³⁶

In scenarios where we are working with DAGs we can use a topological sort to order the edges and relax them in that order – which would result in a linear time algorithm $\Theta(V + E)$. There are no negative-weight cycles possible in a DAG so shortest paths will always be well-defined

³⁶But more or less the approach is to build a contradiction around how the sum of edge weights in a negative cycle is ≤ 0 and the conditions on $v.d$ being equal to $u.d + w(u, v)$

```
DAG-BELLMAN-FORD( $G, w, s$ )
1 sort  $G.V$  topologically
2 INITIALIZE-SINGLE-SOURCE( $G, s$ )
3 for each vertex  $u \in G.V$  taken in topological order
4   for each vertex  $v \in G.Adj[u]$ 
5     RELAX( $u, v, w$ )
```

Recall: topological ordering of DAG is a linear ordering of its vertices such that if $u \rightsquigarrow v$ then u comes before v in the ordering, i.e. a graph traversal such that each node v is visited only after all its dependencies are visited. An algorithm for topologically sorting a graph is just DFS with a colouring concept; WHITE for unvisited, GREY for visiting, and BLACK for visited. Time is simply a shared variable that increments for each call to DFS. The starting time of a vertex is the time when it is first discovered, and the finishing time is the time when it is colored black. The topological ordering is the reverse of the finishing times.

```
TOPOLOGICAL-SORT( $G$ )
1 call DFS( $G$ ) to compute finishing times  $f[v]$  for each vertex  $v$ 
2 as each vertex is finished, insert it onto the front of a linked list
3 return the linked list of vertices
```

Dijkstra's algorithm is then just a special case of the Bellman-Ford algorithm for a graph with no negative edge weights that relaxes greedily edges radially outwards from the source.

```
DIJKSTRA( $G, w, s$ )
1 INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  $S = \emptyset$ 
3  $Q = G.V$ 
4 while  $Q \neq \emptyset$ 
5    $u = \text{EXTRACT-MIN}(Q)$  (min on  $v.d$ )
6    $S = S \cup \{u\}$ 
7   for each  $v \in G.Adj[u]$ 
8     RELAX( $u, v, w$ )
```

SECTION 21

Splay Trees

Consider the weight dictionary problem where we want to maintain a set of keys $K = \{k_1, k_2 \dots k_n\}$ with associated values $V = \{v_1, v_2 \dots v_n\}$ and access frequencies $W = \{w_1, w_2, \dots w_n\}$ and we want to access high-access-frequency item faster than low-frequency

items without sacrificing efficiency on **INSERT**, **SEARCH**, and **DELETE**. If no insert/delete operations are needed across the tree during runtime one can build the static optimal search tree i.e. the one that minimizes $P = \sum_i w_i d_i$ where w_i is the weight of a path and d_i be the depth of node x_i ahead of time in $O(n^3)$ with DP. To motivate the study of splay trees, let's think of the case where we would want to remove from the BST and have different access frequencies during runtime; a more challenging problem. A data structure that achieves an amortized time within a constant multiple of the theoretic lower bound is the **splay tree**; a self-organizing data structure, or one that moves x closer to the entry point (in a BST the root) whenever x is accessed.

Definition 58

A splay tree is an ordered binary tree that satisfies

- every element in the left subtree of x is $\leq x$
- every element in the right subtree of x is $\geq x$

Across which we can perform the **SPLAY** procedure

```

SPLAY( $x$ )
  while  $x$  is not the root do
    let  $p$  be the parent of  $x$ 
    if  $p$  is the root then
      (1)      rotate at  $p$  and stop
                — now  $x$  is at the root
    else if  $x$  and  $p$  are both right children or left children then
      (2)      rotate at parent( $p$ )
      (2)      rotate at  $p$ 
                — now  $x$  is where its grandparent was
    else {  $x$  is a left child and  $p$  a right child, or vice versa }
      (3)      rotate at  $p$ 
      (3)      rotate at the new parent of  $x$  (its former grandparent)
                — now  $x$  is where its grandparent was
  
```

SPLAY is a procedure that moves x to the root of the tree by rotating x and its parent. There are a maximum of $\frac{h}{2}$ steps or h rotations if the height of node x is h (Each step in **SPLAY** is actually two rotations)

The other BST operations are therefore implemented as follows

- **LOCATE(x)**: search for x in the tree and then call **SPLAY** on x
- **INSERT(x)**: insert x into the tree as a leaf and then call **SPLAY** on the newly inserted x
- **DELETE(x)** : Delete an element as usual. Leaf nodes can simply be removed; internal nodes are replaced by their in-order successor and then the successor is removed. Then call **SPLAY** on the parent of the deleted node.

This splay operation causes items that are accessed most frequently to move closer to the root.

Definitions:

PROOF | Prove that **SPLAY** has an amortized cost of $O(\log(n))$.

- Let $WT(x)$ be the weight of node x , i.e the number of nodes in the subtree rooted at x including x itself.
- Let $RANK(x) = \log WT(x)$

Note that there is no explicit balance condition and as such the worst-case complexity can be $O(n)$ but the *amortized* complexity of each operation is $O(\log(n))$

Credit invariant: assert that there are always $\text{RANK}(x)$ credits on x , $\forall x \in \text{tree}$. 1 credit will be charged in each rotate operation and only $O(\log n)$ credits are needed in addition to maintain this invariant.

Claim: for each step in the splay operation the number of additional credits is at most $3(\text{NEW RANK}(x) - \text{OLD RANK}(x))$ and maybe one more for the last step; any other credits can be taken from the tree. Consider the rank of x as the splay operation proceeds; $\text{RANK}_0, \text{RANK}_1, \text{RANK}_2, \dots, \text{RANK}_k$. Therefore the total number of credits needed is

$$1 + 3(\text{RANK}_k - \text{RANK}_{k-1} \dots 3(\text{RANK}_1 - \text{RANK}_0)) \quad (21.1)$$

Cancelling out terms,

$$= 1 + 3(\text{RANK}_k - \text{RANK}_0) \quad (21.2)$$

Since $\text{RANK}_k = \log \text{WT}(x)$,

$$1 + 3(\text{RANK}_k - \text{RANK}_0) \leq 1 + 3 \log n \quad (21.3)$$

Now let's analyze it with respect to the three rotation cases in SPLAY;

1. if x 's parent is the root
2. if x and its parent are both right or left children
3. or if x is a left child and its parent is a right child or vice versa

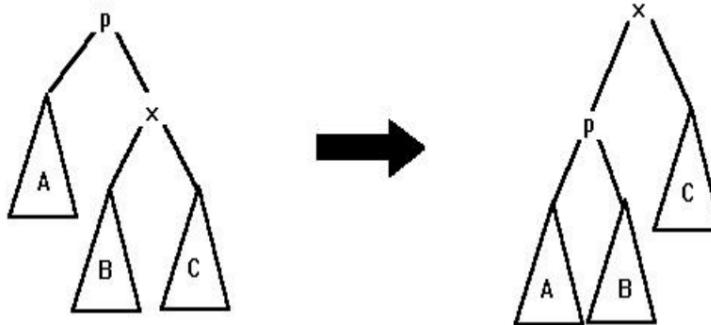


Figure 51. Case 1: x 's parent is the root

This becomes a case of just counting the new and old ranks of x and its parent p

$$\text{OLD RANK}(x) = \log(1 + |B| + |C|) \leq \text{NEW RANK}(x) = \log(2 + |B| + |C| + |A|) \quad (21.4)$$

$$\text{OLD RANK}(p) = \log(2 + |A| + |B| + |C|) \geq \text{NEW RANK}(p) = \log(1 + |A| + |C|) \quad (21.5)$$

p can only decrease in rank. x already has $\text{OLD RANK}(x)$ credits, so we only need $\text{NEW RANK}(x) - \text{OLD RANK}(x)$, which is only 1/3 of the allocated credits. One more credit may be used to pay for the rotation, which can be taken from the +1 taken for the last step.

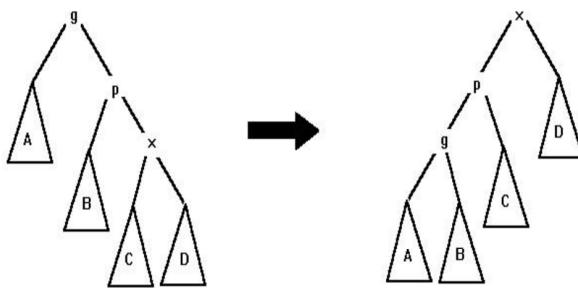


Figure 52. Case 2: x, p both left or right children

We have

$$\text{OLD RANK}(x) + \text{OLD RANK}(p) + \text{OLD RANK}(g) \quad (21.6)$$

Available and we need

$$\text{NEW RANK}(x) + \text{NEW RANK}(p) + \text{NEW RANK}(g) \quad (21.7)$$

Some math can be done to find that this is still bounded by our claim, i.e. $1 + 3(\text{NEW RANK}(x) - \text{OLD RANK}(x))$

Here are the screenshots notes (since I don't want to write it out)

Since $\text{NEW RANK}(x) = \text{OLD RANK}(g)$

$$\text{NEW RANK}(p) + \text{NEW RANK}(g) - \text{OLD RANK}(x) - \text{OLD RANK}(p) \quad (1)$$

credits to what we have. But we note that

$$\begin{aligned} \text{NEW RANK}(p) &\leq \text{NEW RANK}(x) \\ \text{NEW RANK}(g) &\leq \text{NEW RANK}(x) \\ \text{OLD RANK}(p) &\geq \text{OLD RANK}(x) \end{aligned}$$

(by just observing where the subtrees A, B, C, D are moved), so the number of credits we need in this step (Eq. 1) is no larger than

$$2 \text{NEW RANK}(x) - 2 \text{OLD RANK}(x)$$

(maximize the positive quantities and minimize the negative quantities in Eq. 1). That uses only $\frac{2}{3}$ of the credits we have allocated for this step, and allows us to redistribute the credits properly.

Now, we need just one more credit to pay for the rotations we do here. If

$$\text{NEW RANK}(x) - \text{OLD RANK}(x) > 0$$

then we can use the other $\frac{1}{3}$ of the credits we have allocated to this step, and we are done. But if $\text{NEW RANK}(x) - \text{OLD RANK}(x) = 0$, we need to find this credit elsewhere.

So assume that

$$\text{NEW RANK}(x) = \text{OLD RANK}(x).$$

Since

$$\text{NEW RANK}(x) = \text{OLD RANK}(g),$$

this implies that

$$\text{OLD RANK}(x) = \text{OLD RANK}(p) = \text{OLD RANK}(g)$$

and, of course, all of these are the same as the new rank of x . Remember that the rank is related to the log of the weight. So this means that, before the step, more than half of the elements in this tree are below x (in the trees C, D). If not, the weight of g would be twice that of x , and its rank would be larger.

So look at the weights:

$$\text{OLD WEIGHT}(x) + \text{NEW WEIGHT}(g) \leq \text{NEW WEIGHT}(x).$$

On the left hand side of this equation, we have all of the elements in the subtrees A, B, C, D . On the right, is the entire subtree below x after the step is performed — also A, B, C, D . But we already said that more than half of the nodes are under x before the step. This means that after the step, fewer than half of the nodes are under g . In other words, this equation say that

$$(> \frac{1}{2} \text{ of all nodes here}) + \text{NEW WEIGHT}(g) \leq \text{all nodes here},$$

and so $\text{NEW WEIGHT}(g) < \frac{1}{2}$ of all nodes in this subtree, which is $< \frac{1}{2} \text{OLD WEIGHT}(g)$. But then

$$\begin{aligned} \log_2(\frac{1}{2} \text{OLD WEIGHT}(g)) &= \log(\text{OLD WEIGHT}(g)) + \log \frac{1}{2} \\ &= \log(\text{OLD WEIGHT}(g)) - 1 \\ &= \text{OLD RANK}(g) - 1. \end{aligned}$$

So g 's rank has *decreased* by at least one.

Now by the assumptions, the rank of x does not change. p 's rank can't increase at all, so we don't have to add any credits to p . The rank of g does decrease by at least 1, and this frees up one credit which pays for the rotations at this step.

So when $\text{NEW RANK}(x) > \text{OLD RANK}(x)$, we pay for the step using some of the credits we've allocated to pay for (1) rotations and (2) maintaining the credit invariant. When $\text{NEW RANK}(x) = \text{OLD RANK}(x)$, we can do the step for free, using credits in the account at g to pay for the operation while still maintaining the credit invariant.

A very similar proof follows for case 3.

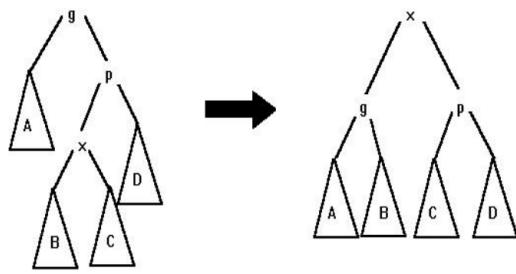


Figure 4: Case 3 of the SPLAY procedure

Figure 53. Case 3

One edge case to consider is when the new rank of x is the same of the old rank x . This can be solved by considering the ranks of g and p

Suppose that

$$\text{OLD RANK}(x) = \text{NEW RANK}(x),$$

so that we have 0 credits to do this step. It is still the case that this implies

$$\text{OLD RANK}(x) = \text{OLD RANK}(p) = \text{OLD RANK}(g),$$

so more than half of the nodes in this subtree are originally under x . But after the rotations, we see that some of the nodes under x are now under g (B) and some are under p (C). There are two possibilities.

Either: One of the nodes g and p now has $> \frac{1}{2}$ of the weight, which means that the other has $< \frac{1}{2}$ the weight and therefore drops in rank. That frees up one credit.

Or: If g and p have about the same weight, then *both* of them have dropped in rank (since each node has lost half of its weight), and we still get at least one credit freed. So, in this case too, we can take the cost of the rotations from the credits already on the tree.

Think of it this way. Let's say that the entire subtree has a bit more than 2^k nodes, and x, p, g originally had rank k . That means that most of these 2^k nodes are under x . Now after the rotation, we split these nodes between g and p . Either one of them has about 2^k nodes — which implies that the other has much less (and so a rank lower than k) — or both of them have about 2^{k-1} nodes — so both of them have a lower rank.

To be sure that you understand the proof, you should try filling in the rest of the details of this case.

Therefore the cost of a splay operation is at most $1 + 3 \log n$ credits, giving it an amortized cost of $O(\log n)$. □

PROOF The time complexities of $\text{LOCATE}(x)$, $\text{INSERT}(x)$, and $\text{DELETE}(x)$ follow from the previous proof and are all $O(\log(n))$ □

As a general comment we use credit immediately when the tree is almost balanced; paying for single rotations or when the rotations make it more unbalanced. In other cases the rotation will use credits to make the tree more balanced.

SECTION 22

Maximum Flow

We can interpret a directed graph as a flow network and use it to answer questions about, well, flow. For example an early application of this problem was to solve an issue the US army was experience: given a train network, what is the minimum set of routes to be bombed to cause maximum disruption. More formally, one can imagine a directed graph³⁷ with vertices representing conduit junctions and each edge having a certain capacity. The maximum flow problem seeks to find the greatest rate at which material can be moved from a source to a sink through a network without violating the capacity constraint of any edge.

³⁷ representing some sort of flow system i.e. pipes in a sewer system or current through a wire

Definition 59

Flow network

A flow network $G = (V, E)$ is a directed graph where each edge (u, v) has a capacity $c(u, v) > 0$. For the time being we will consider only graphs without self-loops and no reverse edges.

The source s and sink t are distinguished vertices in V such that there is a path between them, i.e. $s \rightarrow t$. Also assume that for each vertex $v \in V$ there exists a path from s to v and a path from v to t ; $s \rightarrow v \rightarrow t$

The **flow** in G is then more formally a function

$$f : V \times V \rightarrow \mathbb{R} \quad (22.1)$$

That satisfies the capacity

$$0 \leq f(u, v) \leq c(u, v) \quad \forall (u, v) \in V \quad (22.2)$$

And flow conservation constraints

$$\sum_{v \in V} f(v, u) = \sum_{v \in V} f(u, v) \quad \forall u \in V - \{s, t\} \quad (22.3)$$

$f(u, v)$ is the flow from vertex u to vertex v , and the value of f is defined as the difference between the sum of flow into the source and the sum of flow out of the source.

$$|f| = \sum_{v \in V} f(s, v) - \sum_{v \in V} f(v, s) \quad (22.4)$$

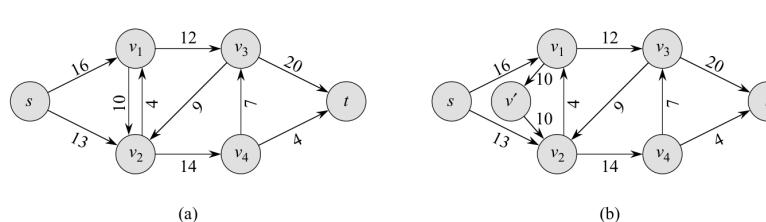


Figure 54. A network with antiparallel (reverse) edges ($v_1 \rightarrow v_2$)

Antiparallel edges can be handled by building an equivalent network with no antiparallel edges (See: 54) via splitting an antiparallel edge into a new vertex and two edges.

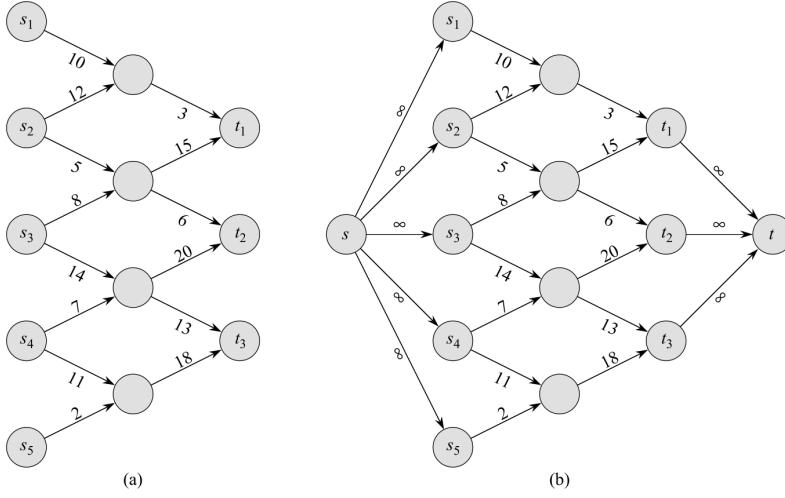


Figure 55. Multiple sources & sinks

Multiple sources and sinks can be modelled by attaching a **supersource** and **supersink** with infinite capacity between each of the existing sources and sinks (55).

One method for solving the maximum flow problem is the **Ford-Fulkerson** method³⁸

³⁸It actually encompasses a number of implementations with different run-times

FORD-FULKERSON-METHOD(G, s, t)

Definition 60

```

1    $f = 0$ 
2   while  $\exists$  augmenting path  $p \in$  residual network  $G_f$ 
3       augment flow  $f$  along  $p$ 
4   return  $f$ 
```

Definition 61

Residual networks G_f are graphs that represent the how much we can still change the flow on the edges of G , i.e $c_f(u, v) = c(u, v) - f(u, v)$. If this capacity is negative the edge is omitted from the residual network. Note that it may be sometimes necessary to send back flow along an edge (or decreasing the flow along an edge). The maximum amount of backflow that can be supported along the edge (u, v) is the current flow along the edge. This can be accomplished by adding a new edge in the opposite direction (v, u) with capacity equal to $f(u, v)$

So, more formally,

$$c_f(u, v) = \begin{cases} c(u, v) - f(u, v) & \text{if } (u, v) \in E \\ f(v, u) & \text{if } (v, u) \in E \\ 0 & \text{otherwise} \end{cases} \quad (22.5)$$

The **augmentation** of flow f by f' , $f \uparrow f'$ is given by

$$(f \uparrow f')(u, v) = \begin{cases} f(u, v) + f'(u, v) - f'(v, u) & \text{if } (u, v) \in E \\ 0 & \text{otherwise} \end{cases} \quad (22.6)$$

Lemma 3 | $f \uparrow f'$ is a flow in G with value $|f + f'| = |f| + |f'|$

Lemma 4 Flow is conserved, i.e.

$$(f \uparrow f')(u, v) = f(u, v) + f'(u, v) - f'(v, u) \leq f(u, v) + f'(u, v) \leq f(u, v) + c_f(u, v) = c(u, v) \quad (22.7)$$

$$\sum_{v \in V} (f \uparrow f')(u, v) = \sum_{v \in V} (f \uparrow f')(v, u) \quad (22.8)$$

Definition 62 An **augmenting path** p is a simple path from s to t in the residual network G_f .

Definition 63 The **Edmonds-Karp** algorithm is an instance of the Ford-Fulkerson method that chooses the shortest (in terms of edge count) augmenting path using BFS and can be shown to terminate after $O(VE)$ augmentations and has $O(VE^2)$ runtime (since BFS is $O(V + E)$)
Suppose p is a path.

$$C_p(p) = \min\{c_f(u, v) : (u, v) \text{ on path } p\} \quad (22.9)$$

SECTION 23

Theory of Computation

SUBSECTION 23.1

Finite Automata

Definition 64 **Deterministic Finite Automata (DFAs)**³⁹

Given a sequence of inputs they can accept (true) or reject (false).
More technically, a DFA is a 5-tuple $(Q, \Sigma, \delta, q_0, F)$ where

- A set of states Q
- An input alphabet Σ (domain of input)
- Transition function $\delta : Q \times \Sigma \rightarrow Q$ which takes in a state and some input $\in \Sigma$ to transition to some other state
- An initial state $q_0 \in Q$
- A set of accepting states $F \subseteq Q$

³⁹Glorified if-statements

Example | Create a DFA that accepts strings of the form "abc"

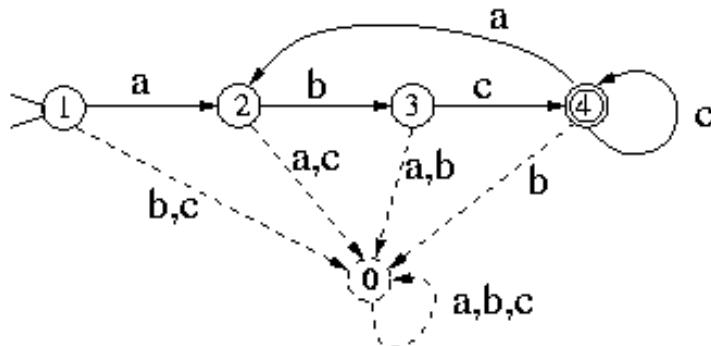


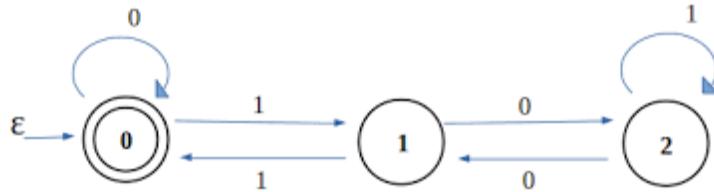
Figure 56. Note non-accepting state 0 w/ loop. Accepting state q_3 indicated via double circle

Example | Create a DFA that accepts all inputs with an even number of 0s. $\Sigma = \{0\}$

$$q_0 \xleftrightarrow{0} q_1 \quad (23.1)$$

where q_0 is the accepting state

Example | Create a DFA on binary strings that only accepts on strings with a number of 0s that is a multiple of 3



Definition 65 Non-deterministic Finite Automata (NFA)

These are the same as DFAs except we can take an additional input ϵ which is the empty string. On ϵ the NFA branches and computes each branch in parallel. If *any* branch reaches an accepting state the NFA accepts.

Example | Create an NFA that accepts on any input with either a multiple of 2 or 3 0s. This is simply the two DFAs we found earlier connected via a branch on q_0 .

Example | Find a NFA that accepts on either an even number of 0s or exactly two 1s

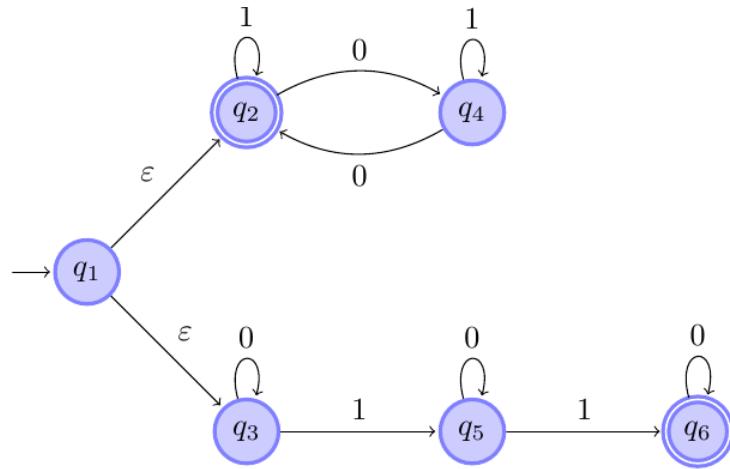


Figure 57. Note how this is just two DFAs connected together

Theorem 11

All NFAs can be converted to an equivalent DFA⁴⁰

⁴⁰NFA doesn't provide any additional computability

DFAs and NFAs are memoryless, for example a DFA or NFA is incapable of accepting on a string e.x

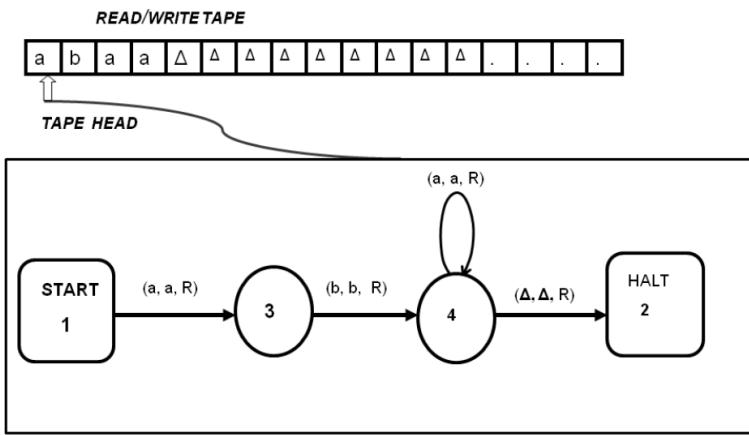
$$0^n I^n \quad \text{for} \quad n \geq 0 \quad (23.2)$$

sequences like 000111 for n = 3, etc

SUBSECTION 23.2

Turing Machines

In order to address this Turing machines were introduced. They are basically just a DFA with infinite memory (commonly represented as a tape)



A Turing Machine for aba^*

Figure 58. A Turing machine for aba

Turing machines can compute anything⁴¹

Definition 66

Nondeterministic Turing Machine (NTM)

Or: a turing machine with a NFA head. Like how a DFA can be converted into a NFA (which would be faster because of parallel computation), a NTM has much better complexity compared to a regular TM.

⁴¹ except for some things like the halting problem etc. TLDR if a problem is not computable by turning machine it is not computable by any finite means. E.g. printing all real numbers between 0 to 1; this is an uncountable infinity but the turning machine tape is countable; larger infinity

SUBSECTION 23.3

Complexity Classes

- P : Polynomial number of steps to compute on a TM
- NP : Non-deterministic Polynomial: takes a polynomial number of steps to compute on a NTM.

Theorem 12

All problems in NP can be verified on polynomial time on a TM, since verifying is like going backwards; since when computing one has to go through all branches on a NTM but when verifying we only have to take one path.

SUBSECTION 23.4

NP Completeness

$$A <_n B \quad (23.3)$$

(23.3) means that A is *polynomially reducible* to B . There are about seven equivalent ways

to interpret this.

1. There exists a polynomial time function f such that $\alpha \in A \Leftrightarrow f(\alpha) \in B$
 - The contrapositive can also apply; $f(\alpha) \notin B \Rightarrow \alpha \notin A$
 - If we can solve it in B -space then we can solve it in A -space
2. We can decide $B \Leftrightarrow$ we can decide A , i.e. that A
3. Runtime of $A \leq$ runtime $B + O(n^k)$ (some polynomial factor)

Hard: Takes a long time to compute

Definition 69

NP-Complete

A decision problem L s.t. $L \in NP$ such that

1. $L \in NP$
2. $L' \leq_p L$ for all $L' \in NP$ (NP-hard)

Suppose $A \leq_p B \Rightarrow A \leq B + O(n^k)$

$$\begin{aligned}
 & \text{if } B \in P \Rightarrow A \in P \\
 & \text{if } A \in P \Rightarrow B \in P \\
 & \text{if } B \in NPC \Rightarrow A \in NPC \\
 & \text{if } A \in NPC \Rightarrow L' \leq_p A \leq_p B \Rightarrow B \in NP\text{-hard}
 \end{aligned} \tag{23.4}$$

Informally, a problem is NP-Complete if it is at least as hard as any other problem in NP

23.4.1 Solving NPC Problems

Given some decision problem A

1. Showing $A \in NP$
 - Provide a *certificate* of $A \in NP$: the evidence that the problem is an instance of A
 - Provide a *verification procedure* to check of the certificate is really an instance of A
 - Argue why the verification procedure is polynomial time
2. Showing $A \in NP\text{-hard}$
 - Determine which NPC problem L' we will use (usually provided on exam)
 - Show that $L' \leq_p A$; transform a known NPC problem to an unknown (i.e. NP-hard)

Showing that A is NP-Complete requires proving both of the above.

Example

- Ham-Cycle: Given $G = (V, E)$, find a simple *cycle* that travels through all vertices.
- Ham-Path: Given $G = (V, E)$, find a simple *path* that travels through all vertices.

1. Show that Ham-Cycle is in NP

- **Certificate:** a cycle $C = \{v_0, v_1 \dots v_n\}$
- **Verification:**
 - Check that all of the edges in the cycle are in the graph; $\forall i, (v_i, v_{i+1}) \in E$. This is $O(VE)$
 - Check that all vertices appear exactly once $O(V)$
 - This verification procedure is clearly of polynomial time

Proof can be fairly hand wavy for these simple complexities

2. Show that Ham-Cycle \in NP-hard via Ham-Path

$$\text{Ham-Path} \leq_p \text{Ham-Cycle} \quad (23.5)$$

We note that we can transform the known NPC problem (that we proved in the previous part of this example) into an unknown by adding an additional vertex x to G to form G' . Then, for each vertex v of G , add an edge to (v, x) . This will guarantee forming a ham-cycle. Construct $G' = (V', E')$, $V' = V \cup x$

$$E' = E \cup (v, x), (x, v) | v \in V \quad (23.6)$$

Claim: $\langle G \rangle \in \text{Ham-Path} \Leftrightarrow \langle G' \rangle \in \text{Ham-Cycle}$

- \Rightarrow proof: Suppose G is an instance of Ham-Path ($\exists V = v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_n \in G$), therefore $x \rightarrow v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_n \rightarrow x$ is a ham-cycle in G . So G' is an instance of a ham-cycle
- \Leftarrow proof: Suppose G' is an instance of Ham-Cycle, then $\exists v_0 \rightarrow v_1 \rightarrow \dots \rightarrow v_n \rightarrow x \rightarrow v_0$; a known cycle in G'
- We can remove x to form a ham-path in G
- And therefore G is an instance of ham-path

Definition 70 Combinatorial Equivalence Checking

Given boolean formulas $A(x_1, x_2, x_3 \dots x_n)$ and $B(x_1, x_2, x_3 \dots x_n)$, does there exist an assignment such that $A \neq B$?

Formula-SAT: Given a boolean formula $F(y_1, \dots, y_k)$, is there an assignment such that $F = 1$?

Example

1) CEC \in NP.

- Certificate: assignment x_1, \dots, x_n
- Verification: Evaluate A, B and check if $A \neq B$

2) CEC \in NP-hard via Formula-SAT

$$\text{Formula-SAT} \leq_p \text{CEC} \quad (23.7)$$

Given a boolean formula F , construct $A = F$ and $B = 0$. Then, make a claim that $\langle F \rangle \in \text{Formula-SAT} \Leftrightarrow \langle A, B \rangle \in \text{CEC}$.

- \Rightarrow proof: there exists an

Example

Prove that Half-3-CNF-SAT is NP-Complete. Given a 3-CNF formula Φ with n variables and m clauses, does there exist a truth assignment to Φ that causes exactly $m/2$ of the clauses to evaluate to true and exactly $\frac{m}{2}$ of the clauses to evaluate to false. Now that Φ can contain duplicate clauses and tautological clauses like $x \vee \neg x \vee y$

1: Show that Half-3-CNF-SAT is in NP

Certificate: an assignment of true, false sequence to Φ . Verification: evaluate all clauses and make sure exactly half are 1 and the other half is 0. This is $O(m)$

2: Show that Half-3-CNF-SAT is NP-hard via 3-CNF-SAT

$$3\text{-CNF-SAT} \leq_p \text{Half-3-CNF-SAT} \quad (23.8)$$

Attempt: construct $\Phi' = \Phi \wedge (p \vee q \vee r)^m$. This is incorrect.

Attempt: construct

$$\Phi' = \Phi \wedge (p \vee q \vee r)^m \wedge (p \vee q \vee r)^{2m} \quad (23.9)$$

Claim:

$$<\Phi> \in 3-SAT \Leftrightarrow <\Phi'> \in Half-3-CNF-SAT \quad (23.10)$$

Example

- Graph-3-colourability: Given $G = (V, E)$, can you assign one of three colours to each vertex such that no two adjacent vertices have the same colour?
- k-clique-cover: Given $G = (V, E)$ and k , can we partition the vertices v into v_1, \dots, v_k such that each v_i is a clique. This is known to be NP-Hard
- 1) Certificate: an assignment of colours to edge vertices. Verification: iterate over each edge and check that the vertices on each end of the vertex are not of the same colour.
 - 2) 3-col is NP-hard via k-clique-cover. In particular we will do this via 3-clique-cover

$$3\text{-clique-cover} \leq_p 3\text{-col} \quad (23.11)$$

Given $G = (V, E)$ is an instance of 3-clique-cover, construct $G' = (V', E')$ is the complement of $G - V' = V, E' = \{(u, v) | (u, v) \notin E\}$.⁴²

Claim:

$$<G> \in 3\text{-clique-cover} \Leftrightarrow <G'> \in 3\text{-col} \quad (23.12)$$

⁴² Note that this is in polynomial time. Remember to put this on the solution for exams!

Forwards: suppose G has a 3-clique-cover. So then we color each set v_i a different color. Then $u.\text{color} = v.\text{color} \Rightarrow (u, v) \in E \Rightarrow (u, v) \notin E'$. So G is 3-colourable

Backwards: suppose G' is 3-colourable. \exists coloring such that $u.\text{color} = v.\text{color} \Rightarrow (u, v) \notin E' \Rightarrow (u, v) \in E$. Therefore let all nodes of the same color be a clique; G .

Example

Partition: Let P be a set of positive integers. P is an instance of Partition if

$$\exists P_1, P_2 \subseteq P \text{ such that } P_1 \cup P_2 = P \text{ and } P_1 \cap P_2 = \{\} \text{ and } \sum_{x \in P_1} x = \sum_{x \in P_2} x. \quad (23.13)$$

Subset-sum: Given a set of positive integers X and positive integer t , does there exist $S' \subseteq S$ s.t. $\sum S' = t$

Showing that this is NP is fairly trivial so we'll skip it. Showing that it is NP-hard is trickier. We will take the approach of showing that Partition is NP-Hard via Subset-sum

$$\text{Subset-sum} \leq_p \text{Partition} \quad (23.14)$$

Given set S and target t , construct $x = \sum S, P = S \cup \{x - 2t\}$.

Claim:

$$<S, t> \in \text{subset-sum} \Leftrightarrow <P> \in \text{Partition} \quad (23.15)$$

Forwards: Suppose $\exists T_1 \subseteq A$ st. $\sum T_1 = t$. Let

PART

V

ECE360: Electronics

SECTION 24

Admin and Preliminary

SUBSECTION 24.1

Lecture 1

Taught by Prof. Khoman Phang

24.1.1 Mark Breakdown

Table 5. Mark Breakdown

Test 1	15
Test 2	20
Homework	10
Labs	12
Final	43

24.1.2 Diodes

Diodes are an electronic valve which causes current to only flow in one direction. An ideal diode is an open circuit in the closed direction and a closed circuit in the other, so the current is always in the direction of the arrow (+'ve @ arrow base, -'ve at arrow point)⁴³.

operation	actual cost	amortized cost
push	$O(1)$	$O(1)$
pop	$O(1)$	$O(1)$
copy	$O(n)$	$O(1)$

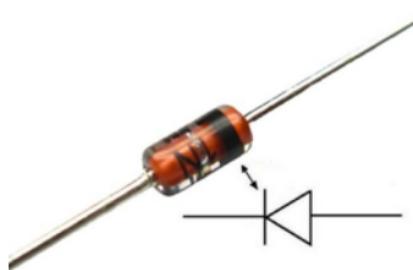
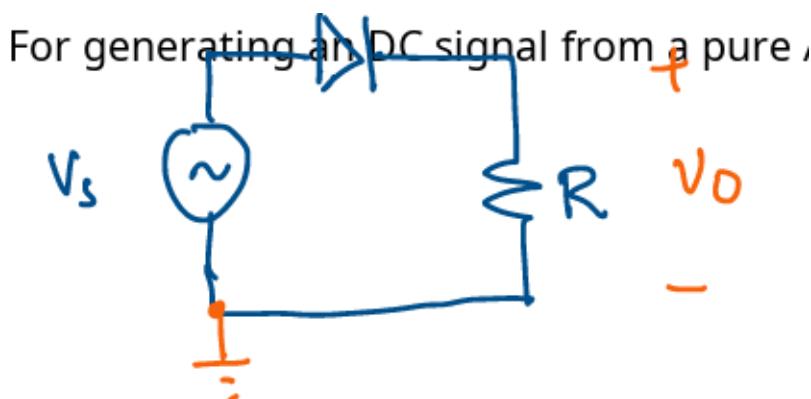
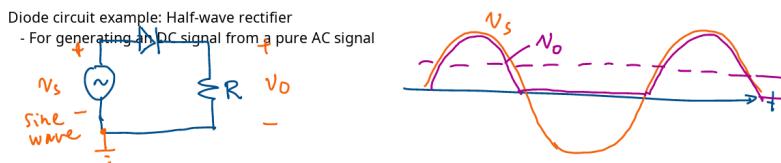


Figure 59. A diode and its symbol

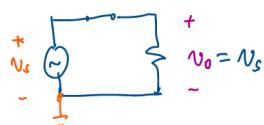
An example of a diode circuit is the half-wave rectifier which turns an AC signal to a DC signal

Can take oscilloscope over resistor to see that a pure DC signal has been generated

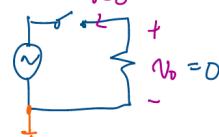




If $N_s > 0$, diode is "ON"



If $N_s < 0$, diode is "OFF"



SECTION 25

Diodes

SUBSECTION 25.1

Lecture 2

More formally, off/on for diodes should be referred to as:

- Off \leftrightarrow reverse bias
 - On \leftrightarrow forwards bias

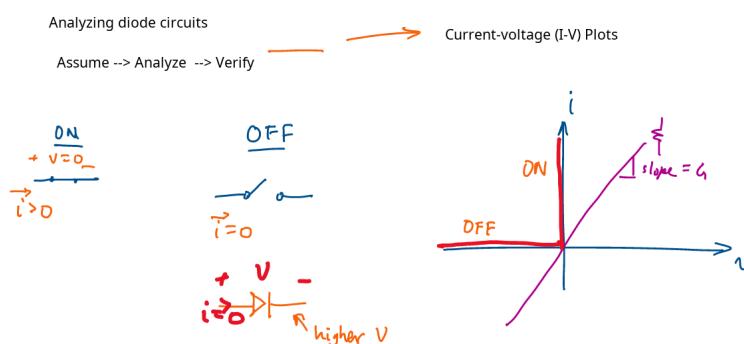


Figure 60. General steps for analyzing non-linear circuits. Note plotting out expected response

An example of how this is used in circuit design is to manage two power sources. Consider an Arduino that could be powered by an AC adapter or by a computer's USB port. This circuit would choose the higher voltage source and prevent backflow into the other power source due to any potential power differentials. It is also effectively an OR gate.

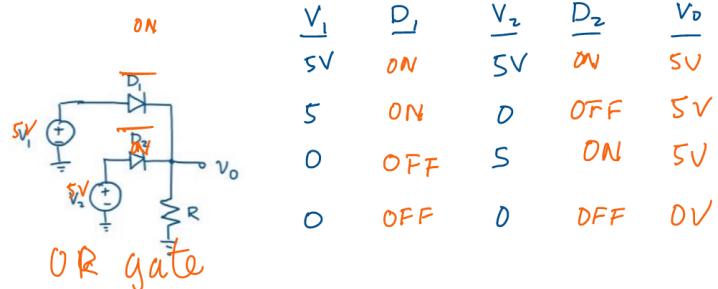
Analysis Examples

'0' '1'

Example 1:

Find output voltage V_o assuming input voltages V_1 and V_2 are either 0V or 5V.

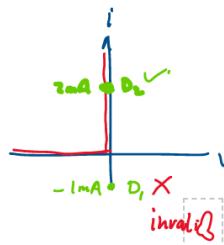
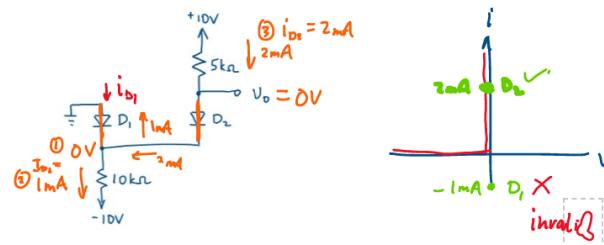
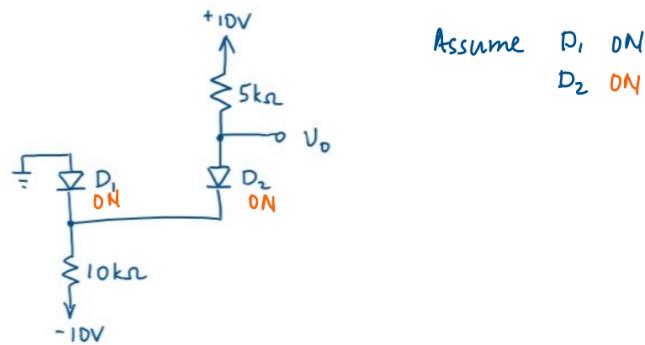
What is the function of this circuit?



Example 2:

Find output voltage V_o

Assume \rightarrow Analyze \rightarrow Verify

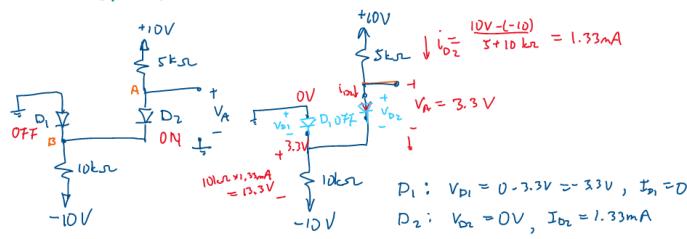


In this example the initial assumption was incorrect.

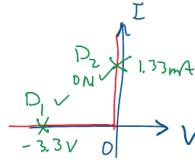
Let's try another analysis with D_1 off and D_2 on:

Second attempt

Step 1: Assume
 D_1 OFF, D_2 ON \Rightarrow Step 2: Analyze circuit

Step 3: verify assumptions

assume $I \rightarrow$ check $V < 0$
 D_1 OFF \rightarrow open $\rightarrow I_1 = 0$
 $D_1 \checkmark$ $V_{D1} = -3.3V < 0 \checkmark$
 Diode D_1 is OFF



assume $V \rightarrow$ check $I_{D2} > 0$
 D_2 ON \rightarrow short $\rightarrow V_{D2} = 0V$
 $D_2 \checkmark$ $I_{D2} = 1.33mA > 0 \checkmark$
 Diode D_2 is ON

If we were to do this brute force we'd have to consider 4 cases, so it's important to build up some sort of intuition for the circuit.

SUBSECTION 25.2

Lecture 3

Today we're going to look at the characteristics of real diodes.

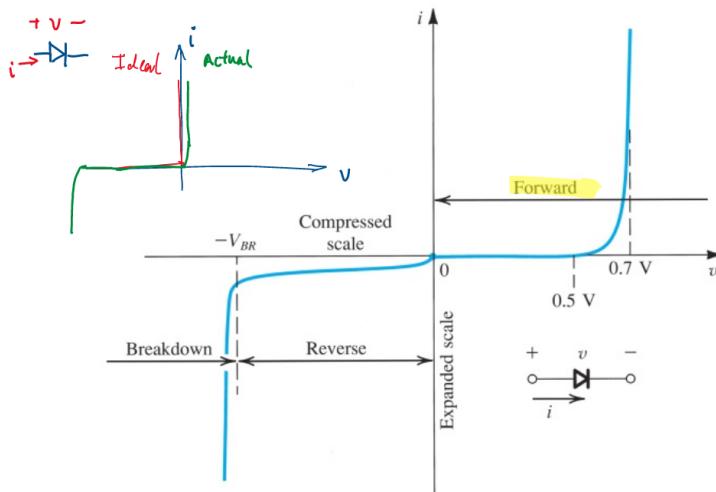


Figure 4.8 The silicon diode i - v relationship with some scales expanded and others compressed in order to reveal details.

Real diodes have a little bit of leakage current and also encounter a breakdown point where they're no longer able to block the current.

$$i = I_s (e^{\frac{V}{V_T}} - 1) \quad (25.1)$$

Where:

$$V_T = \frac{kT}{q} \quad [V] \quad (25.2)$$

Most of the time we can assume that the circuit is at room temperature and that $v_T = 25mV$. Note that this value explodes when $V > V_T$ which is the breakdown point. When encountering a reverse bias $V_s < 0$, the -1 term comes in and causes $i \approx I_s$. The scale current is just a general constant which varies in range from 10^{-9} to $10^{-15} A$ and scales with temperature, doubling with every approximately $5^\circ C$ increase in temperature. Note: the ideal diode equation can be rearranged to find an expression for voltages

$$V = V_T \ln \left(\frac{i}{I_s} \right) = \ln(10) V_T \log_{10} \left(\frac{i}{I_s} \right) \quad (25.3)$$

These expressions turns out to be quite reliable for reasonable diodes to reasonable voltages.

k is Boltzmann's constant, T is temperature in Kelvins, q is the charge of an electron. I_s is the scale current which is usually $\approx 1pA$, which doesn't change much until the breakdown point.

Using the ideal diode equation we can find the relationship between voltages and currents as they pass through the diode.

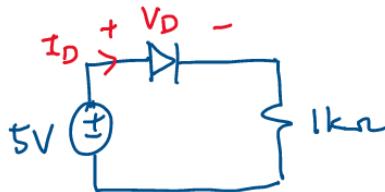
$$\frac{i_2}{i_1} = \frac{I_s e^{\frac{V_2}{V_T}}}{I_s e^{\frac{V_1}{V_T}}} = e^{\frac{V_2 - V_1}{V_T}} \quad (25.4)$$

$$V_2 - V_1 = V_T \ln \left(\frac{i_2}{i_1} \right) \xrightarrow{\text{room temperature}} 60mV \log_{10} \frac{i_2}{i_1} \quad (25.5)$$

Example:

Calculate the diode voltage and current in the circuit below.

Assume that the diode voltage is 0.7V at 1mA and $V_T = 25mV$.



Example

Recall (25.1). Plugging in the given values gives us the scale current.

$$1mA = I_s e^{\frac{0.7V}{25mV}}, I_s = 6.9 \cdot 10^{-16} A \Rightarrow I_o = I_s e^{v_o/v_T} \quad (25.6)$$

Ohm's law can then be applied at the resistor

$$V_r = IR = I_o R = 5V - V_D \Rightarrow 5 - V_D = I_o R \quad (25.7)$$

So we have two equations and two unknowns (since we know $v_T = 25mV$ but v_o was used at first just to find I_s) Solving for the unknowns gives us:

- $V_o = 0.736V$

V_D is the voltage across the diode

- | • $I_D = 4.264mA$

SECTION 26

Lecture 4 & 5: Forward conducting diodes

The exponential model accurately describes the diode outside of the breakdown region, though its nonlinear behaviour makes it difficult to use.

For $V_{DD} > 0.5V$

V_{DD} is DC voltage, v_d is small signal voltage, V_D is the diode voltage

$$I_D = I_S e^{V_D/V_T} \quad (26.1)$$

Where

- I_S is the diode parameter
- V_T is the thermal voltage

Another equation may be produced via Kirchhoff's law

$$I_D = \frac{V_{DD} - v_d}{R} \quad (26.2)$$

The unknown quantities I_D and v_d may be solved for via graphical analysis or iteration.

Example This simple circuit is used to demonstrate the exponential model of the diode.

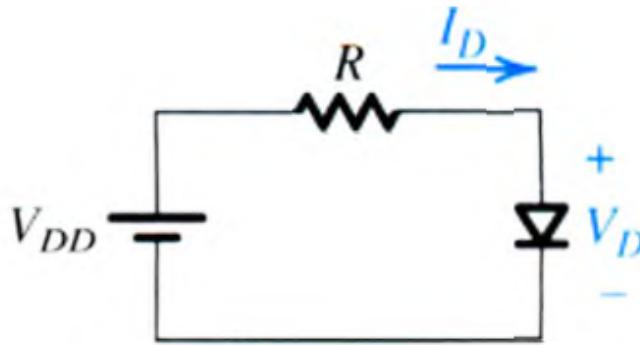
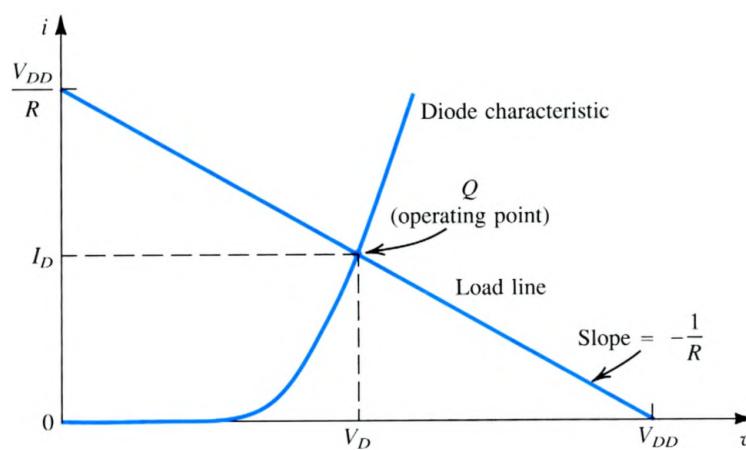


Figure 61. Simple example circuit with diode

Plots of the diode characteristics and Kirchhoff's relation are plotted, the intersection of which gives the solution.



An iterative procedure may also be applied to solve for the unknowns, the procedure for which will be illustrated through an example

Example Find I_D , V_D for the circuit in the previous example (Fig. 61). $V_{DD} = 5V$, $R = 1k\Omega$, and at $V_D 0.7V$, $I_D = 1mA$

1. Assume $V_D = 0.7V$, then use (26.2) to find I_D .

$$I_D = \frac{5V - 0.7V}{1k\Omega} = 4.3mA \quad (26.3)$$

2. Use the diode equation (26.1) to get a better estimate for V_D .

$$V_2 - V_1 = 2.3V_T \log \frac{I_2}{I_1} \Rightarrow V_2 = V_1 + 0.06 \log \frac{I_2}{I_1} \quad (26.4)$$

substituting $V_1 = 0.7V$, $I_1 = 1mA$, $I_2 = 4.3mA$,

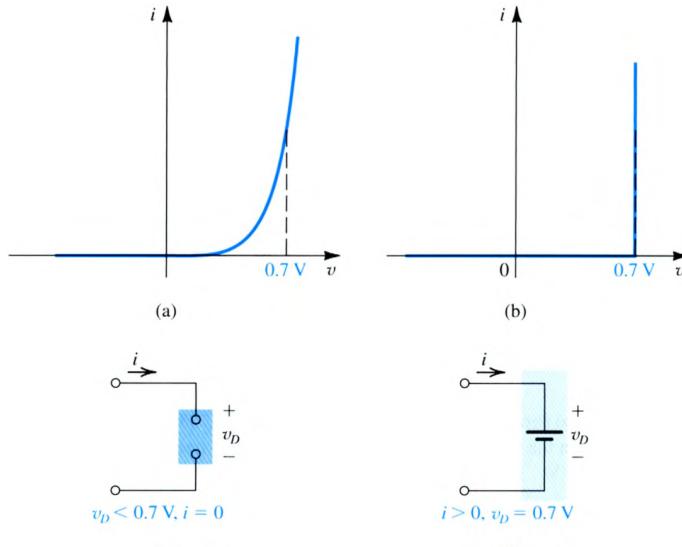
$$V_2 = 0.738V \Rightarrow I_D = 4.3mA, V_D = 0.738V \quad (26.5)$$

- This states that for a decade⁴⁴ change in current the diode voltage drop changes by $2.3(V_T \approx 60mV)$ which is negligibly small for $v < 0.5V$. The voltage at which this behaviour becomes significant is called the **cut-in voltage**

⁴⁴Factor of 10

3. Repeat steps 1 and 2 with the new values until the values more or less become stable

This iterative model is powerful and yields accurate results, but can be computationally expensive especially when calculating by hand. To address this we employ other models such as the *constant-voltage-drop* model which approximates the exponential characteristics via a piecewise linear model. The reason why this is possible is because forward conducting diodes exhibit a voltage drop that varies in a relatively narrow range.



Using the constant voltage drop model in our analysis looks the same as before, but with V_D directly taking on the value of $0.7V$ (as per the prior example) instead of being solved for with the diode equation.

In applications that involve voltages greater than the voltage drop (i.e. usually $\approx 0.6 - 0.8V$) we can neglect the diode voltage drop altogether while calculating the diode current.

$$V_D = 0V$$

$$I_D = \frac{5 - 0}{1} = 5mA \quad (26.6)$$

This is generally good enough for a first estimate, though the previous model isn't that much more work and gives more accurate results. The primary use of this model is to determine which diodes are on or off in a multi-diode circuit

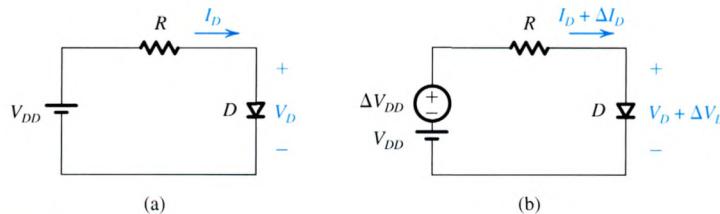
SECTION 27

Small-Signal Model

The small signal method is an alternative model used to describe the nonlinear diode's characteristics with greater accuracy than piecewise linear models.

Consider a small ΔV_{DD} applied to the diode, which would cause a small $\Delta I_D, \Delta V_D$. We want to find a quick way of determining the values of these incremental changes.

Similar methods will be applied to transistors in later chapters



Skipping a bunch of math⁴⁵ the results are as follows: ⁴⁶

⁴³ recall: passive sign convention

⁴⁵ It is 11:17pm and I have two more lectures to catch up to today

Definition 71

Small signal approximation

$$i_D(t) \approx I_D(1 + \frac{v_d}{V_T}) \quad (27.1)$$

This is valid for when variations in diode voltage $|v_d| \lesssim 5mV$.

$$r_d = \frac{V_T}{I_D} \quad (27.2)$$

From this we can define the small signal resistance as the resistance relating i_d to v_d

The steps for calculating the small signal model are as follows:

1. Perform a dc analysis using the exponential, constant-voltage-drop, or piecewise-linear model.
2. Linearise the circuit. For a forward-based diode, find r_d by substitution I_D into (27.2). The small-signal equivalent circuit is found by eliminating all independent dc sources⁴⁶ and replacing the diode with its small-signal resistance r_d
3. Solve the linearised circuit. In particular we would want to find ΔI_D , ΔV_D and check to see if it is consistent with our approximation, i.e. that $\Delta V_D \lesssim 5mV$

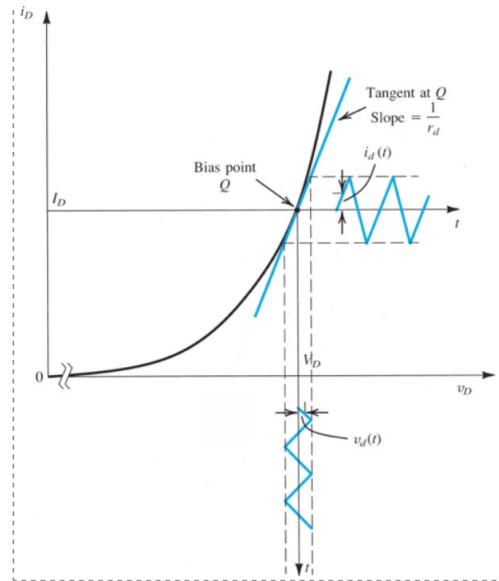
The reason why we linearise these non-linear systems, we, as engineers, try to linearise them because it is convenient to be able to use superposition, phasors, Fourier series, Laplace transforms, and so forth.

⁴⁶ Small signal analysis can be performed separately from the dc bias analysis because of the linearization of diode characteristics in the small-signal approximation

SECTION 28

Lecture 6: Small signal model, cont'd

V_D can be thought of an input to a transfer function that is the diode, with I_D being the output. If the input signal is more or less a triangular wave the output will be as well.



Since we are applying a linear approximation, superposition may apply;

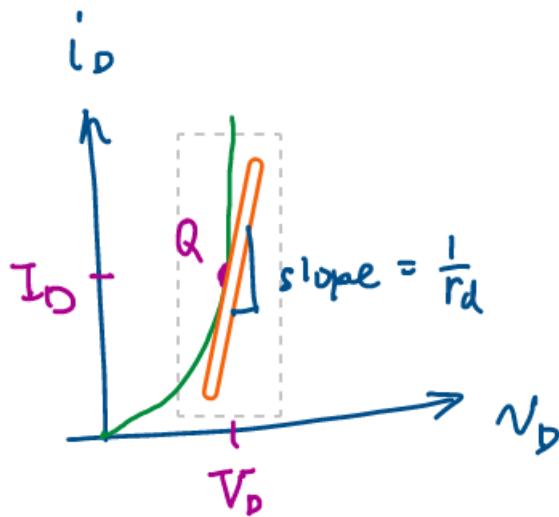
$$f(a + b + c \dots) = f(a) + f(b) + f(c) + \dots \quad (28.1)$$

Here we can think of the function as

$$f(V_o + v_d) = f(V_o) + f(v_d) \quad (28.2)$$

Small signal analysis works because, by superposition, we can zero out the other sources (V_o , etc) and inspect only the effect of the small signal voltage on the system.

28.0.1 Deriving small-signal resistance



Resistance can be found as the slope of the $I - V$ relationship. The following is a proof of (27.2).

PROOF

$$i_D = I_s e^{\frac{V_D}{V_T}} \quad (28.3)$$

$$\text{slope} = \frac{di_D}{dv_o} = I_s \left(\frac{1}{V_T} \right) e^{\frac{V_D}{V_T}} \quad (28.4)$$

Then, substitute values at the operational point Q , i.e. I_D, V_{DD}

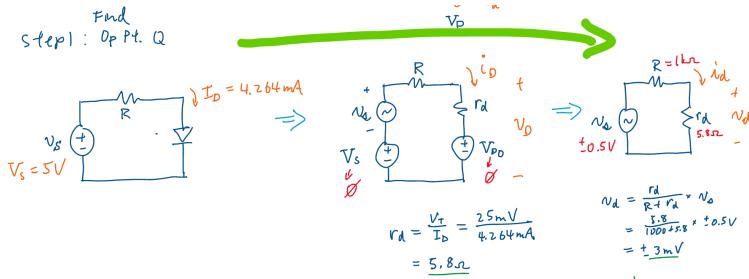
$$\text{slope} = I_s e^{V_D/V_T} \left(\frac{1}{V_T} \right) \Rightarrow \frac{I_D}{V_T} = \frac{1}{r_d} \quad (28.5)$$

And therefore

$$r_d = \frac{V_T}{I_D} \quad (28.6)$$

□

1. Calculate DC bias point Q
2. Derive small-signal circuit
3. Analyze small-signal circuit
4. If required, recombine to arrive at final result



The textbook skips the middle step where we actually apply the small signal approximation and all the values have not been offset by the bias point yet. Note that voltages/values in the 3rd step (small signal approx) are all relative to the bias point Q .

Applying this to our circuit we find

$$r_d = \frac{V_T}{I_D} = \frac{25mV}{4.264mA} = 5.8\Omega \quad (28.7)$$

This was not as accurate as the exponential model but it is more than close enough 99% of the time. But also the $0.7V$ constant voltage drop model is usually sufficient as well.

PROOF The small-signal approximation is valid for voltage variations up to $\pm 5mV$

$$\begin{aligned} i_D &= I_S \exp(V_D/V_T) \\ &= I_S \exp \frac{V_D D + v_d}{V_T} \\ &= I_D \times e^{v_d/V_T} \end{aligned} \quad (28.8)$$

e^x can be expanded to a power series

$$\begin{aligned} e^x &= 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \dots \\ &\approx 1 + x \quad \text{if } \frac{x^2}{3!} \ll x \rightarrow x \ll 2 \end{aligned} \quad (28.9)$$

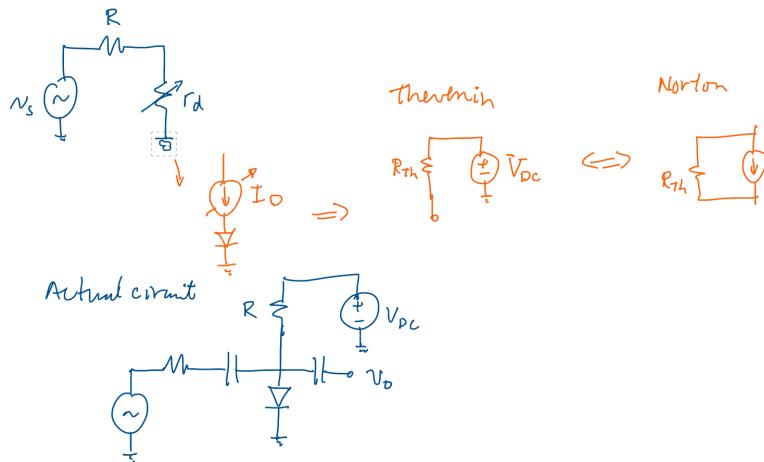
Therefore $\frac{v_d}{V_T} \ll 2$, so $v_d \ll 2V_T = 50mV$
So make

$$|v_d| < 5mV \quad (28.10)$$

□

In the lab we'll be using a variable attenuator circuit, which involves a voltage source. Can make a current source with a voltage source and a large resistor.

variable attenuator circuit

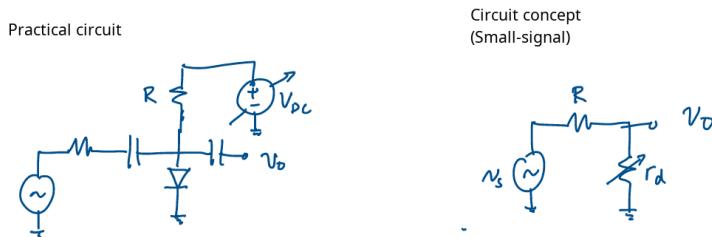


This will be covered more next lecture. TLDR: can use small signals and then build up a circuit to get the intended Q using biases.

SECTION 29

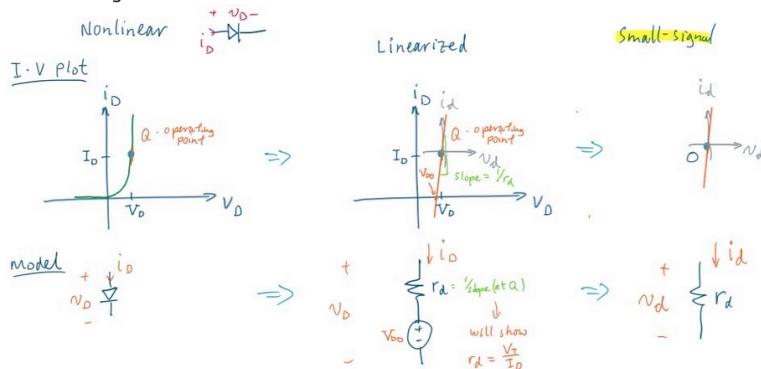
Lecture 7

Let's take a deeper look into the variable attenuator circuit discussed last class.



- The capacitors in this circuit can be treated like short circuits because they are large (which is why we don't see them in the small-signal circuit).
- The constant voltage offset source gets zeroed out in the small-signal circuit because it is a constant, thereby becoming a ground.
- Current sources become open circuits
- Negative currents

What does a negative current mean?

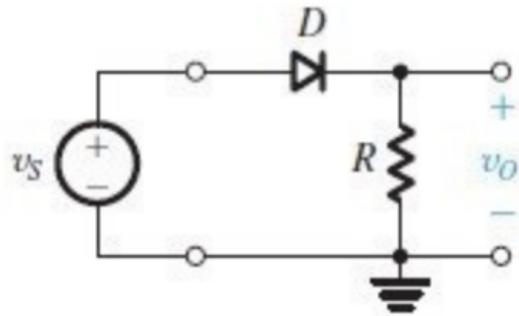


SECTION 30

Rectifiers

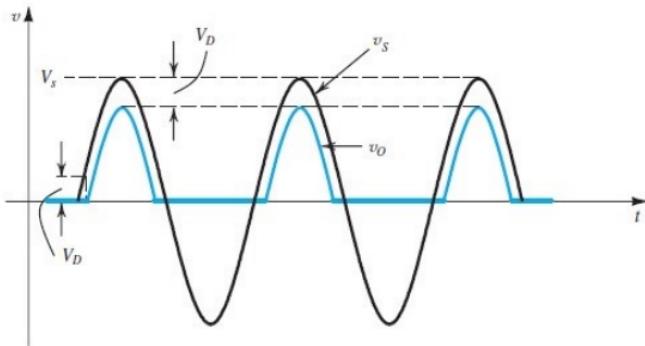
We have previously talked about the half-wave rectifier circuit, which used a single diode to rectify a waveform to only the positive values

recall half-wave rectifier:



This isn't very efficient because it only uses half of the waveform, so it is desirable to build a *full wave* rectifier

Most of the homework, etc we will be practicing working from the practical circuit to the small-signal design, whereas as a designed we will usually be going the other way around.



What we want with a full-wave rectifier is to automatically interchange the wires whenever the input signal goes from negative to positive and vice-versa as to provide the load a wholly positive signal.

A full wave rectifier can be built as follows

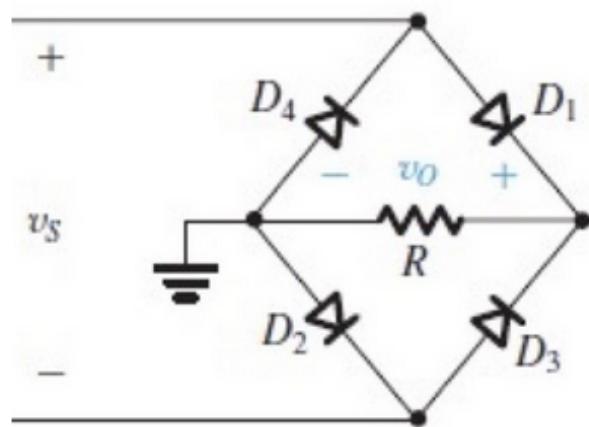


Figure 62. When the signal is positive it will flow from diode 1 to the load and then to diode 2. When it is the negative polarity it will go through diode 3 to the load and then return through diode 4. Note how the direction remains the same, and how we selectively connect the ground to the top or bottom terminal depending on the input polarity.

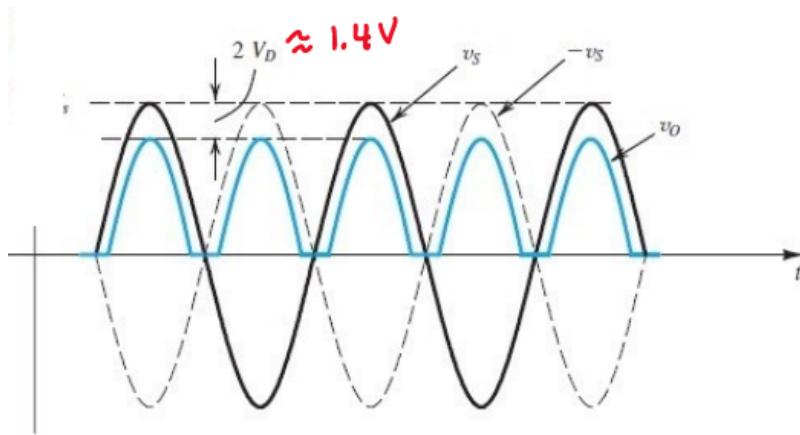


Figure 63. Comparing the two there is a little more loss but at least we get the full wave now

A **peak rectifier** uses a capacitor as an energy storage device to hold the voltage at the peak value.

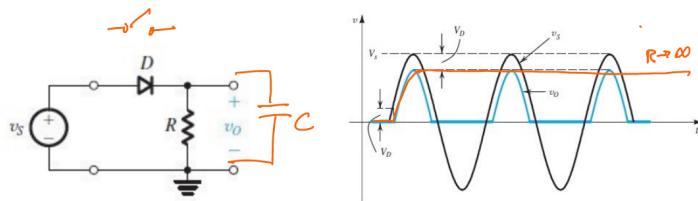
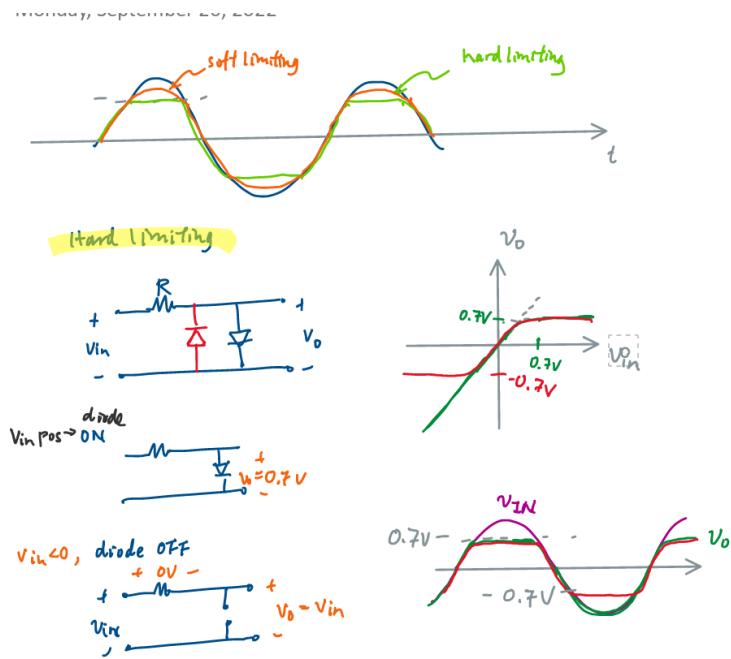


Figure 64. Capacitor size and frequency must be tuned as to avoid voltage drop (recall RC circuits. Note that the time constant $\propto RC$, so a bigger resistor (slower rate of discharge) or bigger capacitor (bigger energy storage) can reduce voltage drop)

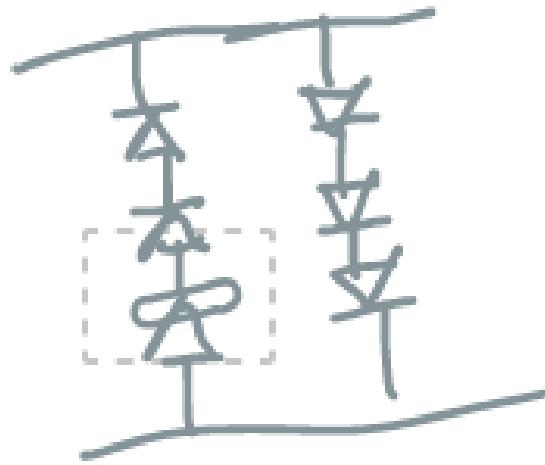
SUBSECTION 30.1

Lecture 8

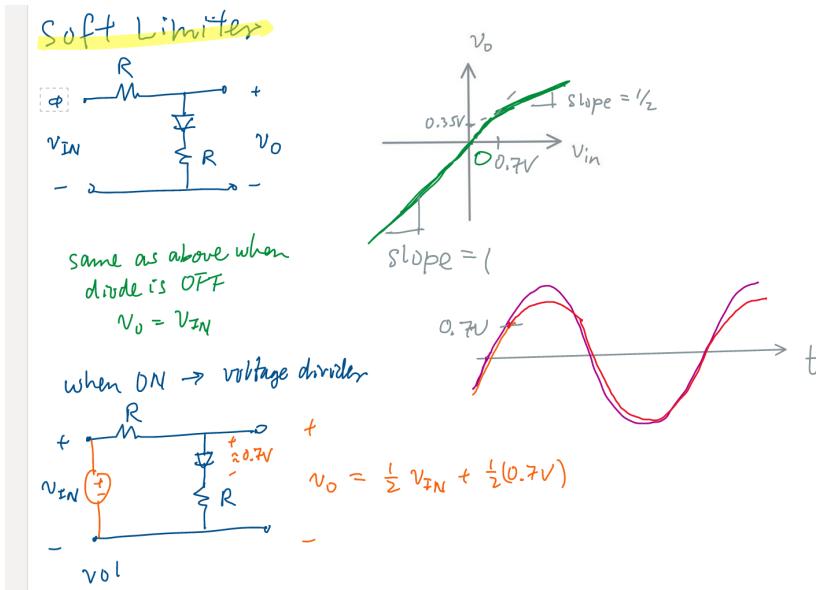
Hard limiting circuits include a pair of diodes which cause only a maximum of $0.7V$ signal to pass through.



Increasing the hard limiting voltage can be done by putting a bunch of diodes in series, i.e.



The hard, squared-off signal response (and multiple of 0.7V) may not be desirable, so we can use a soft limiter instead.



This will act like nothing is happening when the diode is off, i.e. $v_o = v_{in}$. When the diode is on the soft limiting circuit acts like a voltage divider. For this circuit there are two resistors with the same R so the output voltage from just the current divider is

$$v_o = \frac{R}{R+R} v_{in} = \frac{1}{2} v_{in} \quad (30.1)$$

We will also have to account for the voltage of the diode, so

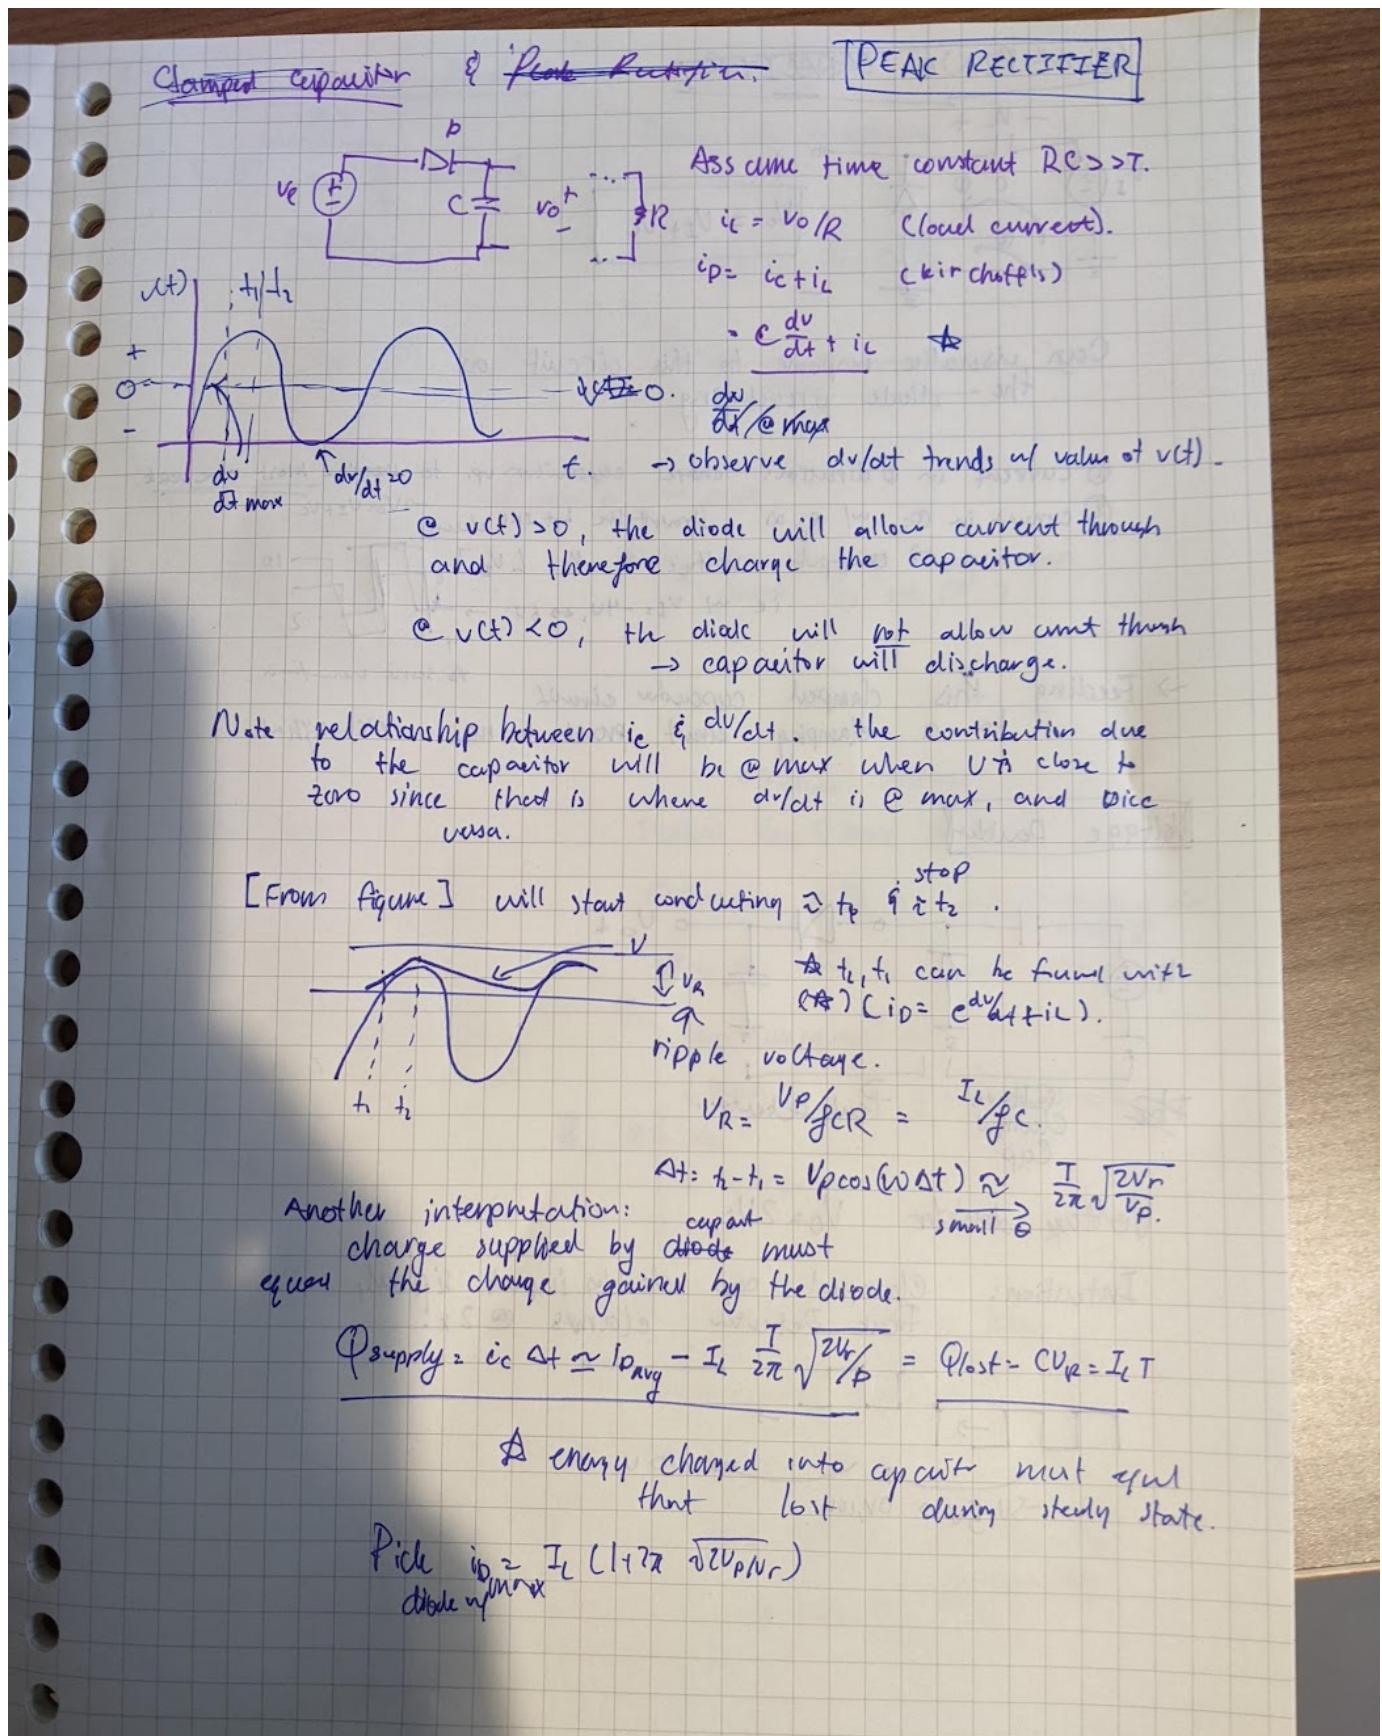
$$v_o = \frac{1}{2} v_{in} + \frac{1}{2} (v_d = 0.7V) \quad (30.2)$$

$$v_o = \frac{1}{2} v_{in} + \frac{1}{2} 0.7V \quad (30.3)$$

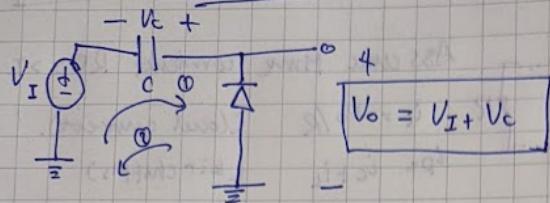
SUBSECTION 30.2

Lecture 9

I missed this lecture, so here are my handwritten catchup notes



CLAMPING CAPACITOR



Can visualize response to this circuit as the diode preventing

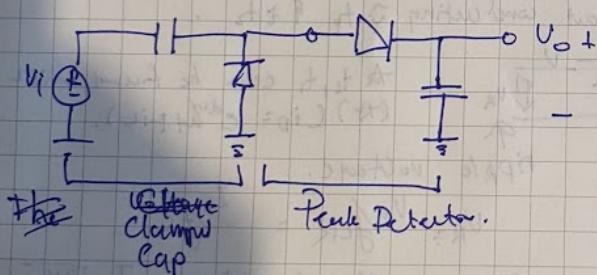
① current in \odot direction: charge capacitor up. to ~~most~~ most -ve peak.

② current in \odot : w/ $R_L \approx \infty$, straight line but zero current.

\rightarrow output voltage $\propto V_C = L V_I$ [6-10] ie w/ $V_I = -4V \rightarrow 6V$ \rightarrow same waveform.

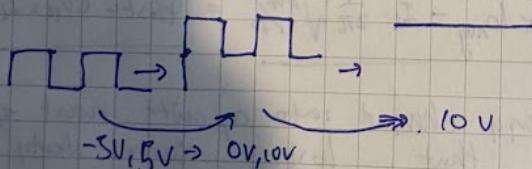
\rightarrow Feeding this clamped capacitor circuit to a clamping circuit provides a nice DC signal.

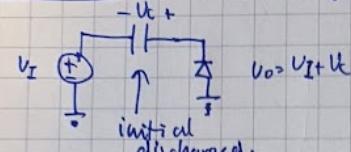
Voltage Doubler



Voltage doubler $V_O = 2V_I$

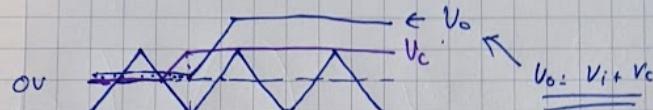
Intuition: Clamped cap doubles input signal, Peak Detector clamps @ 2x.



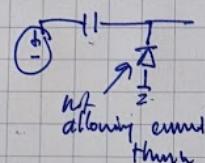
Example Clamped Analysis

Input: 5V triangular.

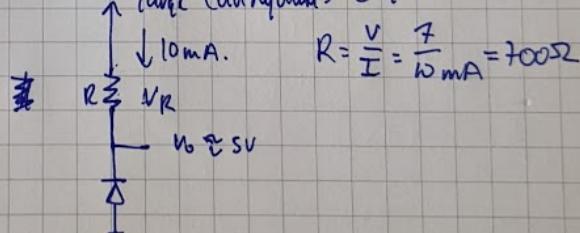
$$V_O = V_I + V_C = V_I + C \frac{dV}{dt}$$



diode off: ON: get value to -1ve peak.

Why does it only charge when circuit G , not D ?Breakdown Region"Zener Voltage"Typical $V_{BZ} \approx 200V = V_Z$ Operation $\approx V_Z$ is not destructive for small I

- can be used to generate voltage.

 V_{PC} large (unregulated) e.g. 1W.

Desired regulated voltage may be created by exploiting the breakdown current

$$(V_O = V_Z) + V_R = V_{DR}$$

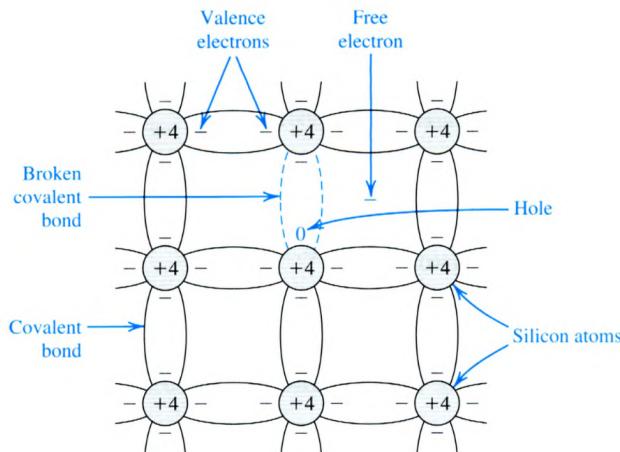
∴ pick V_R s.t. $V_Z = V_{PC} - V_R$

$$\text{so pick } R = \frac{V_R}{I}$$

SUBSECTION 30.3

Lecture 10

The most commonly used semiconductor in electronics is silicon. At a molecular level silicon crystals are linked together via covalent bonds which spontaneously break and reform. This process allows electrons to wander away from its parent and therefore become available to conduct electricity. The process of electrons filling some of these holes is called recombination and the process of electrons leaving the crystal to form holes is called generation. At thermal equilibrium the recombination rate is equal to the generation rate and therefore the concentration of free electrons n is equal to that of the holes p , i.e.



$$n = p = n_i \quad (30.4)$$

Where n_i denotes the number of free electrons in a unit volume of silicon at a given temperature.

Theorem 14

The number of free electrons and holes in a unit volume (cm^3) of silicon at a given temperature T is given by

$$n_i = BT^{\frac{3}{2}} e^{\frac{-E_g}{2kT}} \quad (30.5)$$

Where k is the Boltzmann constant, B is a material-dependent parameter that is $7.3 \cdot 10^{15} cm^{-3} K^{-\frac{3}{2}}$ for silicon, T is the temperature in kelvin, and $E_g = 1.12 eV$ is the bandgap energy or the minimum energy required to break a covalent bond and create an electron-hole pair.

The intrinsic silicon crystal does not have sufficient amounts of free electrons and holes to be appreciably conductive at room temperature, plus their conductivity varies strongly with temperature. **Doping**, or introducing impurities into silicon, can be used to increase the number of free electrons and holes in the crystal. To increase the concentration of free electrons, n Silicon is commonly doped with an element with 5 valence electrons such as phosphorous. Likewise, to increase the concentration of free holes, p Silicon is commonly doped with an element with 3 valence electrons such as boron.

Theorem 15

For n type silicon, the concentration of free electrons is given by

$$p_n n_n = n_i^2 \Leftrightarrow p_n \approx \frac{n_i^2}{N_P} \quad (30.6)$$

For p type silicon, the concentration of free electrons is given by

$$p_p n_p = n_i^2 \Leftrightarrow p_p \approx \frac{n_i^2}{N_A} \quad (30.7)$$

Note that pieces of n or p doped silicon are electronically neutral; the charge of the majority free carriers (electrons in n , holes in p) are neutralized by the bound charges of the impurity atoms.

There are two distinct ways that current can flow in semiconductors; drift and diffusion. **Drift current** is the movement of free electrons or holes in response to an electric field; holes are accelerated in the direction of E and free electrons opposite. Holes acquire a velocity of $v_{p-drift} = \mu_p E$, with μ_p being a constant⁴⁸. If we were to consider a cross section of a silicon bar with a field applied across it, the hole component of the direct current (and current density) flowing through the bar would be

$$I_{s,p} = A q p v_{p-drift} = A q p \mu_p E \Rightarrow J_p = q p \mu_p E \quad (30.8)$$

The current due to free electrons is found similarly and can be used to find the total drift current density

$$I_{s,n} = -A q n v_{n-drift} \Rightarrow J_{s,n} = q n \mu_n E \quad (30.9)$$

$$J_S = J_{s,p} + J_{s,n} = q(p\mu_o + n\mu_n)E = \frac{E}{\rho} \quad (30.10)$$

Where resistivity ρ is given by

$$\rho = \frac{1}{q(p\mu_p + n\mu_n)} \quad (30.11)$$

We now observe that (30.10) is a variant of Ohm's law and can be written as

$$\rho = \frac{E}{J} \quad (30.12)$$

Diffusion current is the movement of holes and electrons from areas of high concentration to low concentration in response to a gradient inside the semiconductor⁴⁹.

The magnitude of the current at any point is proportional to the slope of the concentration gradient, i.e. for holes

$$J_{D,p} = -q D_p \frac{dp(s)}{dx} \quad (30.13)$$

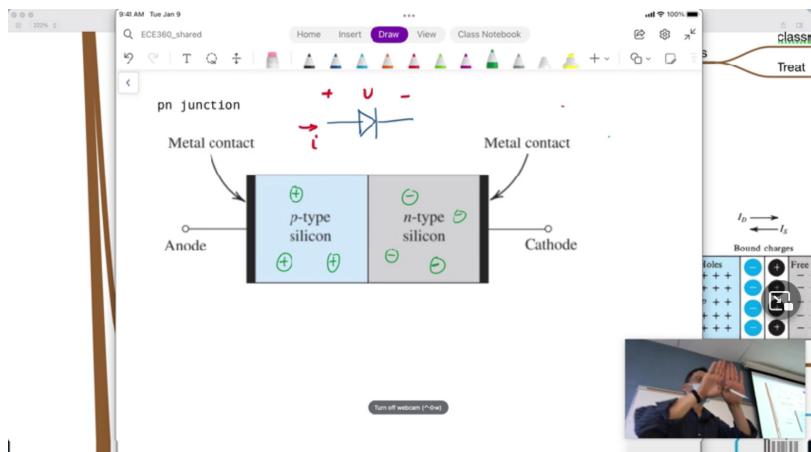
And for electrons

$$J_{D,n} = q D_n \frac{dp(s)}{dx} \quad (30.14)$$

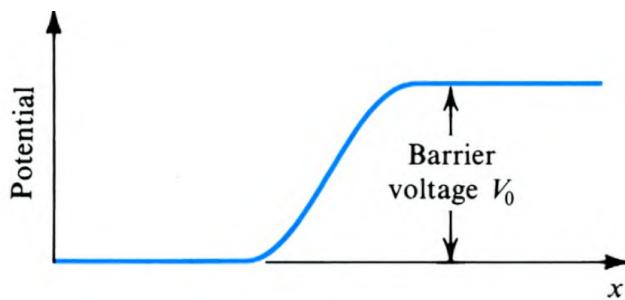
This brings us to our first example of a semiconductor device: the $p - n$ junction diode.

⁴⁷ since we already accounted for them in step 1

⁴⁸ $\mu_p = 480 \text{ cm}^2/\text{V} \cdot \text{s}$ for intrinsic silicon



Looking at a passive device without an applied voltage junction we note that free electrons from *n* side gets caught in the holes in the *p* type, leaving behind a positive charge and vice-versa causing a "depletion region" in the middle with an electric field without any holes or electrons. This causes a potential field and makes it such that work must be done to move charge from the anode to the cathode.



Reverse-biasing the diode causes an even greater potential inside the diode which causes the depletion region to grow and the electric field to increase, thereby prevent current from passing through. A forward bias would counteract the depletion region and therefore allow current to flow through the diode. Drift current⁵⁰ dominates in the reverse bias region, while diffusion current dominates in the forward bias region.

⁴⁹ Think of dropping ink into water

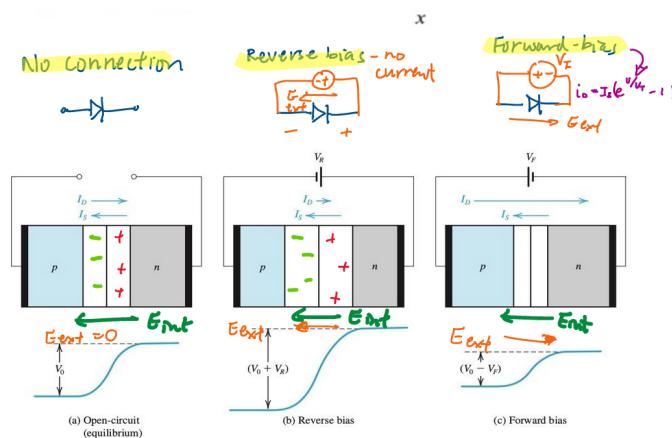


Figure 3.11 The *pn* junction in: (a) equilibrium; (b) reverse bias; (c) forward bias.

Skipping past the derivation, the expression for diode current is given by

$$i_D = I_s(e^{\frac{v_D}{v_T}} - 1) \quad (30.15)$$

This is a result of the Boltzmann distribution and the derivation is given in some Feynman lecture or in the textbook.

SUBSECTION 30.4

Lecture 11: Introduction to transistors and MOSFETs

Transistors are essentially switches that can be turned on and off by applying a voltage, and are implemented in semiconductors by stacking another *p* or *n* layer onto a diode, i.e.



Figure 65. pnp vs npn transistor

There are two main types of transistors, MOSFET (metal-oxide semiconductor field-effect transistors) and BJT (bipolar junction transistor). MOSFETs have found more widespread use due to their simplicity of manufacture and their ability to be scaled down to smaller sizes.

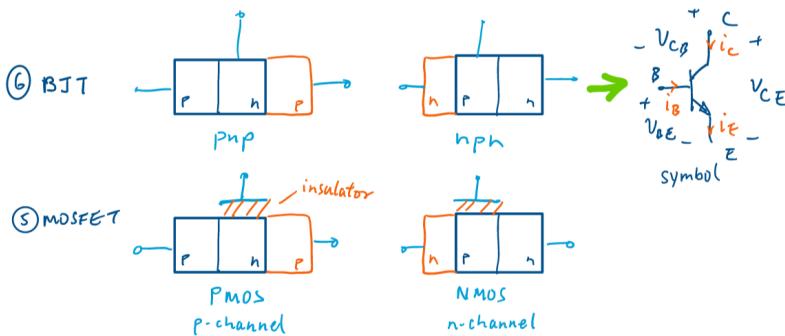


Figure 66. BJT transistors came first but are not commonly used now; they will be discussed later on in the course. Whereas BJTs directly make contact with the semiconductor, MOSFETs apply a field to induce a change in the device.

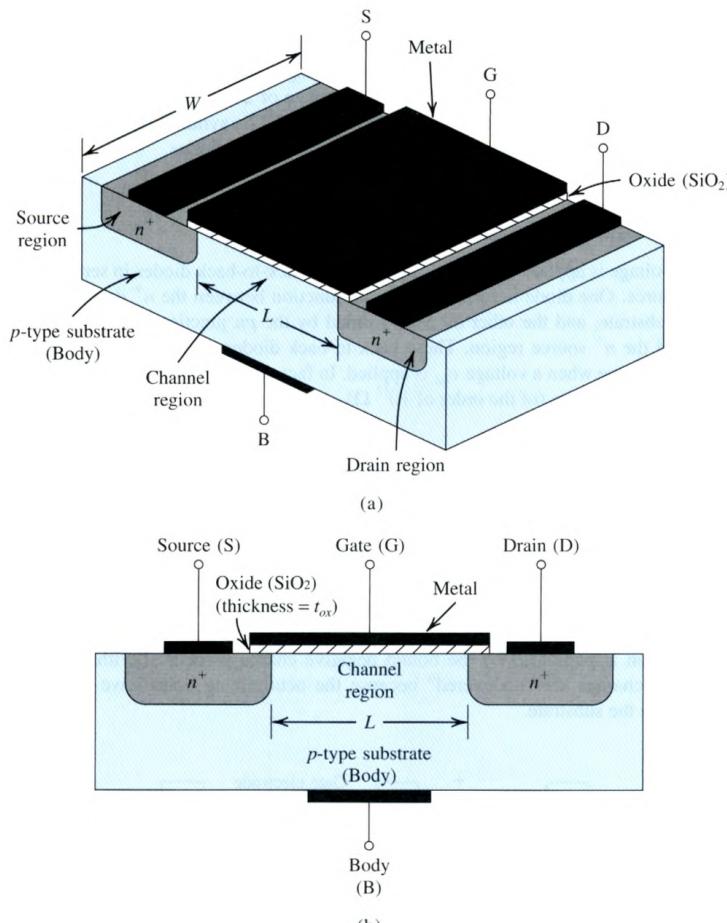


Figure 67. CMOS transistors

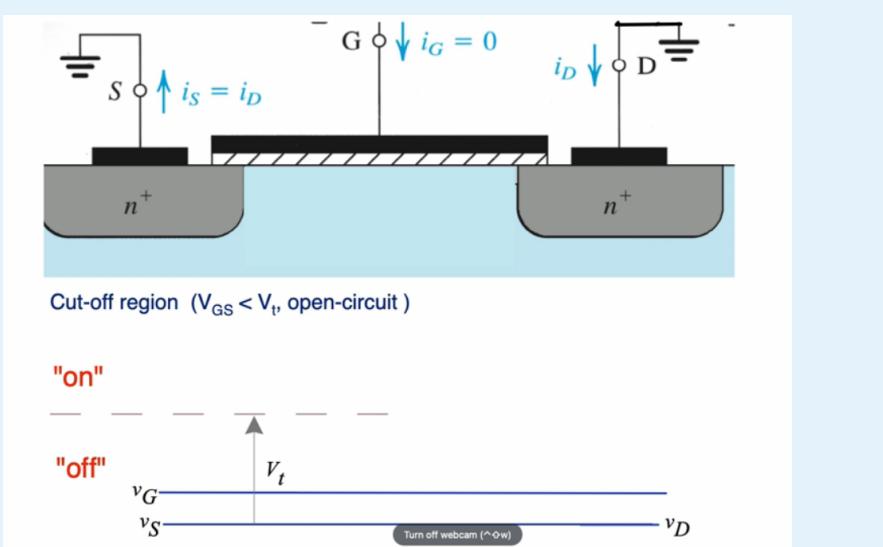
P transistors use holes as their primary charge carrier and allow current to pass if the base is negative with respect to the emitter⁵¹. Conversely, N transistors carry charge using electrons and allow current to pass if the base is positive with respect to the emitter⁵²

Definition 72

Cut-off region

⁵⁰ recall: movement of carriers due to electric field, electrons drifting back to n region

⁵¹ For a PMOS this would be the application of a negative field



There is no charge if the gate voltage is the same as the source. At some point, the threshold voltage, there will be enough voltage to form a "n-channel"⁵³ through which electrons may pass.

⁵³ again, for a n-type

SUBSECTION 30.5

Lecture 12: MOSFET Operating Regions

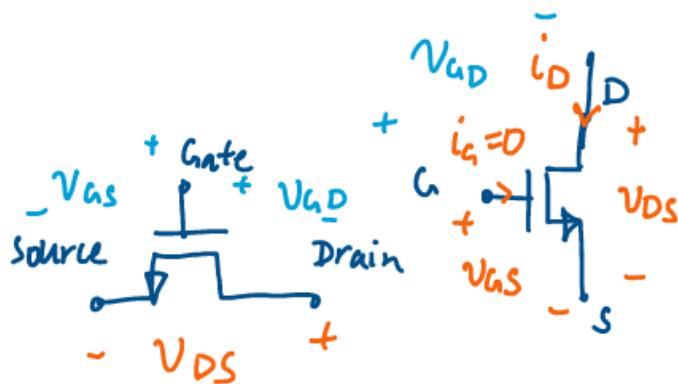
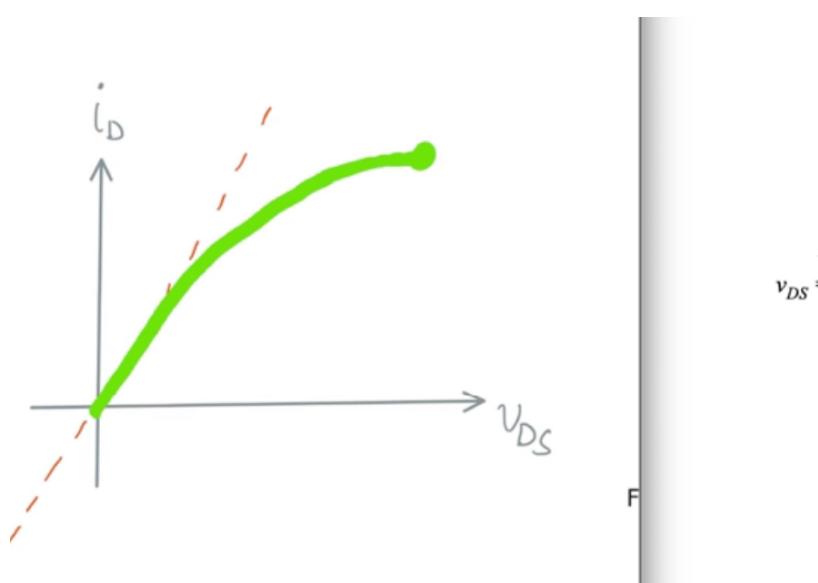


Figure 68. Labelled NMOS. The drain current can be thought of as the output.

There are three MOSFET operation regions; cut-off, triode, and saturation. Regions can be identified by inspecting the various voltages in the device.

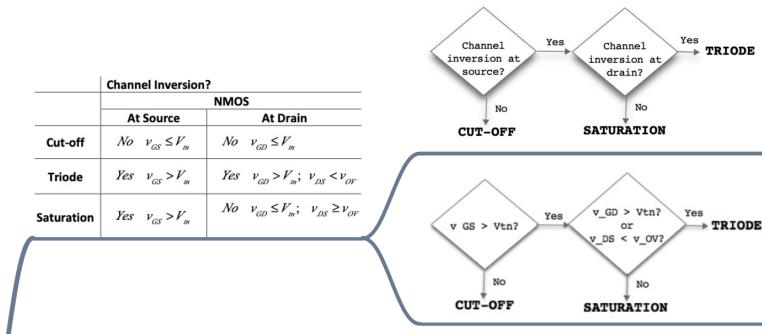
Channel Inversion



V_t is the threshold voltage.
Note that since this is a symmetrical device, we can compare V_{GS} and V_{GD} w.r.t the source and drain

Figure 69. The n -channel will thin out as the drain voltage goes up; at some point we will hit a drain voltage equal to the threshold voltage and cause the channel to disappear

When the drain voltage is equal to the threshold voltage, the overdrive voltage is equal to the V_{DS} . This is the boundary between the triode and saturation region.

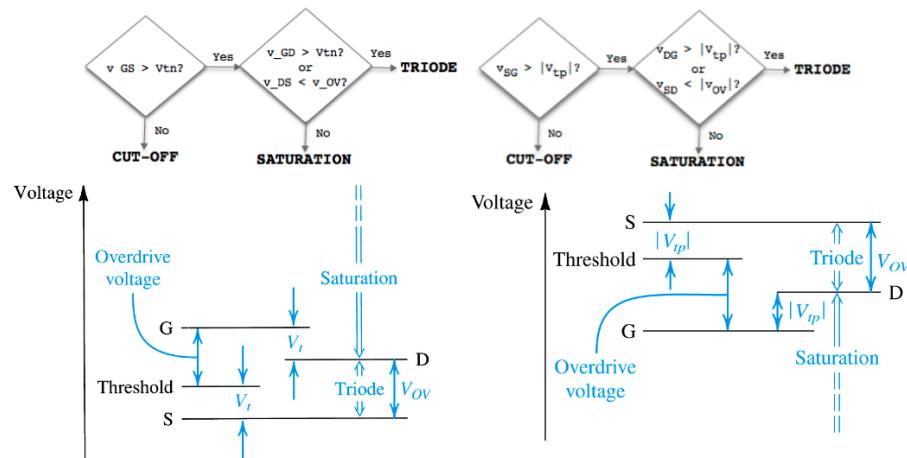


Remark For a NMOS the source is the terminal at the lower voltage and the drain is the terminal at the higher voltage. For a PMOS it is the opposite, i.e. the source is at higher voltage. It is conceivable that the source and drain could be swapped during operation.

MOSFET Regions of Operation

 v_{OV} : Overdrive voltage $v_{OV} = v_{GS} - V_m$

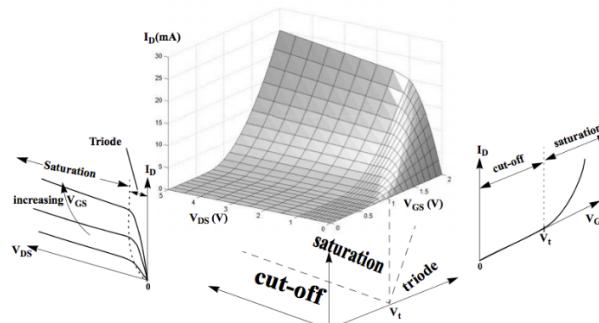
Channel Inversion?				
		NMOS	PMOS	
At Source		At Drain	At Source	
Cut-off	No $v_{GS} \leq V_m$	No $v_{GD} \leq V_m$	No $v_{SG} \leq V_{tp} $	No $v_{DG} \leq V_{tp} $
Triode	Yes $v_{GS} > V_m$	Yes $v_{GD} > V_m$; $v_{DS} < v_{OV}$	Yes $v_{SG} > V_{tp} $	Yes $v_{DG} > V_{tp} $; $v_{SD} < v_{OV} $
Saturation	Yes $v_{GS} > V_m$	No $v_{GD} \leq V_m$; $v_{DS} \geq v_{OV}$	Yes $v_{SG} > V_{tp} $	No $v_{DG} \leq V_{tp} $; $v_{SD} \geq v_{OV} $



MOSFET Device Equations

Operating Region	NMOS	PMOS
Cut-off	$i_D = 0$	$i_D = 0$
Triode	$i_D = \mu_n C_{ox} \left(\frac{W}{L} \right) (v_{GS} - V_m) v_{DS} - \frac{v_{DS}^2}{2}$	$i_D = \mu_p C_{ox} \left(\frac{W}{L} \right) (v_{SG} - V_{tp}) v_{SD} - \frac{v_{SD}^2}{2}$
Saturation	$i_D = \frac{1}{2} \mu_n C_{ox} \left(\frac{W}{L} \right) (v_{GS} - V_m)^2 [1 + \lambda v_{DS}]$	$i_D = \frac{1}{2} \mu_p C_{ox} \left(\frac{W}{L} \right) (v_{SG} - V_{tp})^2 [1 + \lambda v_{SD}]$

n-Channel MOSFET Current-Voltage Characteristic



Lecture 13: MOSFET current-voltage characteristics

The TLDR of this lecture is that the current through a MOSFET is given by those fancy big equations. A cool consequence of the MOSFET is that it acts like a variable resistor in the triode region.

PART

VI

MAT389: Complex Analysis

Complex Numbers

Lecture 1

Consider a 2-vector $\vec{x} = (x, y) \in \mathcal{R}$. As complex numbers correspond to two-vectors

$$\vec{x} = (x, y) \leftrightarrow z = x + iy, i^2 = -1 \quad (31.1)$$

z is, therefore, a complex variable. What are the benefits of a complex number like z ?

Definition 73

Imaginary and Complex Numbers

i is an imaginary number such that

$$i^2 = -1 \quad (31.2)$$

A complex number has the form:

$$z = x + iy \quad (31.3)$$

This prof lectures at the speed of sound and talks *into* the board. Couldn't quite follow during this lecture, hopefully I get better about it in the following ones.

Definition 74

There are a number of operations we can perform on complex numbers.

Addition

$$z + z' = (x + x') + i(y + y') \quad (31.4)$$

Multiplication

$$zz' = (x + iy)(x' + iy') = (xx' - yy') + i(xy' + x'y) \quad (31.5)$$

PROOF

Proof of (31.5):

$$\begin{aligned} zz' &= (x + iy)(x' + iy') \\ &= x + ixy' + iyx' + i^2yy' \\ &= xx' - yy' + i(xy' + yx') \end{aligned} \quad (31.6)$$

□

Magnitude

$$|z| = \sqrt{x^2 + y^2} \quad (31.7)$$

Conjugate

The complex conjugate has the properties:

- $\bar{z}z = |z|^2$
- $\overline{(z + z')} = \bar{z} + \bar{z}'$
- $\overline{z \cdot z'} = \bar{z} \cdot \bar{z}'$

We can define a new operation

$$\forall \text{complex } z, \exists \text{ complementary number } w \text{ such that } zw = wz = 1 \quad (31.8)$$

Denote

$$w = \frac{1}{z} = z^{-1} \quad (31.9)$$

PROOF

Proof of (31.9): Find w s.t. $zw = 1$

$$\begin{aligned} zw &= 1 \\ w\bar{z}z &= \bar{z}z = |z|^2 > 0 \\ |z|^2 w &= \bar{z} \\ w &= \frac{\bar{z}}{|z|^2} \rightarrow Z^{-1} = \frac{\bar{z}}{|z|^2} \end{aligned} \quad (31.10)$$

□

Furthermore, there are operators that we can define on complex numbers.

Definition 75

Real and Imaginary Operators

Given $z = x + iy$, we can define the real and imaginary operators

$$x = \operatorname{Re}\{z\} \quad (31.11)$$

$$y = \operatorname{Im}\{z\} \quad (31.12)$$

Example

$$\operatorname{Im}\left\{(1 + \sqrt{2}i)^{-1}\right\} \quad (31.13)$$

By (31.9), we have

$$\operatorname{Im}\{z^{-1}\} = \frac{-\operatorname{Im}\{z\}}{|z|^2} \quad (31.14)$$

And

$$\operatorname{Re}\{z^{-1}\} = \frac{-\operatorname{Re}\{z\}}{|z|^2} \quad (31.15)$$

Using these, for example, we find that the $\operatorname{Im} = \frac{-\sqrt{2}}{3}$

We can get the real component in a similar way.

Here is an enumeration of absolute value properties for complex numbers:

$$|z \cdot w| = |z||w| \quad (31.16)$$

$$|z + w| \leq |z| + |w| \quad (31.17)$$

$$|\bar{z}| = |z| \quad (31.18)$$

$$|z + w|^2 = (\bar{x} + \bar{w})(z + w) = |z|^2 + |w|^2 + \bar{z}w + \bar{w}z \quad (31.19)$$

PROOF Note that $\bar{z}w + \bar{w}z = 2\operatorname{Re}\{z\bar{w}\}$, by (31.19)
And so

$$|z + w|^2 \leq |z|^2 + |w|^2 + 2|z||w| = (|z| + |w|)^2 \quad (31.20)$$

□

SUBSECTION 31.2

Lecture 2

Whereas a two-vector $\vec{x} \in \mathbb{Z}$, complex numbers exist in the complex plane, $z \in \mathbb{C}$

Theorem 16**Polar Decomposition**

Complex numbers can be expressed in polar form as well

$$z = r(\cos \theta + i \sin \theta) \quad (31.21)$$

Where

$$r = |z| \quad x = r \cos \theta \quad y = r \sin \theta \quad \theta = \tan^{-1} \left(\frac{y}{x} \right) \quad (31.22)$$

This has a number of useful properties

$$z \cdot z' = |z||z'|(\cos(\theta + \theta') + i \sin(\theta + \theta')) \quad (31.23)$$

$$\frac{z}{z'} = \frac{|z|}{|z'|}(\cos(\theta - \theta') + i \sin(\theta - \theta')) \quad (31.24)$$

PROOF Proof for (31.23):

$$\begin{aligned} z \cdot z' &= |z|(\cos(\theta + i \sin \theta)) \times |z'|(\cos \theta' + i \sin \theta') \\ &= |z||z'|(\cos \theta \cos \theta' + i \cos \theta \sin \theta' + i \sin \theta \cos \theta' - \sin \theta \sin \theta') \\ &= |z||z'|[\cos \theta \cos \theta' - \sin \theta \sin \theta' + i(\cos \theta \sin \theta' + \sin \theta \cos \theta')] \end{aligned} \quad (31.25)$$

And the proof follows

□

Lemma 5

A corollary exists

$$z^2 = |z|^2(\cos 2\theta + i \sin 2\theta) \quad (31.26)$$

Theorem 17**Moivre's Theorem**

$$z^n = |z|^n(\cos(n\theta) + i \sin(n\theta)) \quad (31.27)$$

More generally,

$$(\cos \theta + i \sin \theta)^n = \cos(n\theta) + i \sin(n\theta) \quad (31.28)$$

So we may define z to be the n^{th} root of w which implies that

Lemma 6 | Every complex number has a n^{th} root $\forall n$

PROOF

$$\text{Let } z = |w|^{\frac{1}{n}} \left(\cos \frac{\theta}{n} + i \sin \frac{\theta}{n} \right) \quad (31.29)$$

Then

$$w = |w|(\cos \theta + i \sin \theta), \text{ then } z^n = w \quad (31.30)$$

Intuition: define z to be $\frac{1}{n}$ and then take the n^{th} power of both sides to show that $z^n = w$

□

This leads us to the conclusion that representations of complex numbers are not unique⁵⁴.

Table 6. Channel inversion in operation regions

region	source	drain
cut-off	No, $V_{GS} < V_t$	No $V_{GD} < V_t$
triode	Yes, $V_{GS} > V_t$	Yes, $V_{GD} > V_t$
saturation	Yes, $V_{GS} > v_t$	No (pinched off) $V_{DS} > V_{OD}$ or $V_{GD} < V_t$

PROOF | If every z can be written as $z = r(\cos \theta + i \sin \theta)$, then it holds for $\theta + 2\pi n \forall n \in \mathbb{Z}$ since $\sin \theta = \sin(\theta + 2\pi n)$ and $\cos \theta = \cos(\theta + 2\pi n)$. □

31.2.1 Functions on complex planes

Definition 76

Given a domain $\mathbb{D} \in \mathbb{C}$, a function f is a rule such that

$$z \in \mathbb{D} \xrightarrow{f} w \in \mathbb{D} \Leftrightarrow w = f(z) \quad (31.31)$$

Definition 77

We may define \mathbb{D} to be the domain of f

Likewise, range is defined as

$$\text{Ran}\{f\} = \{w \in \mathbb{C} : \exists z \in D : f(z) = w\} \quad (31.32)$$

Example

$$f(z) = \frac{1}{z+i} \quad (31.33)$$

What is the maximum domain of f ?

$$\text{Dom}\{f\} = \{z \in \mathbb{C} : |z| < -i\} \quad (31.34)$$

What is the range of f ?

$$\frac{1}{z+i} = w \quad (31.35)$$

For which values of w can we solve this equation?

$$z = -i + \frac{1}{w} \quad (31.36)$$

So the range of the function is

$$\text{Ran}\{f\} = \{w \in \mathbb{C} : |w| \neq 0\} \quad (31.37)$$

Example

$$f(z) = z^2 + 1 \quad (31.38)$$

It is fairly clear that $\text{Dom}\{f\} = f \in \mathbb{C}$

The range can be found by solving for z in

$$z^2 + 1 = w \quad (31.39)$$

And so

$$\text{Ran}\{f\} = \{w \in \mathbb{C}\} \quad (31.40)$$

31.2.2 Exponential Functions

Definition 78

Given $z = x + iy$,

$$e^z = e^x(\cos y + i \sin y) = e^{\text{Re}\{z\}}(\cos(\text{Im}\{z\}) + i \sin(\text{Im}\{z\})) \quad (31.41)$$

$$1. e^{z+w} = e^z e^w$$

$$2. |e^z| = e^{\text{Re}\{z\}} \neq 0$$

$$3. e^{z+i2\pi n} = e^z$$

PROOF (1) follows from the product rule for complex numbers

(2) follows by definition

(3) follows by definition (recall: writing z w.r.t. \sin, \cos) □

More properties:

- $\text{Dom}\{e^z\} = \mathbb{C}$
- $\text{Ran}\{e^z\} = \{\mathbb{C} \setminus \{0\}\}$
- $e^z = w \quad \text{if } w \neq 0$

arg, or argument is the angle from the real axis to that on the complex plane. Usually denoted by θ

55

⁵⁴ They are part of a cyclic group

$$\begin{aligned} z &= \ln|w| + i \arg w \\ e^z &= e^{\ln|w| + i \arg w} \\ &= e^{\ln|w|} e^{i \arg w} \\ &= |w| \cos(\arg w) + i \sin(\arg w) \\ &= w \end{aligned} \quad (31.42)$$

Remark **Polar representation**

$$w = |w| e^{i \arg w} \quad (31.43)$$

Example | Find polar coordinates of $z = i + 1$

$$r = |w| \quad \theta = \arg w \quad (31.44)$$

$$\begin{aligned}
 |z| &= \sqrt{1+i} = \sqrt{2} \\
 \cos \theta &= \frac{1}{\sqrt{2}} \rightarrow \theta = \frac{\pi}{4} \\
 z &= \sqrt{2}e^{i\pi/4}
 \end{aligned} \tag{31.45}$$

Example Find

$$(1+i)^{\frac{1}{3}} \tag{31.46}$$

Solution: $z = \sqrt{2}e^{\frac{i\pi}{4}} \rightarrow z^{1/3} = 2^{\frac{1}{6}}e^{i\pi/12}$

Definition 79

Trig functions for complex numbers

$$\cos z = \frac{1}{2} (e^{iz} + e^{-iz}) \tag{31.47}$$

PROOF

$$\cos x = \frac{1}{2}(e^{ix} + e^{-ix}) = \frac{1}{2} \left(\cos x + i \sin x + \underbrace{\cos(-x)}_{\text{odd; } = \cos(x)} + \underbrace{i \sin(-x)}_{\text{even; } = -\sin(x)} \right) = \cos x
 \tag{31.48}$$

□

$$\sin z = \frac{1}{2} (e^{iz} - e^{-iz}) \tag{31.49}$$

And a similar proof follows for $\sin z$.

These have the following properties

$$\sin z|_{Im Z=0} = \sin x \tag{31.50}$$

$$\cos(z + 2\pi n) = \cos z \forall n \in \mathbb{Z} \tag{31.51}$$

$$\sin(z + 2\pi n) = \sin z \forall n \in \mathbb{Z} \tag{31.52}$$

PROOF

$$\begin{aligned}
 \cos z + 2\pi n &= \frac{1}{2}(e^{i(z+2\pi n)} + e^{-i(z+2\pi n)}) \\
 &= \frac{1}{2}(e^{iz} e^{i2\pi n} + e^{-iz} e^{-i2\pi n}) \\
 &= \frac{1}{2}(e^{iz} + e^{-iz}) \\
 &= \cos z
 \end{aligned} \tag{31.53}$$

□

The domain of $\{\cos z, \sin z\} = \mathbb{C}$

Range?

Solve $\cos z = w$ for z

$$\begin{aligned}
 \frac{1}{2}(e^{iz} + e^{-iz}) &= w \\
 \dots \times 2e^{iz} \text{ on both sides} \\
 e^{2iz} - 2we^{iz} + 1 &= 0 \\
 \dots \text{Let } S = e^{iz} \\
 S^2 - 2ws + 1 &= 0 \\
 S = w \pm \sqrt{w^2 - 1} &\equiv u
 \end{aligned} \tag{31.54}$$

Now we note that $e^{iz} = u$ can be solved for z for any $u \neq 0$

$$u = 0 \leftrightarrow w = \pm \sqrt{w^2 - 1} \tag{31.55}$$

$$w^2 = w^2 - 1 \text{ impossible for } u \neq 0 \tag{31.56}$$

Therefore:

$$\text{Ran}\{\cos z\} = \text{Ran}\{\sin z\} = \mathbb{C} \tag{31.57}$$

Remark An intuitive way of interpreting this result is thinking of $\{\sin, \cos\}$ being a function that projects values from the complex domain to the real plane; though $\{\sin, \cos\}$ takes on a limited range of values in the real domain, in the complex domain it spans the entire plane. Think: mental model of a complex number spinning around and having that project onto a real line. More formally, see: the [Little Picard Theorem](#)

SUBSECTION 31.3

Lecture 3: Exponent and Logarithm

31.3.1 Exponential

Recall: the complex exponential function \exp is defined as

$$\exp : e^z = e^x(\cos y + i \sin y) \tag{31.58}$$

Where $z = x + iy$.

Properties:

1. $e^{w+z} = e^z e^w$
2. $e^z \neq 0$
3. $e^{2\pi mi} = 1$

The first and third properties imply that the exponential function is a periodic function.

$$e^{z+2\pi mi} = e^z \tag{31.59}$$

Consider the equation

$$e^z = w \tag{31.60}$$

If this has the solution z_* , then $z_* + 2\pi mi, m = 0, \pm 1, \pm 2, \dots$ is also a solution.

31.3.2 Logarithm

Definition 80

$$\log \equiv \log w = \ln |w| + i \arg w \quad (31.61)$$

$$w \neq 0$$

PROOF Proof that $\log w = \ln |w| + i \arg w$

$$\begin{aligned} w &= |w| e^{i \arg w} \\ &= e^{\ln |w|} e^{i \arg w} \\ &= e^{\ln |w| + i \arg w} \\ \rightarrow e^z &= e^{\ln |w| + i \arg w} \\ \Rightarrow z &= \ln |w| + i \arg w \end{aligned} \quad (31.62)$$

□

Note that \arg is a multivalued function, i.e

$$\arg w = \arg(w + 2\pi m_i), i \in \mathbb{Z}, \arg w \in [-\pi, \pi) \quad (31.63)$$

Example

$$e^z = e^5 \quad (31.64)$$

Solve for z

$$z = 5 + 2\pi m i, m \in \mathbb{Z} \quad (31.65)$$

Example

Solve $e^z = i$

$$i = e^{\frac{i\pi}{2}} \quad (31.66)$$

The solution is therefore

$$z = i(\pi/2 + 2\pi m), m \in \mathbb{Z} \quad (31.68)$$

Note that providing only a single solution is wrong; must provide all

$$|i| = 1; \arg i = \frac{\pi}{2} \quad (31.67)$$

The complex logarithm is a multivalued function (like \arg).

$$\log z = \ln |z| + i \arg z, \arg z \in [-\pi, \pi) \quad (31.69)$$

Note that $i \arg z$ denotes the principal branch of \log .

Though it is multivalued, in general, $\log zw \neq \log z + \log w$

Example

Assume $\arg z = \frac{2\pi}{3}$ and $\arg w = \frac{3\pi}{4}$.

$$\arg(zw) = ? \quad (31.70)$$

Typically we would just add them together, i.e. $\frac{2\pi}{3} + \frac{3\pi}{4}$. But this is $> \pi$ which is not allowed as per the definition of \arg , so we must add or subtract something.

Let's try subtracting 2π ⁵⁶

$$\arg(zw) = \frac{17\pi}{12} - 2\pi = -\frac{7\pi}{12} \in [-\pi, \pi) \quad (31.71)$$

⁵⁶since adding $\pm 2\pi$ doesn't change the angle, just rotates it around once

So we proved that in general the arguments don't sum up. But we want to go from here to proving that the logs don't sum up.

$$\begin{aligned}\log zw &= \ln |zw| + i \arg zw \\ &\neq \ln |z| + \ln |w| + i \arg z + i \arg w\end{aligned}\tag{31.72}$$

In general this is not correct because after breaking apart $\arg zw$ $\arg z$ and $\arg w$ when summed can exceed the range allowable for \arg

$$\therefore \log(zw) = \log z + \log w\tag{31.73}$$

Example Compute $\log(\sqrt{3} + i)$.
Just apply the formula.

$$\log w = \ln |w| + i \arg w\tag{31.74}$$

$$\log(\sqrt{3} + i) = \ln \sqrt{4} + i \arg(\sqrt{3} + i) = \ln 2 + i\left(\frac{\pi}{6} + 2\pi m\right), m \in \mathbb{Z}\tag{31.75}$$

31.3.3 Powers

Definition 81

$$\forall a \neq 0, a^z \equiv e^{z \log a}\tag{31.76}$$

Example Complete $(1 + i)^i$

$$\begin{aligned}(1 + i)^i &= e^{i \log(1+i)} \\ \log(1 + i) &= \ln \sqrt{2} + i\left(\frac{\pi}{4} + 2\pi m\right), m \in \mathbb{Z} \\ \dots &= e^{-\frac{\pi}{4} - 2\pi m} e^{i \ln \sqrt{2}}\end{aligned}\tag{31.77}$$

SUBSECTION 31.4

Lecture 4

Definition 82

Analytical Functions

A complex function is a function that maps a complex variable to a complex result. A complex *analytic* function does the same thing, *and* is continuously differentiable over \mathbb{C}

Define \mathbb{D} , an open and connected subset of the complex domain \mathbb{C}

We define

$$f : \mathbb{D} \rightarrow \mathbb{C} \quad \text{to be the analytic of } z_o \in \mathbb{D}\tag{31.78}$$

Now, given f , we can define the complex derivative f'

$$f'(z_0) = \lim_{z \rightarrow z_0} \frac{f(z) - f(z_0)}{z - z_0} \quad \text{exists}\tag{31.79}$$

Noting that

$$z \rightarrow z_0 \leftrightarrow |z - z_0| \rightarrow 0\tag{31.80}$$

(31.79) can be rewritten as

$$f'(z_0) = \lim_{h \rightarrow 0} \frac{1}{h} (f(z_0 + h) - f(z_0))\tag{31.81}$$

$$h = z - z_0; z = z_0 + h$$

Example | Find the complex derivative

$$f(z) = z^n \quad (31.82)$$

PROOF

$$\begin{aligned} f'(z) &= \lim_{h \rightarrow 0} \frac{1}{h} ((z+h)^n - z^n) \\ &= \lim_{h \rightarrow 0} \frac{1}{h} \left(z^n + nz^{n-1}h + \binom{n}{2} z^{n-2}h^2 + \dots + h^n - z^n \right) \\ &= \lim_{h \rightarrow 0} \frac{1}{h} \left(nz^{n-1} + \binom{n}{2} z^{n-2}h + \dots + h^{n-1} \right) \end{aligned} \quad (31.83)$$

... Cancel out terms that go to 0

$$\begin{aligned} &= (z^n)' \\ &= nz^{n-1} \end{aligned}$$

□

$$z = x + iy, \bar{z} = x - iy$$

Example | Find the complex derivative

$$f(z) = \bar{z} \quad (31.84)$$

PROOF

$$\begin{aligned} f'(z) &= \lim_{h \rightarrow 0} (\bar{z} + \bar{h} - \bar{z}) \\ &= \lim_{h \rightarrow 0} \frac{\bar{h}}{h} \end{aligned} \quad (31.85)$$

We then take the limit along the real and imaginary axis separately

$$h = h_1 + ih_2$$

Real:

$$\lim_{h \rightarrow 0} \frac{h_1}{h_1} = 1 \quad (31.86)$$

Imaginary:

$$\lim_{h \rightarrow 0} \frac{-h_2}{h_2} = -1 \quad (31.87)$$

□

So the limit (31.85) does not exist

In the previous two examples we found that z^n is analytic and \bar{z} is not.

Example |

$$f(z) = e^z \quad (31.88)$$

PROOF

$$\begin{aligned}
 f'(z) &= f(z+h) - f(z) \\
 &= e^{z+h} - e^z \\
 &= e^{x+h_1}(\cos(y+h_2) + i\sin(y+h_2)) - e^x(\cos y + i\sin y)
 \end{aligned} \tag{31.89}$$

A Taylor series can be used to expand e^{x+h_1} , $\cos(y+h_2)$, $\sin(y+h_2)$

$$\begin{aligned}
 e^{x+h_1} \cos(y+h_2) \times & \\
 \left(e^x + e^x \frac{h_1}{1!} + \text{higher order terms} \right) \times & \\
 \left(\cos y - \frac{\sin y}{1!} h_2 + \text{higher order terms} \right)
 \end{aligned} \tag{31.90}$$

(31.91)

And then a bunch of terms can be cancelled out to leave us with a couple of terms and a bunch of higher order terms in h_1, h_2

$$= e^x \cos y - e^x \sin y + e^x h_1 \cos y + \dots \text{higher order terms} \tag{31.92}$$

And as a result

$$(e^z)' = \lim_{h \rightarrow 0} \frac{1}{h} e^z h = e^z \tag{31.93}$$

□

So e^z is analytic.

31.4.1 Properties of complex derivative

1. $(f+g)' = f' + g'$
2. $(fg)' = f'g + fg'$
3. $\frac{f'}{g} = \frac{f'g - fg'}{g^2} \cdot f(g(z))' = f'(g(z))g'(z)$

For the 3rd case here, range of $g \in \text{domain of } f$

Example

$$(e^{z^3})' = e^{z^3} (z^3)' = 3z^2 e^{z^3} \tag{31.94}$$

Definition 83

A function f is **entire** if f is analytic in \mathbb{C}

Examples of entire functions include e^z , z^n . Non-analytic functions include $\frac{1}{z}$ since it is not defined at 0 i.e. it is not entire over \mathbb{C}

However, is $\frac{1}{z}$ analytic over the rest of \mathbb{C} ?

PROOF

$$\begin{aligned}
 \frac{1'}{z} &= \lim \frac{1}{h} \left(\frac{1}{z+h} - \frac{1}{z} \right) \\
 &= \lim \frac{1}{h} - \frac{h}{(z+h)z} \\
 &= -\frac{1}{z^2} - \lim -\frac{h}{(z+h)z^2} = -\frac{1}{z}
 \end{aligned} \tag{31.95}$$

Therefore $\frac{1}{z}$ is analytic in $\mathbb{C} - \{0\}$

□

Theorem 18**Cauchy-Riemann equations**

The Cauchy-Riemann equations give us a direct way of checking if a function is differentiable and if it is, it gives us the derivative. It is a consequence of the fact that the limit defining $f(z)$ must be the same no matter what direction z is approached. Namely, if f as defined below is analytic⁵⁷

$$f(z) = u + iv \quad (31.96)$$

Then,

$$f'(z) = \frac{\partial u}{\partial x} + i \frac{\partial v}{\partial x} = \frac{\partial v}{\partial y} - i \frac{\partial u}{\partial y} \quad (31.97)$$

In particular we're interested in the this set of PDEs (which is called the Cauchy-Riemann equations)

$$\begin{aligned} \frac{\partial u}{\partial x} &= \frac{\partial v}{\partial y} \\ \frac{\partial u}{\partial y} &= -\frac{\partial v}{\partial x} \end{aligned} \quad (31.98)$$

The short form is as follows

$$u_x = v_y \quad u_y = -v_x \quad (31.99)$$

If $f = u + iv$ and in \mathbb{D} , then u, v satisfy

$$\frac{\partial u}{\partial x} = \frac{\partial v}{\partial y} \quad (31.100)$$

And

$$\frac{\partial u}{\partial y} = -\frac{\partial v}{\partial x} \quad (31.101)$$

⁵⁷Complex differentiable

PROOF

Since f is analytic, then

$$f'(z) = \lim_{h \rightarrow 0} \frac{1}{h} (f(z+h) - f(z)) \quad (31.102)$$

Which exists and is independent of the way $h \rightarrow 0$.

Take the limits of the real and imaginary parts of (31.102):

Real:

$$\lim_{h_1 \rightarrow 0, h_1 \in \mathbb{R}} \frac{1}{h_1} (f(z+h_1) - f(z)) = \lim \frac{1}{h_1} (f(x+h_1+iy) - f(x+iy)) = \partial_x f(z) \quad (31.103)$$

$$\lim_{ih_2 \rightarrow 0, h_2 \in \mathbb{R}} \frac{1}{h_2} (f(z+ih_2) - f(z)) = \lim \frac{1}{ih_2} (f(x+ih_2+iy) - f(x+iy)) = -\partial_y f(z) \quad (31.104)$$

Since this limit is independent of how $h \rightarrow 0$,

$$\partial_x f(z) = -i \partial_y f(z) \quad (31.105)$$

Recall that $f = u + iv$

$$\partial_x(u+iv) = -i \partial_y(u+iv) = -i \partial_y u + \partial_y v \quad (31.106)$$

Therefore

$$\partial_x u = \partial_y v, \partial_x v = -\partial_y u \quad (31.107)$$

□

Definition 84

Complex derivative:

Using the Cauchy-Riemann equations, we can define the complex derivative of f as

$$\frac{\partial f}{\partial z} \frac{1}{2} \left(\frac{\partial f}{\partial x} - i \frac{\partial f}{\partial y} \right) \quad (31.108)$$

$$\frac{\partial f}{\partial \bar{z}} \frac{1}{2} \left(\frac{\partial f}{\partial x} - i \frac{\partial f}{\partial y} \right) \quad (31.109)$$

Then, via the Cauchy-Riemann equations, we have

$$\frac{\partial f}{\partial \bar{z}} = 0 \quad (31.110)$$

PROOF

Proof of (31.110):

Recall $f = u + iv$

Plug this into the LHS of the expression:

$$\frac{\partial u}{\partial x} + i \left(\frac{\partial v}{\partial y} \right) + i \left(\frac{\partial u}{\partial x} + i \frac{\partial v}{\partial y} \right) \Rightarrow \partial_x u - \partial_y v = 0, \partial_y u + \partial_x v = 0 \quad (31.111)$$

Which is the Cauchy-Riemann equations (31.107)

□

Example

Is $f(z) = e^{z^5} \cdot \sin z \cdot \bar{z}^3$ analytic?

$$\frac{\partial f}{\partial z} e^{z^5} \cdot \sin z e \bar{z}^2 \neq 0 \quad (31.112)$$

| So f is not analytic

Example | Is $f(z) = |z|^6$ analytic?

PROOF

$$\frac{\partial y}{\partial \bar{z}} = \frac{\partial}{\partial z} (z\bar{z})^3 = z^3 \cdot 3\bar{z}^2 \neq 0 \quad (31.113)$$

□

| f is not analytic.

Now that we know that analytic functions satisfy the Cauchy-Riemann equations, we can use this to prove that the converse holds, i.e. that non-analytic functions do not satisfy [the Cauchy-Riemann equations]

Theorem 19 Given $f(x, y)$ continuously differentiable and satisfies the Cauchy-Riemann equations, then f is analytical

PROOF

$$\begin{aligned} f(x + h_1, y + h_2) - f(x, y) &\xrightarrow{\text{Taylor series}} \\ &= f(x, y) + \partial_x f(x, y)h_1 + \partial_y f(x, y)h_2 + \text{H.O.T} - f(x, y) \\ &\dots \text{cancel terms} \dots \\ &= \partial_x = \frac{\partial f}{\partial z} + \frac{\partial f}{\partial \bar{z}} \\ &= \partial_y = \frac{\partial f}{\partial \bar{z}} - \frac{\partial f}{\partial z} \frac{1}{i} \end{aligned} \quad (31.114)$$

Note abbreviation higher order terms \Leftrightarrow H.O.T

And then plugging this back into the initial expression we get

$$\begin{aligned} f(x + h_1, y + h_2) - f(x, y) &= \left(\frac{\partial f}{\partial z} \frac{\partial f}{\partial \bar{z}} \right) h_1 + \left(\frac{\partial f}{\partial \bar{z}} \frac{\partial f}{\partial z} \right) \frac{1}{i} h_2 + \text{H.O.T} \\ &= \frac{\partial f}{\partial z} (h_1 + ih_2) + \text{H.O.T} \\ &= \frac{\partial f}{\partial z} h + \text{H.O.T} \\ f'(z) &= \lim_{h \rightarrow 0} \frac{1}{h} \left(\frac{\partial f}{\partial z} h + \text{H.O.T} \right) \\ &= \frac{\partial f}{\partial z} z \text{ exists} \end{aligned} \quad (31.115)$$

So therefore f is analytic

□

Example | Is $\sin z$ analytic?

$$\frac{\partial}{\partial \bar{z}} \sin z = 0 \rightarrow \sin z \text{ is analytic} \quad (31.116)$$

| In 'slow motion',

$$\begin{aligned}
 \frac{\partial \sin z}{\partial \bar{z}} &= \frac{1}{2} \left(\frac{\partial \sin z}{\partial x} + i \frac{\partial \sin z}{\partial y} \right) \\
 \sin z &= \frac{1}{2i} (e^{iz} - e^{-iz}) \\
 e^{iz} &= e^{ix-y} = e^{-y}(\cos x + i \sin y) \\
 \frac{\partial e^{iz}}{\partial \bar{z}} &= \frac{1}{2} (\partial_x + i \partial_y) \\
 &= \dots = 0
 \end{aligned} \tag{31.117}$$

Therefore $\sin z$ is analytic

TLDR of the lecture; one can find if the function is analytic by checking if the Cauchy-Riemann equations hold. This can be done by taking the complex derivative of the function w.r.t z and \bar{z} . Then, by the theorem we proved, if $\frac{\partial f}{\partial \bar{z}} = 0$, then f is analytic and $f'(z) = \frac{\partial f}{\partial z}$

SUBSECTION 31.5

Power series

Definition 85

A **Power series** is an expression of the form

$$\sum_{n=0}^{\infty} a_n (z - z_0)^n \tag{31.118}$$

Where a_n is a coefficient and z_0 the centre of the series. The power series diverges if it is not converging absolutely, which it does at z_* if

$$\sum |a_n| (z_* - z_0)^n \tag{31.119}$$

converges

Theorem 20

There exists radius of convergence $R \geq 0$ such that

- The series converges absolutely in the disk $|z - z_0| < R$
- The series diverges for $|z - z_0| \geq R$

If

$$\lim_{n \rightarrow \infty} |a_n|^{1/n} \tag{31.120}$$

exists, then

$$\frac{1}{R} = \lim_{n \rightarrow \infty} |a_n|^{1/n} \tag{31.121}$$

Lemma 7

If $\lim |a_{n+1}/a_n|$ exists, then

$$\frac{1}{R} = \lim_{n \rightarrow \infty} |a_{n+1}/a_n| \tag{31.122}$$

Example

Find the radius of convergence for

$$\sum_{n=0}^{\infty} \sqrt{n} z^n \tag{31.123}$$

PROOF | Let $z_0 = 0, a_n = \sqrt{n}$

$$\begin{aligned} \lim_{n \rightarrow \infty} \sqrt{n}^{\frac{1}{n}} &= \lim_{n \rightarrow \infty} n^{1/2n} \\ \ln n^{1/2n} &= \frac{1}{2n} \ln(n) \text{ which } \rightarrow 0 \text{ as } n \rightarrow \infty \\ n^{\frac{1}{2n}} &\rightarrow e^0 = 1 \Rightarrow R = \frac{1}{1} = 1 \end{aligned} \quad (31.124)$$

□

Theorem 21 If $\sum a_n z^n$ and $\sum b_n z^n$ have radius of convergence of at least R

1.

$$\sum a_n z^n + \sum b_n z^n = \sum (a_n + b_n) z^n \quad (31.125)$$

and has radius of convergence of at least R

2.

$$\sum a_n z^n \cdot \sum b_n z^n \quad (31.126)$$

has radius of convergence equal to R

3.

$$\sum a_n z^n \quad (31.127)$$

is complex differentiable (and therefore analytic) in $|z| < R$ and it's derivative is

$$\sum a_n n z^{n-1} \quad (31.128)$$

with radius of convergence R

4.

$$a_n = \frac{1}{n! f^n|_0} \quad (31.129)$$

where

$$f(z) = \sum_{n=0}^{\infty} a_n z^n \quad (31.130)$$

SUBSECTION 31.6

Lecture 5

ECE444: Software Engineering

SECTION 32

Preliminary

SUBSECTION 32.1

Lecture 1, 2

Taught by Prof. Shurui Zhou

PART

VII

- Software engineering is different from what coding is; design, architecture, documentation, testing, etc v.s. just script kiddie-ing
- *Vasa syndrome*
- Rockstar engineers are a myth

SECTION 33

Project Management

SUBSECTION 33.1

Lecture 3

Definition 86

Conway's law states that 'Any organization that designs a system (defined broadly) will produce a design whose structure is a copy of the organization's communication structure'.

The waterfall method is slow and costly and defects can be extremely costly, especially early on in the development lifecycle.

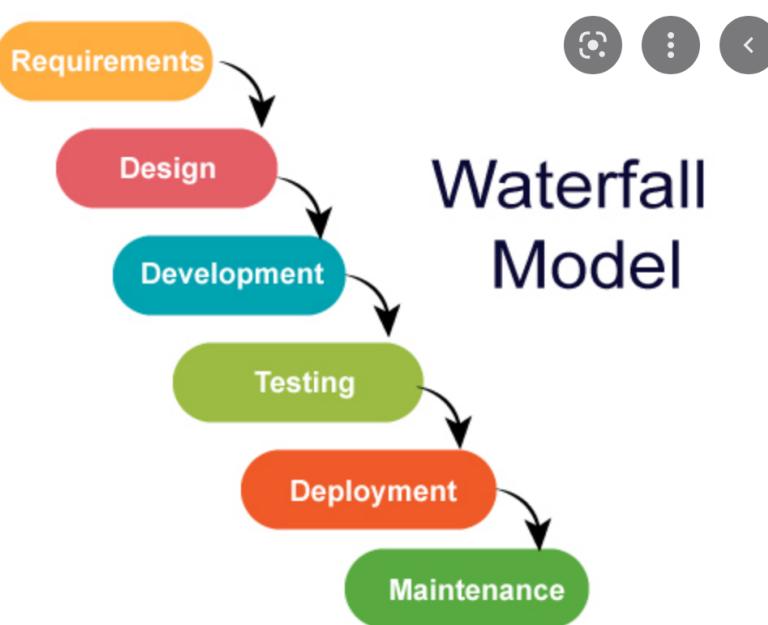


Figure 70. Waterfall method

In order to address this the V model was introduced which increases the amount of testing to reduce the possibility of having to rework everything

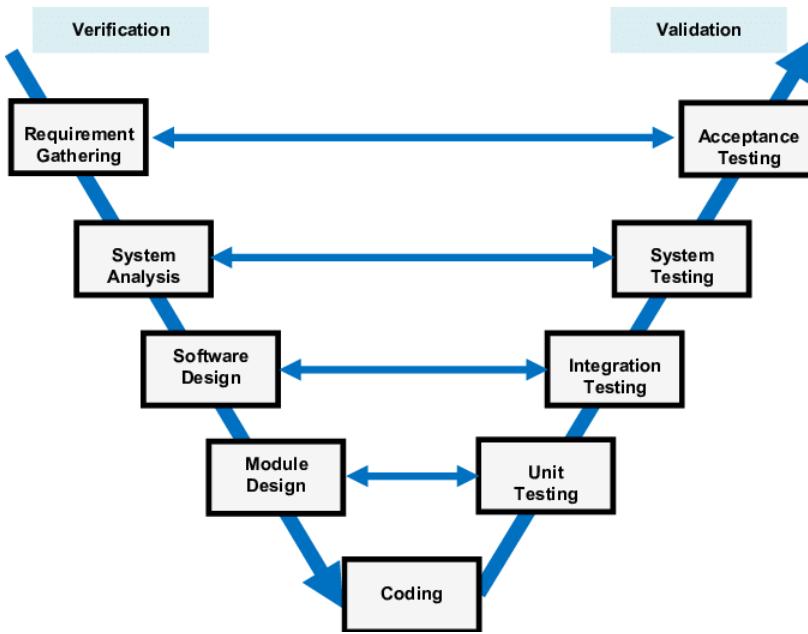


Figure 71. V model

Generally speaking the waterfall model isn't used much anymore due to the reality that software specifications change on a near daily basis.

Recall: aUToronto Spring 2022 integration hell

33.1.1 Agile

Agile is a project management approach which, in most general terms, seeks to respond to change and unpredictability using incremental, iterative work (sprints). This allows for a balance between the need for predictability and the need for flexibility. Some agile methods include:

- Extreme programming: really really fast iteration (think days)
- Scrum: 2-4 week sprints with standups and backlogs; sticky notes for tasks, etc. Think kanban boards. Daily scrum meetings to unblock ASAP. Development lifecycle is therefore a series of sprints.
- On-site customer; frequent interaction with end users to figure out what exactly they need.

SUBSECTION 33.2

I dropped this course

I decided to drop this course because the courseload was a little too much to handle between EngSci ECE, clubs, design teams, work, and trying to have a life.

PART

VIII

What are notes?